

ALMA MATER STUDIORUM — UNIVERSITÀ DI BOLOGNA

---

DEIS - Dipartimento di Elettronica, Informatica e Sistemistica  
PhD Course in Electronics, Computer Science and Telecommunications

Cycle XXVII  
Competitive Sector: 09/H1  
Disciplinary Sector: ING-INF/05

## ENGINEERING COMPLEX COMPUTATIONAL ECOSYSTEMS

*Candidate*

Dott. Ing. DANILO PIANINI

*Supervisor*

Chiar.mo Prof. Ing. MIRKO VIROLI

*Tutor*

Chiar.mo Prof. Ing. ANTONIO NATALI

*Coordinator*

Chiar.mo Prof. Ing. ALESSANDRO VANELLI CORALLI

---

FINAL EXAMINATION YEAR 2015



*Lorem ipsum dolor sit amet.  
  consectetur adipiscing elit.  
  Sed fringilla quis mauris id sagittis everything.*

## **Acknowledgements**

Curabitur commodo dictum risus laoreet tincidunt. Sed dapibus nec ex sit amet consequat. Quisque cursus est sit amet lectus tempor, nec egestas sapien rutrum. Sed dapibus consequat egestas. Quisque blandit, tellus et molestie interdum, augue est molestie lorem, tristique congue metus massa sed eros. Vivamus fermentum erat a faucibus porta. Praesent sit amet risus leo. Integer venenatis lectus sed euismod euismod. Ut pulvinar fermentum sagittis. Aliquam maximus nisl velit, ac varius tellus dignissim in. Fusce nibh dolor, blandit vel nisl non, vehicula tincidunt lorem. Suspendisse fringilla magna ac justo fermentum, nec accumsan odio sollicitudin. Maecenas consectetur, nulla sit amet ultricies pretium, metus sapien posuere turpis, in lobortis tellus turpis ac ipsum. Aliquam sollicitudin augue a aliquam volutpat. Mauris bibendum nunc id est ullamcorper, nec feugiat elit dapibus.

Duis tincidunt maximus justo, id convallis mauris mattis congue. Maecenas ullamcorper laoreet lacinia. Praesent luctus dictum metus, sed ultrices dui. Fusce fringilla eu est sit amet porta. Aliquam ornare eleifend congue. Aliquam orci urna, accumsan nec metus a, tincidunt pulvinar augue. Nunc eu vulputate lectus. Morbi placerat varius purus at scelerisque. Mauris malesuada ut massa non porta. Quisque ac efficitur odio. Aliquam scelerisque dapibus felis in ullamcorper. Pellentesque ullamcorper massa quis nibh suscipit elementum. Suspendisse tincidunt, sem non porta consectetur, arcu mauris mollis mi, ut commodo libero diam ut risus.

*Danilo Pianini, March 2015*



# Contents

<b>Abstract</b>	<b>ix</b>
<b>I Background and Motivation</b>	<b>1</b>
<b>1 Pervasive computing</b>	<b>3</b>
1.1 Smart, portable devices . . . . .	3
1.2 Communication technologies . . . . .	7
1.2.1 International Mobile Telecommunications . . . . .	7
1.2.2 WiFi . . . . .	9
1.2.3 Bluetooth and Bluetooth LE . . . . .	9
1.2.4 NFC . . . . .	10
1.3 Towards a P2P pervasive continuum? . . . . .	12
<b>2 Self-organisation</b>	<b>15</b>
2.1 Software ecosystems . . . . .	15
2.2 Nature inspiration . . . . .	16
2.2.1 Physical-inspiration . . . . .	16
2.2.2 (Bio)chemical-inspiration . . . . .	17
2.2.3 Stigmergy . . . . .	17
2.3 Spatial patterns . . . . .	18
2.3.1 Gossip . . . . .	18
2.3.2 Gradient . . . . .	18
2.3.3 Gradcast . . . . .	19
2.3.4 Voronoi partition . . . . .	19
2.4 Tuple-based coordination . . . . .	19
2.4.1 SAPERE . . . . .	20
2.5 Aggregate programming . . . . .	21
2.5.1 Proto . . . . .	23
2.5.2 Field Calculus . . . . .	24
2.6 Engineering and tools . . . . .	25
2.6.1 General purpose frameworks . . . . .	27
2.6.2 Specific simulators . . . . .	27
<b>3 Shortcomings of coordination infrastructures</b>	<b>29</b>
3.1 Engineering emergence . . . . .	29
3.2 Local to global . . . . .	29

<b>4 Shortcomings of existing tools</b>	<b>31</b>
4.1 Specific tools: expressiveness . . . . .	31
4.2 General purpose tools: performance . . . . .	31
<b>II Contribution</b>	<b>33</b>
<b>5 An integrated toolchain for pervasive ecosystems</b>	<b>35</b>
5.1 Chemical-inspired engine . . . . .	35
5.1.1 Gillespie's SSA as an event-driven algorithm . . . . .	35
5.1.2 Gillespie's optimised versions . . . . .	36
5.1.3 From SSA to full fledged DES . . . . .	37
5.2 Meta-meta model . . . . .	39
5.3 Architecture . . . . .	44
5.3.1 Writing a simulation . . . . .	46
5.3.2 Implementation details . . . . .	46
5.4 Performance . . . . .	47
5.5 Collocation in literature . . . . .	48
5.5.1 ALCHEMIST as a DES . . . . .	48
5.5.2 ALCHEMIST as a MABS . . . . .	49
5.6 Meta-model agnostic features . . . . .	51
5.6.1 Distributed statistical analysis . . . . .	51
5.6.2 Real world maps . . . . .	59
5.7 Chemical meta-model . . . . .	61
5.7.1 Example scenario: Morphogenesis . . . . .	61
5.8 SAPERE meta-model . . . . .	69
<b>6 Self-organisation patterns</b>	<b>71</b>
6.1 Crowd evacuation . . . . .	71
6.2 Crowd steering . . . . .	71
6.2.1 Reference scenario . . . . .	71
6.2.2 Steering strategy . . . . .	72
6.2.3 Simulation in ALCHEMIST . . . . .	74
6.3 Self composition . . . . .	76
6.4 Semantic resource discovery . . . . .	76
6.5 Anticipative adaptation . . . . .	76
6.6 Crowd disasters prediction . . . . .	76
6.7 Crowd steering at the urban scale . . . . .	76
6.7.1 Devices and physical configuration . . . . .	76
6.7.2 Distributed crowd steering . . . . .	77
6.7.3 Simulation in ALCHEMIST . . . . .	78

<b>7 Aggregate programming languages</b>	<b>81</b>
7.1 Higher order functions in field calculus . . . . .	81
7.1.1 Impact on alignment . . . . .	81
7.2 Protelis: practical aggregate programming . . . . .	81
7.2.1 Syntax . . . . .	82
7.2.2 Ordinary Language Features . . . . .	82
7.2.3 Special Field Calculus Operators . . . . .	84
7.3 Architecture of Protelis . . . . .	86
7.4 Application examples . . . . .	86
7.4.1 Rendezvous at a Mass Event . . . . .	88
7.4.2 Network Service Management . . . . .	90
<b>III Conclusion</b>	<b>93</b>
<b>8 Results achieved</b>	<b>95</b>
8.1 Integrated toolchain for pervasive ecosystems . . . . .	95
8.2 Aggregate programming languages . . . . .	95
8.2.1 Higher Order Functions in Field Calculus . . . . .	95
8.2.2 Protelis . . . . .	95
<b>9 Future and ongoing work</b>	<b>97</b>
9.1 Aggregate programming . . . . .	97
9.2 Biochemical meta model for ALCHEMIST . . . . .	97
<b>Bibliography</b>	<b>98</b>



# Abstract

This work presents advancements of the latest three years in the engineering techniques for self-organising pervasive ecosystems of devices and services. The inherent complexity of such systems poses new challenges to those who try to dominate the complexity by applying the principles of engineering.

The recent growth in number and distribution of devices with decent computational and communicational abilities, that got suddenly accelerated with the massive diffusion of smartphones and tablets, is envisioning a world with a much higher density of devices in space. This already high device density is probably going to consistently rise if the diffusion of wearable devices gets momentum. Also, communication technologies seem to be focussing on short-range device-to-device (P2P) interactions, with technologies such as Bluetooth® Low Energy and Near-Field Communication getting more and more diffused.

Locality and situatedness become key to provide the best possible experience to users, and the classic model of a centralised, enormously powerful server gathering data and processing it is likely to get less and less efficient with device density. Accomplishing complex global tasks without a centralised controller responsible of aggregating data from devices, however, still is a challenging task. In particular, it is hard to understand which device-local programs could properly interact and guarantee a certain global service level. Such local-to-global issue makes the application of engineering principles challenging at least.

In this work, I lay the foundations of my contribution by first analysing the state of the art in coordination systems, namely in those software frameworks devoted to control and promote interactions among independent software entities. I then motivate my work, by describing the main issues of pre-existing tools and practices and identifying the improvements that would benefit the design of such complex software ecosystems. My contribution is described in Part II, and can be divided in three main branches: i) a novel simulation tool for pervasive ecosystems, ii) introduction of novel and improvements over existing self-organisation patterns, iii) the creation of a new language and interpreter based on “field calculus” and its integration with the previously mentioned simulator. Finally, I draw conclusions and future works.



# **Part I**

## **Background and Motivation**



# 1

## Pervasive computing

It is no mystery that, with the huge progresses of miniaturisation, computational-capable devices are populating the world. The diffusion got momentum with affordable “personal computers” that made their way in a one-device-per-family world. Laptops boosted the process, offering user a personal and mobile device.

The pervasive revolution, however took place when the phones became “smart”, and with the subsequent improvements in the communication technologies. In this chapter, I try to briefly walk the path of success of personal mobile devices, describing also the probable newcomer, namely the wearable devices. I also focus on the current status of the communication protocols, their range and usage, and I try to foresee which world are we going to build if the current trend continues.

### 1.1 Smart, portable devices

When Apple in 2007 released the first iPhone (Figure 1.1), a mobile revolution started. Even though other manufacturers proposed products similar to iPhone under the point of view of communication technologies and computational capabilities in the same time frame (e.g. Nokia N810), the Apple’s smartphone was the first gaining widespread adoption. It is not really relevant for this work to understand if the branding, the design, the multi-touch finger based interaction UI or the feature set was the key of its commercial success: what really matters is that, starting 2007, every person began to carry with her a personal device featuring both the abilities to communicate and compute. The reason, besides the success of iPhone in the higher segment of the phone market, is mainly to attribute to the widespread diffusion of similar but cheaper and less powerful devices in the lower segments. This kind of devices ultimately pushed the market share of the dominating today’s mobile operating system, Google’s Android.

This trend towards a higher diffusion got another leap forward three years later, in 2010, when a device with a feature set similar to iPhone, but with no phone abilities and a bigger screen was released: the iPad. As iPhone gave new vitality and perspectives to the phone arena, iPad revitalised a market that was languishing: the tablets. Tablet devices as they were conceived before iPad launched were nothing more than small laptops with a screen that could be rotated or detached from the keyboard, and a touch-screen normally used with the help of a pen. They



Figure 1.1: Apple iPhone is one of the first devices without any physical keyboard, and the first smartphone to gain worldwide success. *Source: Wikimedia.*

changed from devices designed for a professional niche to widespread tools, up to the point that Gartner forecasts their shipments to overtake in 2015 those of desktop PC and laptops aggregated<sup>1</sup>.

The new frontier of pervasive is probably the wearable technology. Smartphones were precursors, they substituted mobile phones introducing new features and they have potential to substitute our wallets (see, for instance, payments through NFC technology) and keys, becoming the only object we need to carry with us in our pockets. Still there are other accessories which are hard to replace, above all watches and glasses. Yes, smartphone can easily show the current time accurately (and, to be honest, also feature phones had this feature well before 2007), but they require the user to pick them from the pocket, turn them on, and sometimes, depending on privacy settings, also unlock them. The whole operation, takes a much longer time and higher effort with respect to just rotating a wrist. This might be the reason why, despite the explosive expansion of smartphones, the wristwatches market did not decline, as an analysis from MarketWatch points out<sup>2</sup>.

There is room for manufacturers to create new portable devices. Sony, starting 2012, has produced a series of smartwatches, such as those in Figure 1.2 that pair with a smartphone and provide quick access to some of its functionalities. The success of such solution is not huge, but despite that many other companies are interested in this market: Samsung, Motorola and

---

<sup>1</sup><http://www.gartner.com/newsroom/id/2791017>

<sup>2</sup><http://www.marketwatch.com/story/the-watches-time-isnt-up-2013-07-01>



Figure 1.2: Smart watches. From left to right: Sony Smartwatch (*Source: Alex S.H. Lin*), Samsung Galaxy Gear (*Source: Karlis Dambrans*), Motorola Moto 360.

Apple presented devices meant to replace the classic wristwatch, a clear sign that this market is in expansion. At the time of writing, the main issues that slow the widespread adoption of such solutions are battery duration and dependence on a smartphone.

Another notable attempt to make a common accessory smarter is Google Glass project, depicted in Figure 1.3. Their goal is to enhance the experience of wearing glasses by attaching a device with a camera, an optical head-mounted display, and the abilities to locate itself and communicate with other devices. Google Glass, at the time of writing, are way too expensive (with the kit sold at \$1500) for being able to penetrate the general public, but they are an interesting anticipation of possible future devices.

On the same line of such wearable devices, a discrete success is being achieved by the so called “fit bands”. They are bracelets equipped with low energy sensors, mainly accelerometers and gyroscopes, which are used to keep track of user’s activity. Depending on the model, they can be used to monitor some user’s health parameters, such as the number of steps walked per day or heartbeats. They normally work along with another device, a smartphone or a tablet. Such devices, due to their precise market niche and reasonably low price (the Chinese manufacturer Xiaomi recently introduced a low-end wristband at around \$15) are having a notable success.

The wearable devices segment also includes less common devices such as “smart shoes” and materials that can be used to make clothing, such as e-textiles. It is a market in expansion, greatly benefiting from recent increases in performance per watt efficiency. If the trend continues, it is likely that we will more and more powerful wearable devices on sale at cheaper and cheaper prices, and a consequent widespread diffusion. The same sort may occur to other parts of our life: kitchen gear, indoor lights and many other objects are getting more and more “smart” around us. We may, literally, end up with a world where every single object embeds computational and communicational abilities.

A problem arises: how can software engineers deal with such a complexity?

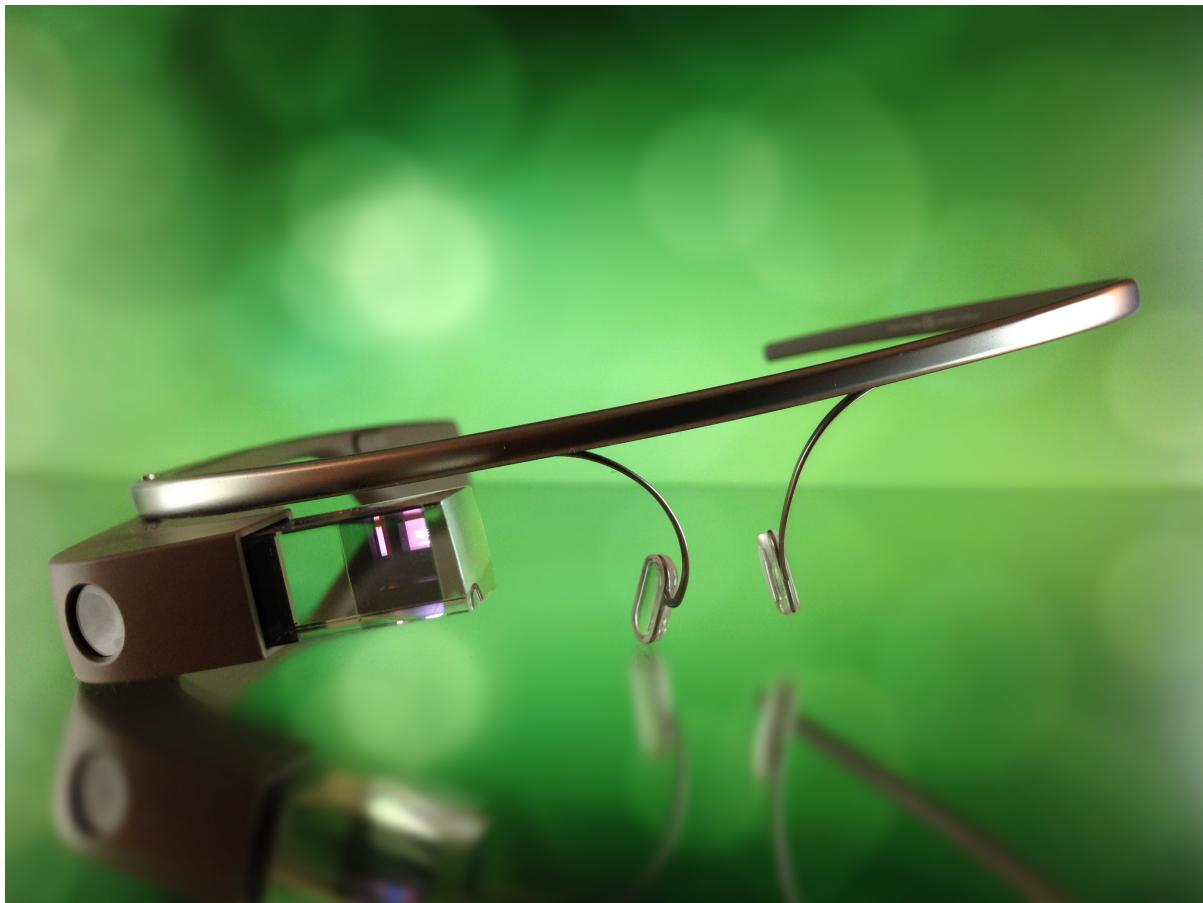


Figure 1.3: Google Glass. In this image, it is possible to see both the camera (on the left hand side) and the semi-transparent head-mounted display. Are those devices going to be part of our everyday life? *Source: Wikimedia*.

## 1.2 Communication technologies

Besides miniaturisation, and as a consequence the increase of computational density in space, another factor played a fundamental role in the world of pervasive computing: the ability to communicate, and in particular the ability to rely on wireless communication, which is of paramount importance when considering mobility.

In later years, many communication means arose. They largely differ in terms of range, protocols, and availability. In this section, I try to resume the most diffused technologies available on today's devices, but the reader is warned: keep in mind that such technologies are evolving very quickly, and the scenario is incredibly fluid.

### 1.2.1 International Mobile Telecommunications

This first mean of communication is designed to allow mobile devices to access the Internet from anywhere in the world, relying on the existing mobile phone infrastructure. Such technologies are meant to be used with the standard IP protocols, and they are normally used to get access to public services, in particular to the world wide web. They are not designed for a local peer-to-peer (P2P) communication, and as a consequence they provide no mean to exploit locality. The diffusion of such communication protocols is widespread, in particular among smartphones. Due to the fact that they rely on the mobile phone network, they require a contract with a mobile telecommunications provider, and as such they are much less diffused in tablets and other portable devices.

The possibility of accessing the Internet from everywhere is probably one of the key bricks that allowed for the huge success of smartphones in today's world. As Figure 1.4 shows, the bandwidth available grew exponentially with time, to the point that in some countries (e.g. in Italy, at the time of writing) the best available mobile connections offer a higher performance than the best available home connection<sup>3</sup>. Such performance unlock the possibility of fully exploiting the possibilities of the world wide web, including cloud services and fruition of multimedia content.

If bandwidth is not currently an issue for international mobile telecommunications, the situation is well different when it comes to device density. Any of us probably experienced network availability issues when participating crowded events, such as concerts or sport events. The current technology, in fact, makes all the network user share the same physical resources: when the device density is too high, there is simply not enough space in the frequency spectrum to grant a decent bandwidth to everyone. Future networks (5G, and presumably those that will follow) are focussing toward this issue among others [ASS<sup>+</sup>11]. In particular, a so-called “spectrum crunch” is expected due to the expected traffic increase (thousand fold over this decade and still growing into the next), that could not be faced simply with the foreseen steady increase

---

<sup>3</sup>At the time of writing, Telecom Italia Mobile offers mobile connections on LTE with a download bandwidth up to 225Mb/s. Fastweb, the company offering the faster solutions for fiber-to-home connections, goes up to 100Mb/s.

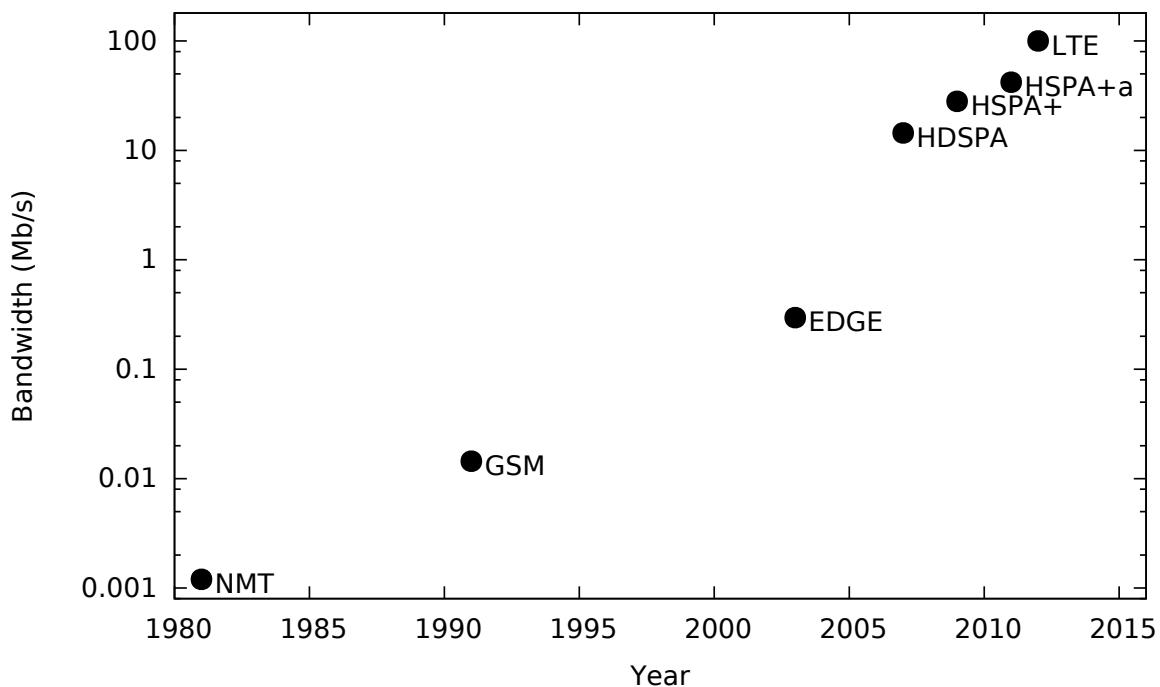


Figure 1.4: Maximum download bandwidth available for mobile devices with time. Each point is labelled with the specific communication technology name. For the most recent multiple-input-multiple-output (MIMO) technologies, such as LTE, a conservative single input channel was used.

of the spectrum allocated for mobile communication, and will require technological advances [AARC<sup>+</sup>12].

### 1.2.2 WiFi

WiFi technology is the most diffused technology for wireless local networking. It is widely diffused, integrated in all smartphones and tablets and also in other devices, such as printers, gaming consoles and TVs. WiFi devices communicate on a distance that ranges from 20 to 100 meters, depending on the condition of the wireless medium and on the power of the communication devices. The communication speed between two linked devices ranges from 56Mb/s to 300Mb/s.

WiFi was designed to provide wireless access to a local area network. In the most classic “infrastructure mode”, wireless devices get connected to a so called access point, which is responsible to route packets among wireless devices and bridge the wireless local area network to the wired backbone. Multiple access point that share the network name (SSID) may be connected using wired network technologies, and they will appear as a single, bigger access point. It is also possible to drop the wired backbone, but specific access points are required.

Also, some WiFi devices provided “ad-hoc mode”, allowing multiple devices to directly communicate without an intermediate access point. This working mode was problematic, mainly due to the fact a standard communication protocol for peer-to-peer WiFi communication was missing. This lack was filled with WiFi Direct, which provides a protocol by which one of the devices that want to communicate directly becomes the access point, allowing for direct communication. The most common usage of such a feature are direct file sharing between devices and connection to peripheral devices such as printers or scanners.

### 1.2.3 Bluetooth and Bluetooth LE

Bluetooth is a technology designed for building energy efficient personal area networks (PANs). Bluetooth devices are assigned a class which identifies the maximum permitted power and, consequently, the maximum operating range. For the most powerful (and power hungry) devices the communication range can go up to 100m. The communication speed ranges from 1Mb/s of the earliest 1.0 version to the 24 Mb/s of version 3.0 and later.

The most interesting features of Bluetooth are not the bandwidth nor the range (WiFi performs better on both), but rather the simple association process and the low power consumption. Thanks to those features Bluetooth found widespread diffusion as a mean to connect low consumption peripherals, such as headsets. Also, it is diffused in cars, and allows user to use the car’s audio system as a speakerphone for making calls or listen to music.

A technology which is often associated with Bluetooth but that is actually a separated and not compatible protocol is Bluetooth LE. The reason why such technologies get associated is that, since the radio frequency used is the same (2.4GHz), dual mode devices can share a single radio antenna. LE stands for Low Energy, and it is the main difference between the two protocols: at



Figure 1.5: A iBeacon, compared to a 2 € coin.

the expense of some bandwidth, Bluetooth LE consistently reduce the amount of energy required. Bluetooth LE applications are particularly interesting, and range from health care to fitness to alerts to proximity sensing.

Proximity can be estimated using the received signal strength indicator (RSSI), and the very low power consumption of Bluetooth LE allowed for the realisation of electronic leashing systems, namely systems where an electronic device is paired to an object and can be used in order to compute the relative position. The applications are, for instance, finding of misplaced, out-of-sight devices (when the electronic device is paired with a movable object) and indoor localisation (if the electronic device is located on a still standing object). Relying on this technology, Apple created iBeacon, namely very small (coin sized, see Figure 1.5) electronic devices consisting basically of a battery and a Bluetooth LE device. iBeacons can be attached to objects, and they send a universally unique identifier (UUID) to enabled smartphones in range. If the smartphone can associate the UUID with a position, it can deduce its location relative to the iBeacon. The low energy feature plays a fundamental role in this kind of applications: beacons whose battery would last few hours would be of little practical use. With current technologies, a beacon device can be powered by a standard, rather cheap battery for several months, up to a couple of years. This technology is probably the prelude to precise indoor localisation.

### 1.2.4 NFC

Near Field Communication, or NFC, is a technology designed for low energy communication between two devices in proximity (typically few centimetres). It is designed for low energy

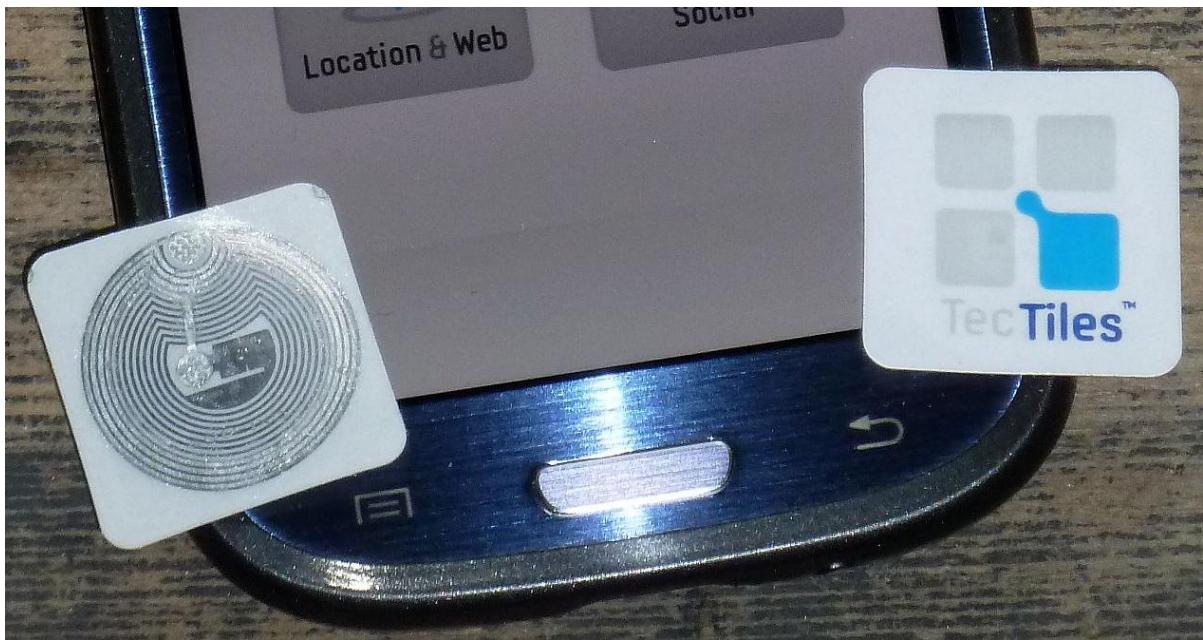


Figure 1.6: Two NFC tags, mounted on stickers. In order to understand the size, they are placed above a Samsung Galaxy S3, a 4-inches smartphone. Moreover, smaller versions of such tags exists, the models pictured here are very common-sized.

consumption rather than high bandwidth: its speed (depending on the specification) ranges from 126kb/s to 242kb/s.

One of the most interesting features of NFC is that one of the two devices (the so-called “tag”) can be completely passive, and still carry a small amount of information within (currently between 96 and 4096 bytes). No battery or energy source is required, the information included can be read by active NFC devices in proximity. Having no need of battery at all, NFC tags can really be tiny, even smaller than iBeacons. Figure 1.6 shows two NFC tags: compared with iBeacons, they can be much lighter and thinner. Brought to the world of humans, it is something like creating a sticker with very small sized text: those who have a powerful enough magnifying glass and are close enough to use it properly can read what it says.

The range of applications of such technology is rather broad. The one which was probably most sponsored is contact-less payment, namely the ability to pay just tapping the phone close to the check-out counter. This is a very interesting possibility, and indeed in 2007 there were enthusiastic forecasts [OP07] about its quick diffusion, that did not happen as quick as expected. A number of studies tried to understand the reasons behind this slow adoption, and it appears that reasons are more correlated to marketing and management rather than technological maturity [PH14, Apa13, TOCH14]. Similarly, if the phone stores identity or access tokens, NFC is a very suitable technology for effectively using such tokens: in this case, its very low range is a nice feature. NFC can also be used as a technology enabler, namely as a mean to securely bootstrap

another connection, or join a local network of devices. An example of such use is the Android Beam technology, that relies on NFC in order to establish a Bluetooth connection between two Android devices, transfer a file, then close the Bluetooth connection. Technically, NFC could be used to directly transfer files, but both WiFi and Bluetooth offer much wider bandwidths and range, and as a consequence are preferred for such task. Another interesting application field is mobile device automation: it is achieved by attaching Generally speaking, NFC comes in handy when there is need of a communication mean whose range should be very limited.

### 1.3 Towards a P2P pervasive continuum?

We are living exciting times. In about five years from the introduction of the technology on the market, almost everybody got a personal smart device always with her. Miniaturisation and power efficiency is constantly growing at a stunning rate, allowing data, communication systems and computation to be spread around in our physical world.

In few decades, we will probably witness the diffusion on computation on everyday object. In such a scenario, the device density will be much higher if compared to the current, up to the point that the aggregation of devices participating the system could be seen as a “pervasive continuum” [ZCF<sup>+</sup>11]. This continuum is studied under a number of names, including pervasive computing, smart cities, and the Internet of Things, in order to provide a wealth of services in an un-intrusive manner [ZOA<sup>+</sup>14, CDB<sup>+</sup>12, Zam12, HDBDM12].

One of the possible strategies is to connect every single device to the Internet, aggregate its information in a remote server, do the necessary computation, then send back eventual results where they are needed. This is the strategy behind cloud computing, which is achieving great success. In particular, we argue, this strategy is interesting when the information could or should be aggregated with information from other, distant sources, or conserved for historical purposes. This path, however, gets harder and harder to follow with device density: besides the obvious increasing on the total information produced, and consequently of the information to transmit and process, there are two other problems: the saturation of the wireless medium, and the locality of information.

Who tried to use its own smartphone in a very crowded environment has probably experienced connection or network issues. The problem, as discussed in Section 1.2.1, is that current technology must share a common medium among all the devices in the same area. Increasing the maximum number of devices per area is one of the goals of the next generation of international mobile telecommunication technologies. One of the proposed approaches is to switch to a very dense array of very small cells by deploying multiple antennas at a very short distance one another, e.g. inside the public illumination poles. Devices nearby the local antenna would connect to it, and the antenna would then connect them to the rest of the network transparently, in a way somehow similar to the current WiFi “infrastructure mode”. Clearly, diffusing such antennas can possibly represent a major infrastructural upgrade, and, potentially, cost.

Another observation is that not every device needs direct access to the Internet to be able

to accomplish its task, and this is increasingly true with increased density. Thinking about today's devices, let's consider the current smartwatches and fitbands: the former relies on a smartphone or tablet in order to provide Internet-based services, and the latter uses no Internet connection at all, but just sends data to the smartphone to be processed. Along the line of favouring locality, there is a second advantage which relates to privacy issues: there is no reason to send personal information away, if the system does not need data from distant points nor requires more computational power of the amount available locally. Privacy issues gain great attention recently, especially after the leaks of classified information started in 2010 on Wikileaks and continued with the more recent leaks by Edward Snowden. The content of such documents raised greater attention to privacy issues from general public.

A possible path which would help in both those directions (reduce wireless medium usage and keep data as local as possible) is the usage of local, possibly peer-to-peer interactions. This way of organising communication is already exploited by existing applications.

One notable example is Firechat<sup>4</sup>, which got particularly spotlighted during the “Umbrella revolution”, namely the sequence of protests that took place in Hong Kong in 2014. Similarly to what was done during Arab springs in 2011 [HDF<sup>+</sup>11], protesters relied on Internet services to organise and coordinate themselves. The Chinese government policy on internet is not exactly a bright example of openness and neutrality [Mac08], and services such as Facebook and Twitter, widely exploited during Arab springs, were already effectively blocked in the land. In short time, other social network were closed (such as Instagram), in order to cut protesters' communication means. At that point, protesters had to find a communication system free of centralisation in order to prevent targeted Internet filters, and Firechat was the answer. Firechat is a messaging application for mobile phones that, when the smartphone has access to the remote Firechat cloud, works as other more famous alternatives, such as WhatsApp<sup>5</sup> and Telegram<sup>6</sup> do. When no access to Firechat servers is available, then the software tries to reach the destination by spreading the message hop-by-hop, building a de-facto mesh network.

Another interesting experiment is Serval Mesh [GSP11]. Serval Mesh accomplishes similar tasks, but it also supports calls and file transfers besides messaging. Its main goal is to provide a networking among users who are in an area where there is no Internet access at all, for instance because of a disaster event. It relies on WiFi to create a ad-hoc peer-to-peer network among devices. Due to this lower level aspect, Serval Mesh requires privileged access to the hardware and some higher skill than Firechat, which is easier to setup but requires the availability of a Internet connection. In [AGS11], Serval Mesh is used to build an alternative, purely peer-to-peer telephony network.

Despite the existence of such mesh-oriented applications, however, a general and widespread approach for easily design and program the devices that compose our pervasive continuum is still under investigation [RCKZ13]. The most consistent contribution of this thesis is devoted to the research of general, well engineered approaches to build such systems.

---

<sup>4</sup><https://opengarden.com/firechat>

<sup>5</sup><http://www.whatsapp.com/>

<sup>6</sup><https://telegram.org/>



# 2

## Self-organisation

In Chapter 1 we took a look to the world of pervasive devices, also describing their communication means and hypothesising the near future development. In this chapter, we focus on the software, and in particular on the challenges of engineering the development of software that will run on an ensemble of pervasive devices. We first discuss the issue of coordination and self-organisation: how do we make all those possibly devices collaborate together in order to achieve a global goal, without a centralised decision-maker? Which software platform may we devise to ease this operation? We will see that similar problems have already been successfully solved in nature: the mechanisms underlying such natural behaviours, can, if properly mimicked, help to realise solutions in software systems. We will then run through the existing literature on the issue, analysing the existing platforms supporting pervasive computing, and the tools that can be used to test and debug applications prior to deployment. Finally, we will discuss the shortcomings of the existing technology, and pave the way for the contribution of this PhD thesis.

### 2.1 Software ecosystems

Regardless the name that we want to use, being it “Internet of Things” rather than “Smart cities” or “Pervasive computing”, in all cases we are talking about a system that is expected to host many computations that feature, according to [ZOA<sup>+</sup>14]:

**Situatedness** — Pervasive services are typically time- and space-dependent, and feature physically- or socially-situated activities. Components of pervasive systems should be able to interact with the surrounding physical and social world by adapting their behaviour accordingly.

**Autonomy and self-adaptivity** — While individual components should be autonomous in the face of the inherent dynamics of their operational environment [GPO13], pervasive systems should also feature *system-level autonomy* to deal globally with the unpredictability of the environment, providing properties such as self-adaptation, self-management, and self-organization [MMTZ06].

**Prosumption and diversity** — Infrastructures for pervasive systems must promote open models of component integration, to be able to take advantage of the injection of new services and components [ZO04]. This is particularly true in the context of *socio-technical systems* [?], where human users and software agents act as *prosumers* – both consumers and producers – of devices, data, and services.

**Eternity** — As well as short-term adaptation, a pervasive systems infrastructure should allow for the long-term evolution of organizations, components, and patterns of usage, in order to accommodate technological advances as well as the mutable needs of users without requiring extensive re-engineering effort [Jaz05]. In fact, pervasive systems are better conceived as *eternal* systems, engineered for continuous, unlimited service, upgrading, and maintenance over time.

At the same time, “traditional” networks are also increasing in scale and importance for enterprises both large and small. In an increasingly information-dependent and interconnected world, the rising cost of managing and maintaining such systems is driving a search for solutions that can increase the autonomy of computing systems, enabling them to act more as collective services than individual machines [EAWS12, HQL<sup>+</sup>11], and continuing working over time. In both of these cases, and a number of other areas facing similar challenges (e.g., large-scale sensor networks, multi-UAV control), there is a growing recognition that new paradigms are needed to face the challenge of coordination, namely, engineering the space of interactions [Weg97]. The goal of such engineering paradigms is to find reliable processes that lead to self-organising aggregates that act more as a collective service than as individual machines [CLMZ03, BMP<sup>+</sup>11, LPF12, SRM<sup>+</sup>13].

## 2.2 Nature inspiration

Natural systems are good in dealing with the *complexity* of coordinating large-scale software ecosystems [Sha04, KR08, LT06], since they natively feature key properties such as autonomy, openness, fault tolerance, situated behaviour, self-adaptation and robustness. Many specific models have been proposed in literature, taking inspiration from natural ecosystems, social insects, biochemistry, chemistry and also physics [ZV11].

### 2.2.1 Physical-inspiration

To the category of physical inspired systems belong all those works that take inspiration on the way physical particles move and self-organise according to gravitational and electromagnetic forces. For instance, in [MZ06a], computational “force fields” – generated either by coordinated components or by the coordination middleware – propagate across a spatial environment, leading to distributed data structures that affect the actions and motions of the agents in that environment. A similar approach is proposed in Co-fields [MZL03], exploiting composite computational fields

to coordinate the motion of users and robots in an environment. Physical inspiration also has influenced self-assembly, for instance in [GJM12] the adoption of virtual force fields is suggested in order to control the material’s shape.

### 2.2.2 (Bio)chemical-inspiration

Chemical reactions can be seen as an ensemble of myriads of simple laws that generate and regulate the evolution of extraordinarily complex molecular structures. Those rules coordinate in some way the behaviours of a huge number of components, up to the point where they reductionistically drive the evolution of complex assemblies such as biological organisms and meteorological systems.

Gamma [BLM90] was the first and the most prominent example of a chemically-inspired model. In Gamma, coordination is conceived as the evolution of a space governed by chemical-like rules, globally working as a rewriting system [BFLM01]. In the CHAM (chemical abstract machine) model [Ber92], states are interpreted as chemical solutions, where floating molecules (representing coordinated entities) interact according to some reaction rules, and where *membranes* constrain the execution of reactions.

In chemical tuple spaces [VCNO10] data, devices, and software agents are uniformly represented in the form of chemical reactants, and system behaviour is expressed by means of full-fledged chemical-like laws that are both time-dependent and stochastic rather than in form of rewriting rules. *Biochemical tuple spaces* [VCMZ11] enhance chemical tuple spaces by shaping coordination through distribution and topology. Biochemical coordination models appear very flexible in enabling the spatial formation of both localised and distributed activity patterns, and have been exploited in many special-purpose pervasive frameworks, including crowd mobility management [WBYI08] and participatory sensing [LMG<sup>+</sup>09]. The amorphous computing model can be considered an example of biochemical coordination model [AAC<sup>00</sup>].

Beside having been conceived as general computational model, (bio)chemically-inspired computing can be an effective starting point to realize schemes of dynamic service composition [FSS10] or knowledge aggregation [MO13]. Network protocols for data distribution and aggregation according to chemical models have also been explored, as in the Fraglets approach [MYT07, MMTL13]. Several proposals exist to support service composition based on chemical models [BP09], including proposals specifically conceived for adaptive pervasive services [HAES09], or to support the adaptive organization of knowledge [MO13].

### 2.2.3 Stigmergy

In the computer science literature, the term “stigmergy” refers to a set of nature-inspired coordination mechanisms mediated by the environment [TB99, Bon99]. Agents deposit data into a distributed, shared, environment so as to collectively (yet implicitly) build distributed data structures that can help them navigate in such environments. The notion of *stigmergy* was introduced in [Gra59] as the fundamental coordination mechanism in termite societies.

Nowadays, the most widely-studied example of stigmergic coordination in insect societies is probably that of ant colonies [DS04]. The basic mechanism is based on *pheromones* that are released in the environment by the ants that find food on their way back to the nest, thus building pheromone trails towards food that other ants are then stimulated to follow. The pheromones act as environment markers for specific social activities, driving both the individual and the social behaviour of ants. For instance, digital pheromones [PBS02, Par06] have been fruitfully exploited as the basic mechanism for coordinating the movements of robot swarms and modular robots [MZ05], for helping people find directions in an unknown environment [MZ07], and for efficiently cluster information in networks [CVG09, PVM<sup>+</sup>10]. In the area of networking, stigmergy has been exploited to realize effective routing mechanisms in dynamic networks [BDT99, DDF<sup>+</sup>06].

## 2.3 Spatial patterns

Vedere se c'è qualcosa di utile in NACO [FMDMSM<sup>+</sup>12] o in [GVO07], o in [ZV11] o in [BCD<sup>+</sup>06]

### 2.3.1 Gossip

### 2.3.2 Gradient

A simple though paramount data structure that can be built upon this paradigm and that is a building block of many of the more advanced patterns is the spatial gradient, which assigns to each node a value  $\Gamma$  depending on its position in time and space and on its context [MZ09, BBVT08, VCMZ11]. This structure originates in one or more devices called sources. In every source device,  $\Gamma = 0$ . In every other device, let  $N$  be the set of devices connected with it, and let  $n$  be the  $n$ -th neighbouring device. For this device,  $\Gamma = [\min(f(n)) | n \in N]$ : namely, the value of the device is the minimum of a function which operates on the neighbours. For instance, if  $\forall n: f(n) = \Gamma_n + 1$ , where  $\Gamma_n$  is the value of the gradient in  $n$ , then the local value will reflect the minimum hop count towards the nearest source. If  $f(n) = \Gamma_n + d(n)$  where  $d(n)$  measures the actual distance from the device towards  $n$ , then the value of the gradient will approximate the distance from the nearest source. Finally, along with the local value, a gradient can carry more information, e.g. some strings, the position of  $n$ , or numeric values.

[BBVT08]

### 2.3.3 Gradcast

### 2.3.4 Voronoi partition

## 2.4 Tuple-based coordination

In a tuple-based coordination model [RCD01], software agents synchronize, co-operate, and compete based on *tuples*, which are simple data structures representing information chunks. Tuples are made available in *tuple spaces*, which are shared information spaces working as the *coordination medium*. Coordination occurs by accessing, consuming, and producing tuples in an *associative* way, relying on the actual content of tuples and not on any form of naming, addressing, or indexing. An interesting survey of the technologies and platforms [MT03] analyses the suitability of tuple-based coordination systems in supporting openness, unpredictable changes in distributed environments, and several aspects related to adaptiveness: requirements that are of primary importance in pervasive systems.

Several implementations from both academia and industry exist.

**Anthill** [BMM02] is a framework meant to support design and development of adaptive peer-to-peer applications, in which each node is provided with a local tuple space, agents can travel the network and interact indirectly reading, writing and retrieving tuples.

**Biochemical tuple spaces** [VC09]

**GigaSpaces**<sup>1</sup>

**JavaSpaces** [FHA99, NZ01]

**Lime** [MPR06]

**Linda** [Gel85] is the common ancestor of every tuple based coordination framework.

**Molecules of knowledge** [MO13]

**SwarmLinda** [TM03]

**T Spaces** [WMLF98]

**TOTA** [MZ09, MZ06b] is a tuple-based middleware explicitly conceived to support field-based coordination for adaptive context-aware and spatially-aware activities in pervasive computing scenarios. In TOTA, each tuple also carries with it a diffusion rule and a maintenance rule in addition to a content. Tota inspired other works, such as the evolving tuples model [SJ07], which adds to the tuples a form of context awareness, in form of possible evolution and adaptation to environmental changes.

---

<sup>1</sup><http://www.gigaspaces.com>

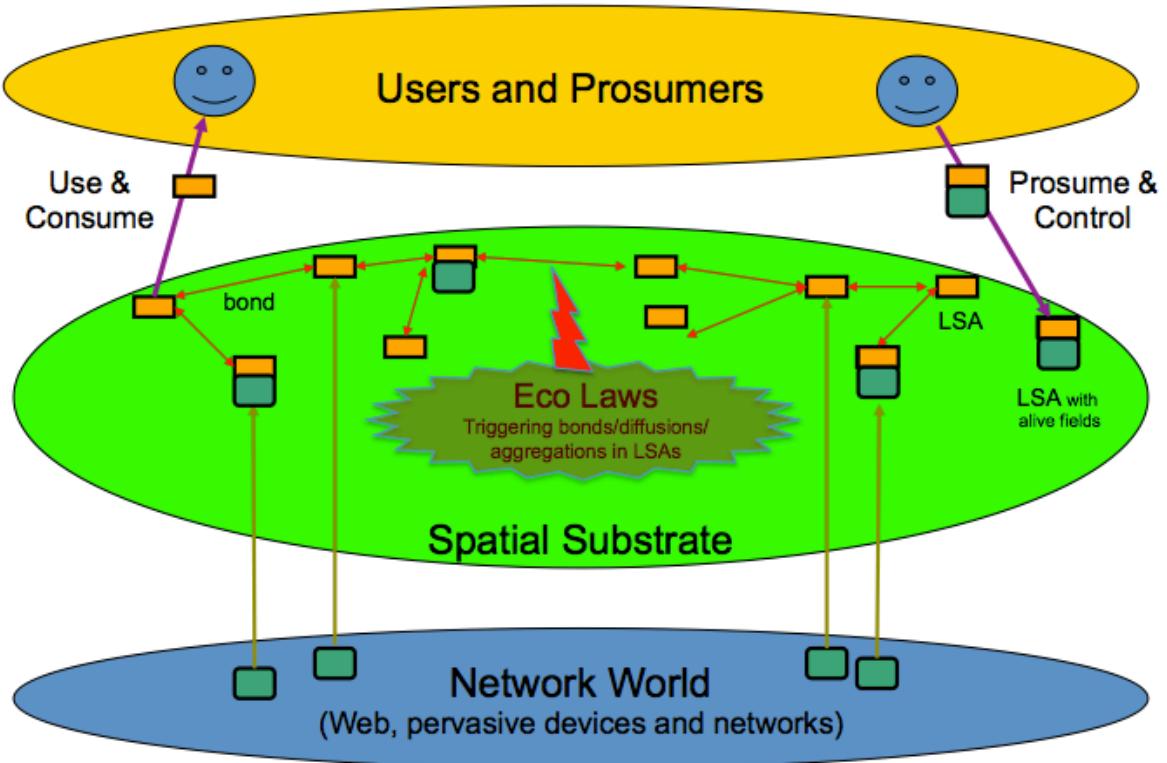


Figure 2.1: The SAPERE Reference Architecture.

**TuCSoN** [OZ99]

SELFMAN? [RHR<sup>+</sup>08]

MARS?

Plastic? [ABI09]

CArtAgO? [RPV11]

### 2.4.1 SAPERE

SAPERE<sup>2</sup> [ZCF<sup>+</sup>11] is a nature-inspired coordination model and framework to support the design and development of composite pervasive service systems. Its reference architecture and coordination model synthesize from existing nature-inspired approaches, such as [VC09, VCO09, VCMZ11, PVM<sup>+</sup>10], and is based on an assumption of spatial, local interactions (to be realized via a network of distributed tuple spaces), which is in line with all nature-inspired

<sup>2</sup>The model was developed as part of an EU-funded research project. See <http://www.sapere-project.eu/>.

approaches. Its coordination laws make it possible to express and deploy general nature-inspired distributed algorithms and coordination patterns. SAPERE abstracts a pervasive environment as a non-layered *spatial substrate* deployed upon a dense network of connected heterogeneous ICT devices (Figure 2.1). SAPERE acts as a shared coordination medium embodying the basic laws of coordination. Its core components are [VPMS12]:

**LSA** Because of the need to tolerate diversity, a cornerstone of pervasive ecosystems is that a uniform representation is required for the various software agents living within them (whether they run on smartphones, sensors, actuators, displays, or any other computational device). Such a representation needs to expose any information about the agent (state, interface, goal, knowledge) that is pertinent for the ecosystem as a whole or for any subpart of it. In SAPERE this description is called “Live Semantic Annotation” for it should continuously represent the state of its associated component (live), and it should be implicitly or explicitly connected to the context in which such information is produced, interpreted and manipulated (semantic)—possibly relying on standard technologies and techniques of the Semantic Web, like RDF [MM04].

**LSA-space** To handle situatedness, the behaviour of each agent strictly depends on the local context in which it runs, that is, on the state of other agents living in the same locality (intended as network neighbourhood). As such, the LSAs of each agent are reified in a distributed space (called an “LSA-space”) acting as the fabric of the ecosystem, where “context” is simply defined and represented as the set of LSAs stored in a given locality.

**LSA bonding** Additionally, and in order to make any agent act in a meaningful way with respect to the context in which it is situated, special mechanisms are needed to provide a fine-tuned control of what to each agent is visible/modifiable and what is not. SAPERE tackles this issue by allowing an LSA to include bonds (i.e., references) to other LSAs in the same context. It is only via a bond that an agent can inspect the state/interface of another agent and act accordingly, while modifications are allowed only to the LSAs an agent injected itself.

**Eco-laws** Because of adaptivity, while agents enact their individual behaviour by observing their context and updating their LSAs, global behaviour (i.e., global system coordination) is enacted by self-organising manipulation rules of the LSA-space, called eco-laws. They can execute deletion/update/movement/re-bonding actions applied to a small set of LSAs in the same locality. SAPERE structures such eco-laws as chemical-resembling reactions over LSAs—similarly to other approaches like [BP09, VC09, VCMZ11].

## 2.5 Aggregate programming

Aggregate programming is founded on the observation that in many cases the users of a system are much less concerned with individual devices than with the services provided by the collection of

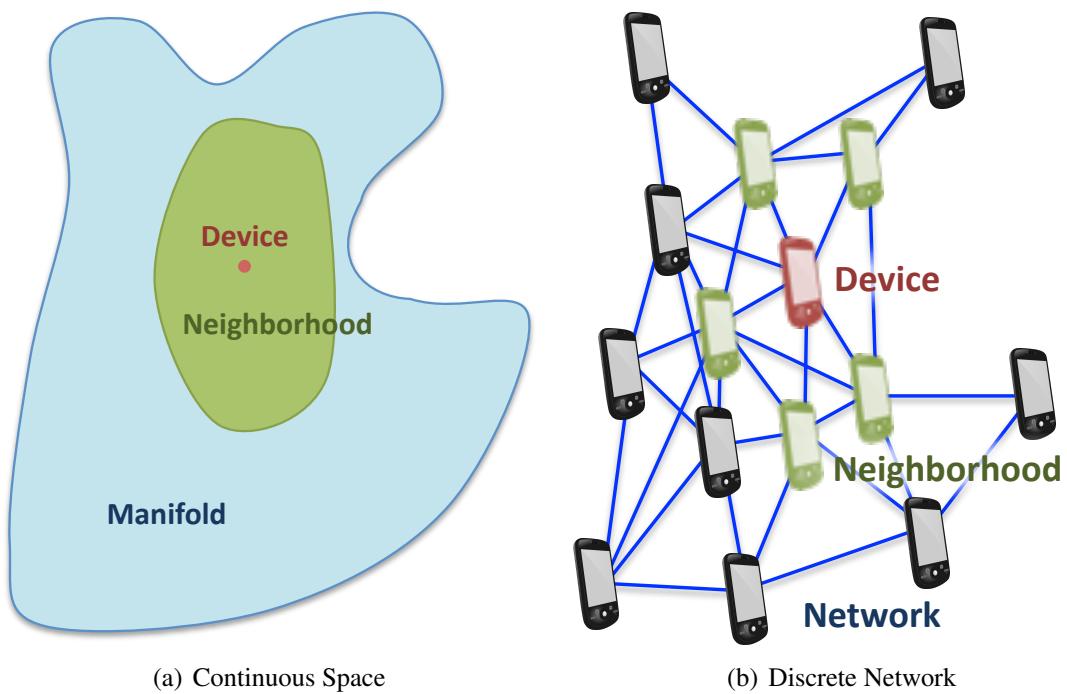


Figure 2.2: Computational field models originate from approximation of continuous space (a) with discrete networks of devices (b).

devices as a whole. Typical device-centric programming languages, however, force a programmer to focus on individual devices and their interactions. As a consequence, several different aspects of a distributed system typically end up entangled together: effectiveness and reliability of communications, coordination in face of changes and failures, and composition of behaviours across different devices and regions. This makes it very difficult to effectively design, debug, maintain, and compose complex distributed applications.

Aggregate programming generally attempts to address this problem by providing composable abstractions separating these aspects:

1. device-to-device communication is typically made entirely implicit, with higher-level abstractions for controlling efficiency/robustness trade-offs;
2. distributed coordination methods are encapsulated as aggregate-level operations (e.g., measuring distance from a region, spreading a value by gossip, sampling a collection of sensors at a certain resolution in space and time); and
3. the overall system is specified by composing aggregate-level operations, and this specification is then transformed into a complete distributed implementation by a suitable mapping.

In short, the key idea behind aggregate programming is to provide languages and APIs that allow a distributed collection of devices to be programmed in terms of their collective behaviours, keeping outside the sight of the software designer details on how such coordination is actually implemented.

A large number of very diverse aggregate programming approaches have been proposed, including abstract graph processing (e.g., [GGG05]), declarative logic (e.g., [ARGL<sup>+</sup>07]), map-reduce (e.g., [DG08]), streaming databases (e.g., [MFHH05]), and knowledge-based ensembles (e.g., [NLPT14])—for a detailed review, see [BDU<sup>+</sup>13].

Despite the fact that all of them are based on viewing the collection of devices as an approximation of continuous space (as depicted in Figure 2.2), most aggregate programming approaches, however, have been too specialized for particular assumptions or applications to be able to address the complex challenges of these emerging environments.

### 2.5.1 Proto

Proto language and middleware [BB06] exploits field-based coordination to orchestrate the activities of sensor-actuator networks

Proto is based on the notion of a *computational field*—a map from devices comprising the system to (possibly structured) values, which is treated as unifying first-class abstraction to model system evolution and environment dynamics.

The programming constructs are combined together to form programs, whose semantics is defined in terms of a sequence of synchronous rounds of evaluation by a discrete network

of devices called “computational rounds”. In practice, however, there is no requirement for synchrony, and each device can evaluate its own computational rounds independently.

Despite its qualities, Proto still lacks many features expected in a modern programming language and has an implementation encumbered by a number of obsolete considerations that make it difficult to maintain and extend.

[BB06]

## 2.5.2 Field Calculus

Field Calculus is an attempt to find a unifying model for programming *computational fields* as a generalization of a wide range of existing approaches (Proto in particular, but also [MZ09, NW04, VCMZ11, MZ06b, Nag01, Yam07, NW04]). Formalized as the computational field calculus [VDB13], this universal language appears to provide a theoretical foundation on which effective general aggregate programming platforms can be built.

Critically, although originally derived from continuous-space concepts, the calculus does not depend on them and is applicable to any network. Field calculus is expressive enough to be a universal computing model [BVD14] but terse enough to enable a provable mapping from aggregate specifications to equivalent local implementations.

Field calculus [VDB13] hence provides a key theoretical and methodological foundation for aggregate programming. Its aim is to provide a universal model that is suitable for mathematical proofs of general properties about aggregate programming and the aggregate/local relationship, just as  $\lambda$ -calculus [Chu32] provides for functional programming,  $\pi$ -calculus for parallel programming [Mil99], or Featherweight Java [IPW01] for Java-like object-oriented programming.

In the field calculus, everything is a field: computational fields are used to model every aspect of distributed computation, including input from sensors, network structure, environment interactions, distributed computations (e.g. progressive aggregation and spreading processes), and output for actuators. In particular, field calculus is constructed using five basic constructs:

1. function definition and evaluation;
2. “built-in” operations for stateless local computation, e.g., addition, multiplication, reading a sensor;
3. a time-evolution construct which allows for stateful computation;
4. a neighbour-value construct that creates a field of values from a device’s neighbours;
5. a restriction operator to select which computations to perform in various regions of space and time.

As well as in Proto, in Field Calculus the semantics of programs created by combining such constructs is defined in terms of a sequence of (possibly synchronous) “computational rounds”.

The minimal syntax of field calculus has allowed its semantics, including proper coherence of device interactions, to be proven correct and consistent [VDB13]. Additionally, despite its definition in terms of discrete semantics, field calculus is also space-time universal [BVD14], meaning that it can approximate any field computation, either discrete or continuous, with arbitrary precision given a dense enough network of devices.

This, then, is the key contribution of field calculus: any coordination method with a coherent aggregate-level interpretation is guaranteed to be expressible in field calculus. Such a method can then be abstracted into a new aggregate-level operation, which can be composed with any other aggregate operation using the rules of built-in functions over fields. Moreover, it can have its space-time extent modulated and controlled by restriction, all while guaranteed that the relationship between global specification and local implementation will always be maintained.

Such a core calculus, like any other, is more a theoretical framework than a practical programming language. In practice, in fact, effective aggregate programming for real-world distributed applications is likely to require a layered approach such as the one depicted in Figure 2.3, of which field calculus represents the first block. Some of the prior approaches on which field calculus is based provide very similar semantics (most notably Proto), but they all suffer from some combination of design and implementation problems that render them impractical for widespread adoption. Besides Proto, which is probably the closest to a practical programming environment and whose problems have already been described in Section 2.5.1, others, such as [VPB12], have only minimal implementation.

## Alignment

## 2.6 Engineering and tools

Despite the efforts of addressing specific aspects of self-adaptability and openness in pervasive contexts [MMTZ06, BCD<sup>+</sup>06, KR08], yet a comprehensive engineering framework is missing. The hardest challenge is due to the fact that most of the complexity of such systems does not come from the individual behaviours, but rather from interaction [Weg97, GSW06, ORV06].

Multiple methodologies have been proposed to tackle complex pervasive systems [MCOV11, BW08, BMV09a, MN10, ZCF<sup>+</sup>11], and they always include *simulation* as a key step to realise what-if analysis prior to actual development, to both assess the general validity of the designed mechanisms and to fine tune system parameters. There are many kinds of simulation tools available: they either provide programming/specification languages devoted to ease construction of the simulation process, especially targeting computing and social simulation (e.g. as in the case of multi-agent based simulation [BMV09a, SGJ07, BMV09b, NHC07, Skl07]), or they stick to quite foundational computing languages to better tackle performance, mostly used in biology-oriented applications [Pri95, Mur89, UP05, EMRU07]. None of the existing tools, however, aims at bridging the gap between these approaches, trying to extend the basic computing model of chemical reactions – still retaining its high performance – toward ease applicability to complex situated computational systems.

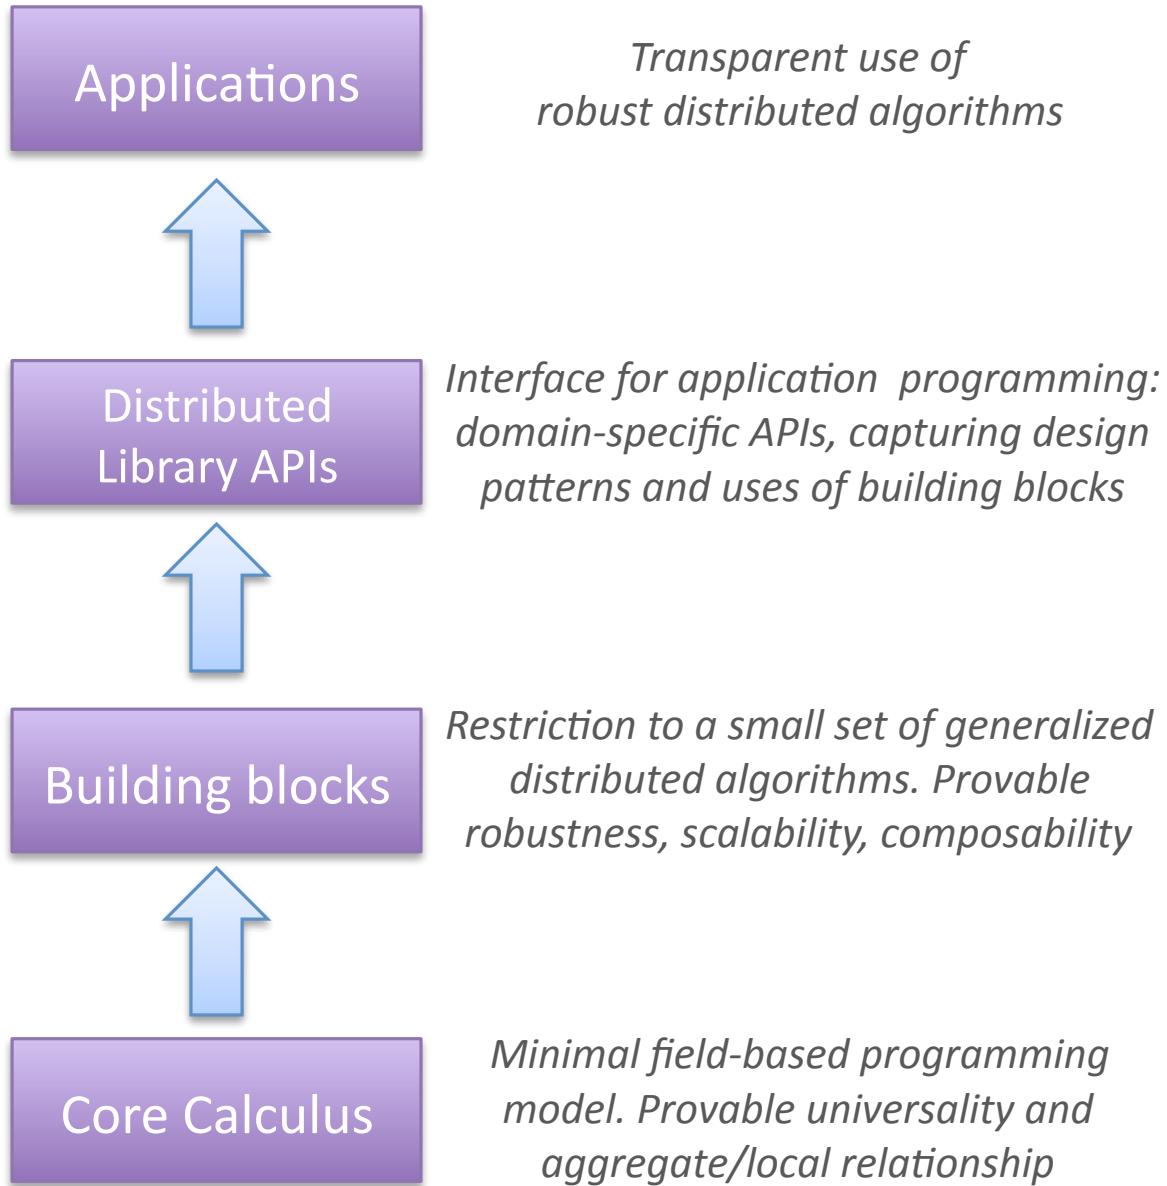


Figure 2.3: Layered approach for development of spatially-distributed systems via aggregate programming.

### 2.6.1 General purpose frameworks

MASON [LCRP<sup>+</sup>05], Repast [NHCV07], NetLogo [Skl07] and Swarm

### 2.6.2 Specific simulators

network simulators (the one, ns2) biological simulators (ask Sara)

Specific per-use simulators, gener



# 3

## **Shortcomings of coordination infrastructures**

**3.1 Engineering emergence**

**3.2 Local to global**



# 4

## **Shortcomings of existing tools**

**4.1 Specific tools: expressiveness**

**4.2 General purpose tools: performance**



## **Part II**

# **Contribution**



# 5

## An integrated toolchain for pervasive ecosystems

We discussed in Section 2.6 the importance of simulation in the engineering process of designing pervasive ecosystems. In this chapter, we will discuss one of my contribution to this research branch, namely the development of an integrated tool named ALCHEMIST, which has been inspired by chemical oriented simulators. We first talk about chemical inspiration, and how existing chemical-oriented stochastic simulation algorithms (SSAs) can be extended towards higher flexibility, up to the point that they become suitable to simulate potentially complex environments, still retaining their high performance. To do so, our first effort will be to take an efficient SSA implementation and convert it to a general-purpose discrete event simulator (DES). We then devise a meta-meta-model, complex and flexible enough to support the definition of middleware-specific meta-models. Finally, some meta-models are presented.

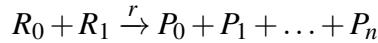
### 5.1 Chemical-inspired engine

In the following, we present a discrete event simulation engine derived from the popular and successful Gillespie's stochastic simulation algorithm [Gil77], and in particular from its notable and more efficient extensions developed by Gibson-Bruck [GB00] and by Slepoy [STP08]. Gillespie's SSA is a discrete and stochastic method that intrinsically owns the event-driven properties [BBL<sup>+</sup>11]: introduced to model chemical systems, nowadays it represents the basis of many simulation platforms, in particular those developed for the investigation of biochemical systems [PRSS01, Kie02, CG09, VB08, MV10, GPOS13, HSG<sup>+</sup>06]. With significative extensions, it was recently used to model artificial systems grounding on computational models inspired to chemical natural systems and ecology [MVFM<sup>+</sup>13]. Moreover, its optimised extensions [GB00, STP08] are very efficient, thus allowing really fast simulation runs.

#### 5.1.1 Gillespie's SSA as an event-driven algorithm

First of all, we summarise the idea that the Gillespie's SSA is grounded on: a chemical system is modelled as a single space filled with molecules that may interact through a number of reactions

describing how they combine. The instantaneous speed of a reaction is called propensity and depends on the kinetic rate of the reaction and on the number of molecules of all the reagents involved. For a reaction  $i$  of the form:



the propensity  $a_i$  is defined as:

$$a_i = r \cdot [R_0] \cdot [R_1]$$

where  $[M]$  is the number of molecules of species  $M$ . Given that, the algorithm relies on the idea that the system can be simulated by effectively executing the reactions one by one and changing the system status accordingly. Every algorithm follows four main steps:

1. select the next reaction  $\mu$  to be executed;
2. calculate the time of occurrence of  $\mu$  according to an exponential time distribution and make it the current simulation time;
3. change the environment status in order to reflect this execution;
4. update the propensities of the reactions.

Such algorithm is event-driven in that it executes only one reaction/event at a time: it changes the state of the system and consequently the event list *i.e.*.. which other reactions/events can be executed from there.

### 5.1.2 Gillespie's optimised versions

This algorithm has been improved in various works in literature, particularly notable is the work in [GB00] and [STP08]. Both works optimise the base algorithm in two phases: the selection of the next reaction to execute and the update of the reaction pool – pending event list – once an event has been executed. For the latter, they both rely on the concept of “dependency graph”. A dependency graph (DG) is a statically created directed graph in which nodes are all the reactions in the simulated system, and arcs connect a reaction  $r$  to all those that depend on it, namely, those whose triggering time should be updated as  $r$  is executed. For instance, if  $r$  changes the concentration of some molecule  $m$ , all those reactions that use  $m$  are to be properly re-scheduled as soon as  $r$  is fired. Even though this optimisation does not affect the execution time in the worst case, it offers great benefits in the average case, since most of the reactions are not interdependent.

The selection of the next reaction to execute is where those improved algorithms differ most. In Slepoy's work, the author divides reactions in groups based on their propensity: if  $p_{min}$  is the lowest propensity, the first group contains those reactions whose propensity ranges from  $p_{min}$  and  $2p_{min}$ , the second those between  $2p_{min}$  and  $4p_{min}$ , and so on. Give those groups, an algorithm called “Composition-Rejection” (CR) is applied. In the composition phase, a group is

randomly selected in logarithmic time; in the rejection phase a reaction is chosen within a group by throwing two random numbers: the first one, between 0 and the number of reactions in the group, identifies a reaction; the second one, between 0 and the maximum propensity allowed for the group, is used to decide whether or not to actually execute the selected reaction. In case the reaction is rejected, the rejection procedure is re-applied. Given the way groups are built, there is at least a 0.5 probability that the rejection procedure concludes at each attempt. On average, the CR algorithm requires five random numbers to be thrown. This algorithm requires logarithmic time to select the next reaction with respect to the total number of groups, but the author argues which, since in most biological scenarios this number is constant, then the algorithm runs effectively in constant time. In Gibson-Bruck, the same selection operation is made by computing a putative execution time for each reaction, and then using a binary heap data structure to sort them. This way, the next reaction to execute is always the root node of the tree, and the selection is performed in constant time. However, once the reaction has been executed, all the dependent reactions must be updated and re-sorted in the tree. In the worst case, this takes logarithmic time with respect to the number of reactions.

### 5.1.3 From SSA to full fledged DES

Both the Gisbon-Bruck's and Slepoy's schedulers are granted to correctly simulate a Poisson process. However, in order to build a full-fledged DES engine, we must offer the possibility to schedule also non-Markovian events. Imagine, for instance, that we want to simulate an agent that does an action every fixed time interval (e.g. a man walking). This, clearly, is not a memoryless process. We argue that Gibson-Bruck offers a more suitable base for a general purpose DES: in fact, its next reaction choosing mechanism is orthogonal to the way the putative times are computed. This intrinsic feature allows to neatly separate the generation of times (namely, the time distribution of each event) and the actual scheduling (choice of which event should run next). We chose the Next Reaction Method for ALCHEMIST, and, consequently, the main simulation algorithm follows the basic steps listed in Algorithm 1.

A second feature that we need in order to shift the paradigm from pure chemistry towards higher expressiveness is the possibility to simulate multiple, separate, interacting and possible mobile entities. This requirement can be partly addressed by the notion of intercommunicating compartments [CG09, VB08, MV10, GPOS13], in a way that allows to also model systems characterised by a set of connected volumes and not only a unique chemical solution. The hardest challenge in simulating multiple compartments with an efficient SSA is in improving the dependency graph: reactions that could be interdependent but happen on separated compartments should not be marked as interconnected within the dependency graph. Mobility makes everything even harder, since it may lead to the creation and disruption of communication channels between compartments, and consequently to simulation-time changes in the structure of such dependency graph. Summarising, supporting dynamic environments with multiple mobile compartments require the dependency graph to become a dynamic data structure, which cannot be pre-computed at the simulation initialisation and kept static.

---

**Algorithm 1** Simulation flow in ALCHEMIST

---

```

1: cur_time = 0
2: cur_step = 0
3: for each node n in environment do
4:   for each reaction nr in n do
5:     generate a new putative time for nr
6:     insert nr in DIPQ
7:     generate dependencies for nr
8:   while cur_time < max_time and cur_step < max_step do
9:     r = the next reaction to execute
10:    if r's conditions are verified then
11:      execute all the actions of r
12:      for each reaction rd which depends on r do
13:        update the putative execution time
14:      generate a new putative time for r

```

---

### Dynamic Dependency Graph

There are multiple ways to conceive such dynamic data structure. Ideally, it would be possible to just drop the optimisation or reuse the classic definition, in which case the triggering of a reaction in whichever compartment would cause the recalculation of the status of each potentially dependent event in every compartment. This approach would lead to a massive performance impact, since the dependency computation is the most expensive operation. As evidence, in [STP08] some charts show the difference between Gibson-Bruck and Slepoy et al. algorithms with and without a dependency graph: the result is very strongly in favour of the former.

A possibility for efficiently adapting a dependency graph to a network of compartments could be to define the input and output contexts for each reaction, namely the places where reactions respectively “read” their reactants and “write” their products. Multiple contexts could be defined, we propose to adopt three levels: local, neighborhood and global. In a purely chemical simulation, all the reactions have a local input context and may have either neighborhood or local output context, depending on whether or not they send molecules towards other compartments.

Let  $dep : R^2 \rightarrow \text{boolean}$  be the function that is used to build the static dependency graph in every SSA: given two reactions  $r_1$  and  $r_2$ ,  $dep(r_1, r_2)$  returns true if the propensity of  $r_2$  may be influenced by the execution of  $r_1$ . In a multi-compartment scenario,  $r_1$  influences  $r_2$  (there is an oriented edge in the dependency graph connecting  $r_1$  to  $r_2$ ) iff  $dep(r_1, r_2)$  is true and:

- $r_1$  and  $r_2$  are on the same node OR
- $r_1$ 's output context is global OR

- $r_2$ 's input context is global OR
- $r_1$ 's output context is neighborhood and  $r_2$ 's node is in  $r_1$ 's node neighbourhood OR
- $r_2$ 's input context is neighborhood and  $r_1$ 's node is in  $r_2$ 's node neighbourhood OR
- $r_1$ 's output context and  $r_2$ 's input context are both neighborhood and the neighbourhoods of their nodes have at least one common node.

The filters listed above greatly compact the number of edges of a dependency graph in most scenarios, with great benefits on the engine performance. On top of this finer-grain locality concept, if the model supports compartment mobility, the dependency graph must support the dynamic addition and removal of reactions.

Adding a new reaction implies to verify its dependencies against every reaction of the system. In case there is a dependency, it must be added to the dependency graph. Removing a reaction  $r$  requires to delete all dependencies in which  $r$  is involved both as influencing and influenced. Moreover, in case of a change of system topology, a dependency check among reactions belonging to nodes with modified neighbourhood is needed. It can be performed by scanning them, calculating the dependencies with the reactions belonging to new neighbours and deleting those with nodes which are no longer in the neighbourhood.

### **Dynamic Indexed Priority Queue**

An issue that arises with addition and removal of nodes from the simulation is the possible unbalancing of the scheduling queue, that in the original work is realised as a binary tree of reactions, whose main property is that each node stores a reaction whose putative time of occurrence is lower than each of its sons.

Our idea is, for each node, to keep track of the number of descendant per branch, having in such way the possibility to keep the tree balanced when adding nodes. In figure 5.1 we show how the same IPQ drawn in [GB00] would appear with our extension. Given this data structure, the procedures to add and remove a new node  $n$  are described respectively in Algorithm 2 and Algorithm 3, in which the procedure `UPDATE_AUX(n)` is the same described in [GB00].

Using the two procedures described above, the topology of the whole tree is constrained to remain balanced despite the dynamic addition and removal of reactions.

## **5.2 Meta-meta model**

The complexity of the systems we want to design is achieved by the following set of common key properties:

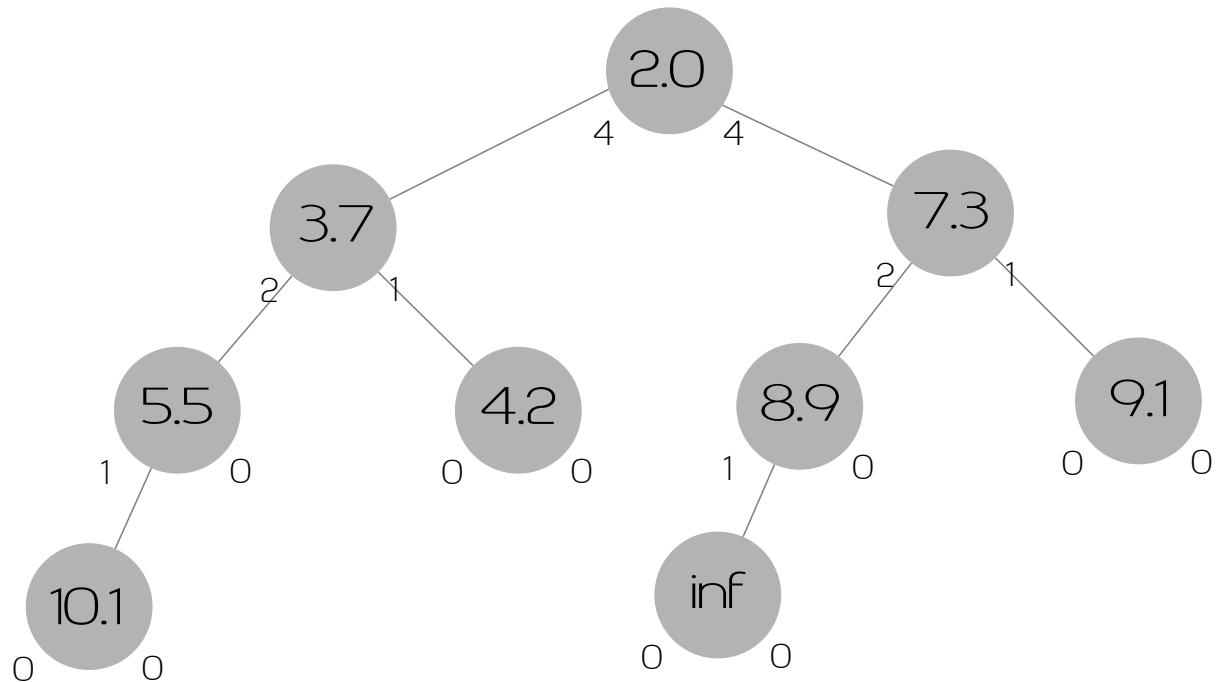


Figure 5.1: Indexed Priority Queue extended with descendant count per branch

---

**Algorithm 2** Procedure to add a new node n

---

```

1: if root does not exist then
2:   n is the new root
3: else
4:   c ← root
5:   while c has two descendants do
6:     if c.right < c.left then
7:       dir ← right
8:     else
9:       dir ← left
10:    add 1 to count of dir descendants
11:    c ← c.dir
12:    if c has not the left child then
13:      n becomes left child of c
14:      set count of left nodes of c to 1
15:    else
16:      n becomes right child of c
17:      set count of right nodes of c to 1
18: UPDATE_AUX (n)

```

---

---

**Algorithm 3** Procedure to remove a node n

---

```

1: c ← root
2: while c is not a leaf do
3:   if c.left > c.right then
4:     dir ← left
5:   else
6:     dir ← right
7:   subtract 1 to count of dir descendants
8:   c ← c.dir
9: if c ≠ n then
10:   swap n and c
11:   remove n
12:   UPDATE_AUX (c)
13: else
14:   remove n

```

---

- situatedness – they deal with spatially- and possibly socially-situated activities of entities, and should therefore be able to interact with a limited portion of the surrounding world and contextualise their behaviour accordingly;
- adaptivity – they should inherently exhibit properties of autonomous adaptation and management to survive contingencies without external intervention, global supervision, or both;
- self-organisation – spatial and temporal patterns of behaviour should emerge out of local interactions and without a central authority that imposes pre-defined plans.

Among the many natural metaphors one can use as inspiration for modelling and developing artificial systems with the above properties [ZV11], we consider chemistry following a series of work in the field of pervasive computing [VCMZ11, VZ10, ZCF<sup>+</sup>11]. We argue that there are three main issues to be resolved in order to build a meta-model that can be sufficiently expressive for our purpose starting from a purely chemical model:

1. the concept of environment where agents are situated and can move is missing in a model that considers only intercommunicating chemical compartments;
2. the only available mean for changing the system status is the execution of a reaction;
3. the only data item that chemical reactions can manipulate are molecules' concentrations, namely numbers connected to a particular token.

In the following discussion, we will use interchangeably compartment/agent/node and reaction/events as synonyms. As first step, we introduce the environment, absent in chemistry-derived

SSAs, as first class abstraction. The environment has responsibility to provide, for each compartment, a set of compartments that are its neighbours. The rule which is applied to determine whether or not a node belongs to another node's neighbourhood can be arbitrarily complicated. Also, it is responsible of exposing possible physical boundaries, namely, to limit the possible movements of compartments situated within the environment.

The fact that reactions are the only abstraction the modeller can rely upon in order to let the simulated system progress is not a difficult problem by itself. In fact, nothing prevents to widen the generality of a reaction by defining it as: "a set of conditions that, when matched, trigger a set of actions on the environment".

With this definition in mind, a condition is a function that associates to each possible state of the environment a numeric value ranging from zero to positive infinity. If such value is zero, the event can not be scheduled; otherwise, it is up to the reaction to interpret the number: it can influence or not the time at which the reaction will be scheduled, depending on the specific reaction implementation. In case we desire to re-build the original chemical model, we would define a condition for each of the molecules on the left-hand side of the chemical reaction that return the number of molecules currently available in the local compartment. Also, we would define the reaction in such a way that it correctly interprets the number returned by the conditions as concentration of each reactant and correctly applies the rules for computing a propensity to be used to influence the reaction speed [Gil77].

In this framework, actions are arbitrary changes of the environment. In case of pure chemistry, the actions of a reaction would be one for each reactant (that must be removed from the local compartment) and one for each product (that must be added to the local compartment). In case of an extended model considering also multiple compartments, an action should be programmed to be responsible of transferring molecules from a node to a neighbouring one.

Both conditions and actions must expose the set of possible data items (molecules) that they may read or modify: this is necessary in order to allow the dependency graph to be built. Also, both conditions and actions must expose a context of the type `local`, `neighborhood` or `global`; it will be used internally to determine the input and output contexts for the reaction itself.

The reactions are responsible of computing their expected execution time. The engine may require such putative time to be updated in two cases: i) the reaction has just been executed or ii) a reaction on which this reaction depends on has been executed. In case of update required, the reaction should leverage a separately defined time distribution to compute the next putative execution time, possibly feeding the time distribution with a summary of the data gathered from conditions. In case of a Poisson process, a negative exponential time distribution initialised with  $\lambda = r$  should be used, for instance. In case of a repetitive event, such as a timer, a Dirac comb may be used.

In this model the atomicity of the reactions represent a double edged sword: on the one hand, they allow for arbitrarily complex behaviours to be ordered and executed within a DES, on the other hand they make it difficult to model events that last in time, e.g. to simulate devices with limited computational power on which some complex task takes a not negligible amount of time

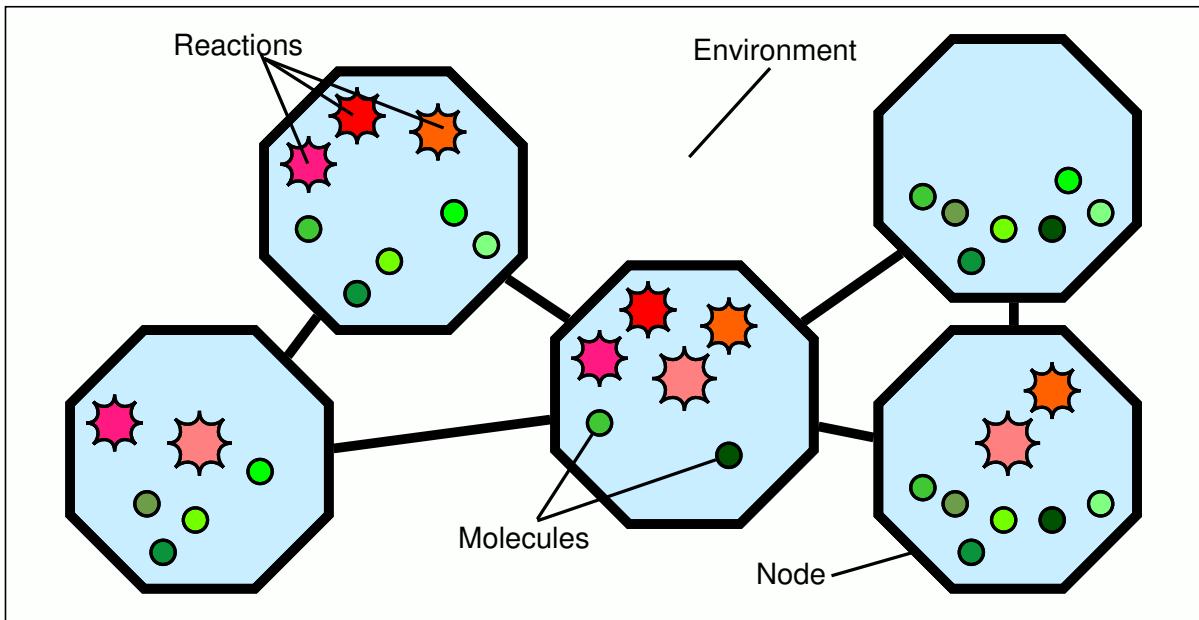


Figure 5.2: ALCHEMIST computational model: it features a possibly continuous space embedding a linking rule and containing nodes. Each node is programmed with a set of reactions and contains a set of structured molecules.

to get completed. To support something similar, two reactions should be defined: one to trigger the start of the computation, and another to actually run it and complete.

The low expressive power of the classical concentration is probably the hardest challenge to tackle when trying to extend a chemical-born meta model towards a richer world, where data items can be complex structures and not just simple numbers. We have found no trivial solution to this issue; instead, we propose to make the definition “concentration” depend on the actual meta-model: let “concentration” be “the data items agents can manipulate”. Besides the trivial example of chemistry, where data items are integer numbers, let’s consider a distributed tuple spaces model: in this case, the molecules would be tuples, and the concentration would be defined as “the set of tuples matching a certain tuple”. Clearly, such flexibility comes with a cost: since the conditions and actions operating on different concentrations dramatically change their behaviour, for any possible class of data items the meta-model must be instanced with a proper set of conditions and actions that can act on such “concentration”. We call this set of concentration-specific instances of conditions and actions an “incarnation”. This sort of model inheritance justifies our double “meta” level: the model described is a meta-meta-model, an incarnation is a meta-model, and, finally, a specific scenario instancing one of such meta-models is a model.

A pictorial representation of the underlying meta-meta-model is shown in Figure 5.2. In this vision of the world, an environment is a multi dimensional space, continuous or discrete,

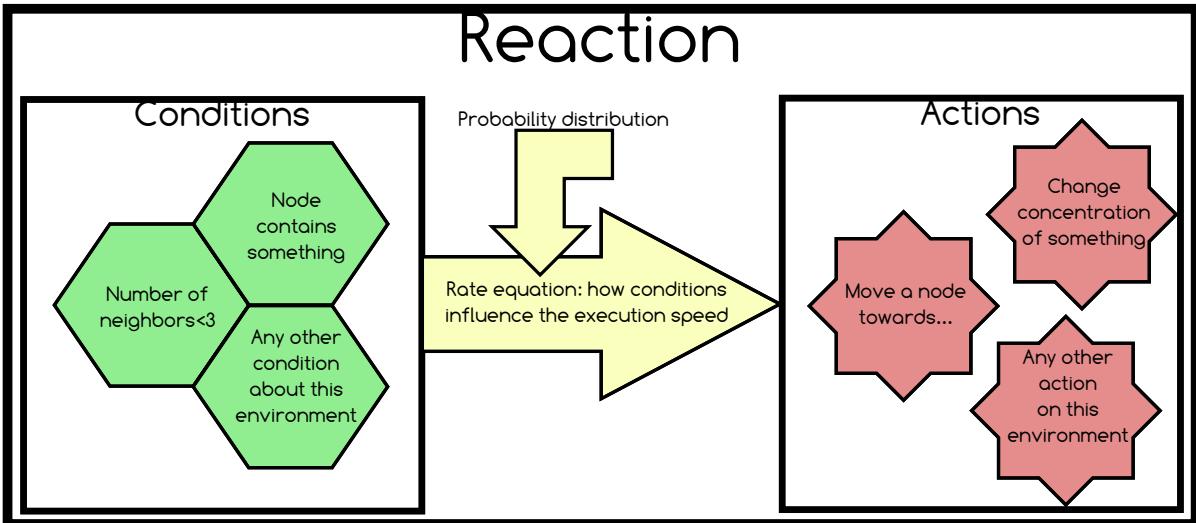


Figure 5.3: ALCHEMIST model of reaction: a set of conditions on the environment that determines whether the reaction is executable, a rate equation describing how the speed of the reaction changes in response to environment modifications, a probability distribution for the event and a set of actions, which will be the neat effect of the execution of the reaction.

which is able to contain nodes and which is responsible of linking them following a rule. The environment may or not allow nodes to move. Nodes are entities which can be programmed with a set of reactions, possibly changing over time. They also contain molecules, each one equipped with a data structure generalising on the concept of “concentration”.

The concept of reaction is graphically depicted in Figure 5.3. It allows, for example, to model reactions which are faster if a node has many neighbours, or also reactions that resemble complex biological phenomena such as the diffusion of morphogenes during embryo development as described in [LIDP10]. It also allows to define which kind of time distribution to use to trigger reactions: this enables us to model and simulate systems based on Continuous Time Markov Chains (CTMCs), to add triggers, or also to rely just on classical discrete time “ticks”.

## 5.3 Architecture

The whole framework has been designed to be fully modular and extensible. The whole engine or parts of it can be re-implemented without touching anything in the model, and on the other hand the model can be extended and modified without affecting the engine.

It is important to note that there is no restriction about the kind of data structure representing the concentration, which can in fact be used to model structured information: by defining a new kind of structure for the concentration, it is possible to incarnate the simulator for different specific uses. For example, by assessing that the concentration is an integer number, representing

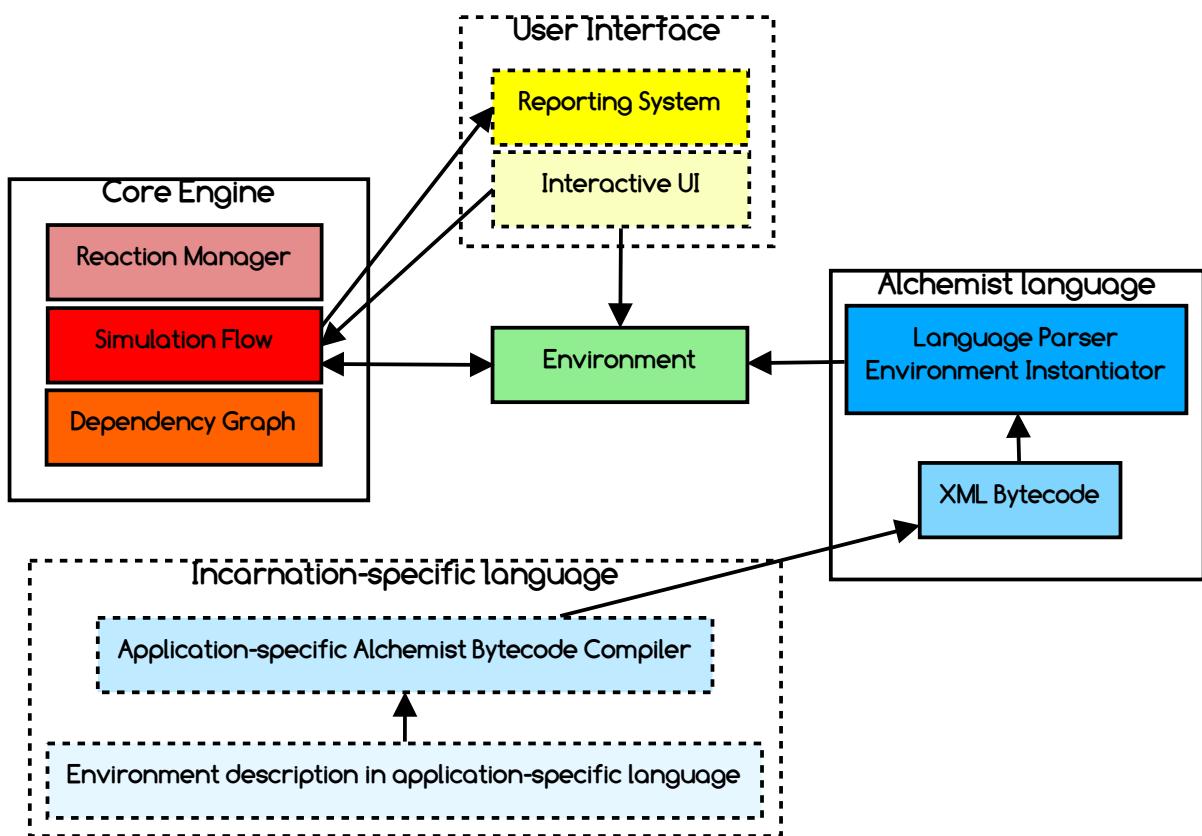


Figure 5.4: ALCHEMIST architecture. Elements drawn with continuous lines indicates components common for every scenario and already developed, those with dotted lines are extension-specific components which have to be developed with a specific incarnation in mind.

the number of molecules currently present in a node, ALCHEMIST becomes a stochastic simulator for chemistry. A more complex example can be the definition of concentration as a tuple set, and the definition of molecule as tuple template. If we adopt this vision, ALCHEMIST can be a simulator for a network of tuple spaces. Each time a new definition of concentration and molecule is made, a new “incarnation” of ALCHEMIST is automatically defined. For each incarnation, a set of specific actions, conditions, reactions and nodes can be defined, and all the entities already defined for a more generic concentration type can be reused.

### 5.3.1 Writing a simulation

In order to write a simulation, the user must have, or implement herself, an incarnation of ALCHEMIST, as described in Section 5.3.

As shown in Figure 5.4, the simulations are written in a specific XML language containing a complete description of environment and reactions. This code is interpreted in order to produce an instance of an environment: once it is created, no further interpretation is needed in order to run the simulation. This XML code is not meant to be directly exploited by users: the XML format itself is not exactly human friendly, and XML file is often considerably big (a few megabytes are considered to be normal) However, it is a very standard way of describing environments in a machine-friendly format. The idea behind this choice is that ALCHEMIST is flexible enough to be used in various contexts, each one requiring a personalised language and a different instantiation of the meta-model. It is up to the extensor to write a translation module from its personalised language to the ALCHEMIST XML. Of course, it is also possible to code the simulation behaviour directly with Java, although this way exposes the user to many more low level details.

### 5.3.2 Implementation details

The framework was developed from scratch using Java. Being performance a critical issue for simulators, we compared some common languages in order to evaluate their performance level. Surprisingly, Java performance are at same level of compiled languages such as C/C++ [BSB<sup>+</sup>03, ORAI11]. The Java language was consequently chosen because of the excellent trade off between performance, easy portability and maintainability of the code, and the high-level support for concurrent programming at the language level. The COLT Java library<sup>1</sup> provided us the mathematical functions we needed. In particular, it offers a fast and reliable random number generation algorithm, the so called Mersenne Twister [MN98].

ALCHEMIST is still actively developed and currently consists of about 200 classes for about 20'000 lines of code. Though still in beta version, it is released with GPL license as open source<sup>2</sup>.

---

<sup>1</sup><http://acs.lbl.gov/software/colt/>

<sup>2</sup><http://alchemist.apice.unibo.it>.

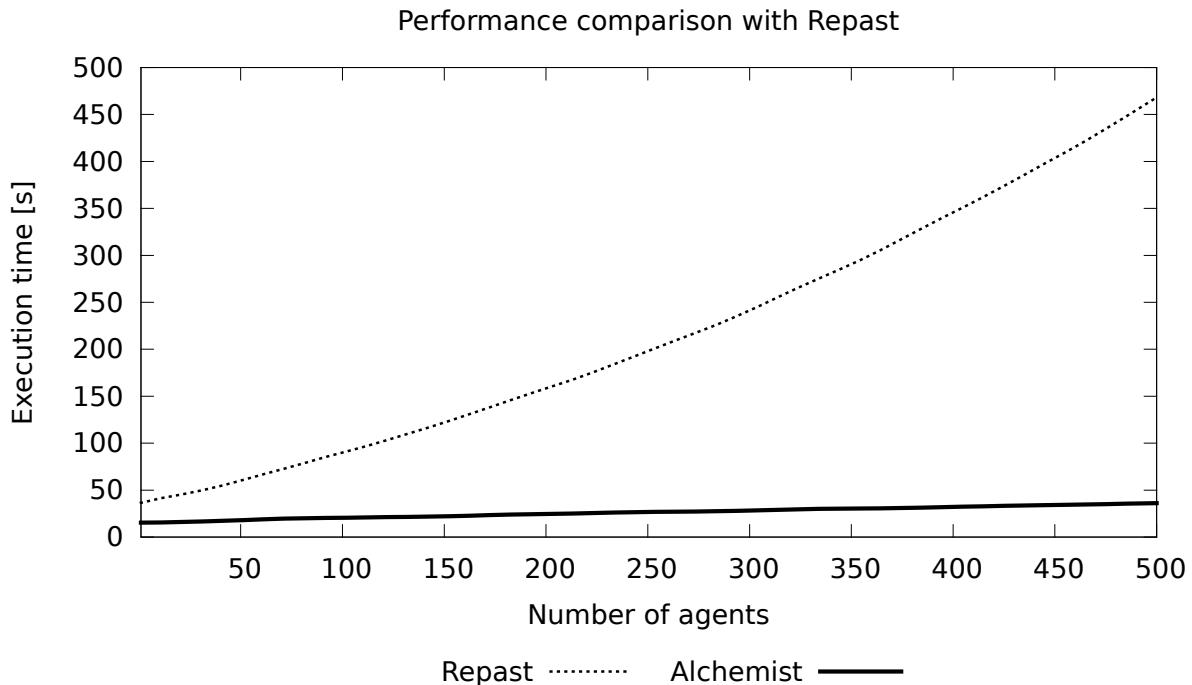


Figure 5.5: Chart showing the performance scaling of ALCHEMIST

## 5.4 Performance

It is possible to evaluate and compare the performances of ALCHEMIST with respect to some known MABS. We exemplify it considering Repast, which we used to developed an alternative simulation for the case study presented in Section 6.2<sup>3</sup>. There are some important facts that deserve discussion here.

First, since there is no built-in support for stochastic simulation in Repast, we choose to collapse the last five laws of Figure 6.1 into a single code path, and the same was made for ALCHEMIST by defining a new action. In that way, we were able to avoid for this very specific case the need of a full fledged dependency graph, because there will always be exactly one **source** and one **field** per sensor, and no reactions need to be enabled or disabled. This crippled most of the effectiveness of ALCHEMIST’s dependency graph, which is indeed a source of optimisation not natively existing in Repast—developing a dependency graph for stochastic simulation in Repast is out of the scope of this thesis, though it would be an interesting subject for future work. On the other hand, it would have been unfair to compare our optimised version against just a plain Gillespie’s algorithm built upon Repast.

The second important point is that we choose to encode all the data in both Repast and

<sup>3</sup>The source code of the simulation we developed is publicly available at <http://apice.unibo.it/xwiki/bin/view/Alchemist/JOS/>.

ALCHEMYST as an array of double values instead of real tuples. For ALCHEMIST, this meant that we developed a new incarnation for the precise scope of this performance test. The choice of encoding data that way made things faster (no matching required), but also less general. This was done because the SAPERE incarnation of ALCHEMIST, which allows us to write the laws as in Figure 6.1, requires a matching system that is not easily portable to Repast.

We choose the configuration of Figure 6.3 and we run multiple simulations varying the number of agents per group, in order to evaluate how the two systems scale with the problem size. We measured the running time required to our testbed to run the simulation from the time zero to the time 100. No graphical interface were attached to the simulators while running the batch, in order to evaluate only the raw performance of the engine. The system we used was an Intel Core i5-2500 equipped with 8GB RAM, with Sabayon Linux 7 as operating system and Sun Java HotSpot™ 64-Bit Server version 1.6.0\_26-b03 as Java Virtual Machine. Results are shown in Figure 5.5.

Results show the simulator built upon the ALCHEMIST framework to be at least twice faster and to scale better than the one built on Repast. Being the dependency graph optimisation cut off as explained above, the reasons for such a difference can lay on the internal scheduler of the engine or in the optimisations in the model. For the first point, we used the default scheduler of Repast Symphony, which is a binary heap implemented through a plain Java array, while our implementation relies on the algorithm and data structures already presented in Section 5.1, so there is not a substantial efficiency difference. For the latter point, a big difference in terms of performances is due to the heavy optimizations of the neighbourhoods of the default ALCHEMIST continuous environment. Since the concept of neighbourhood was part of the computational model, it was possible to adopt caching strategies in order to ensure a fast access to the neighbourhood and a quick execution of the operations on it. This is probably the component which gives the highest impact in this case, since most interactions occur among an agent and its neighbourhood.

## 5.5 Collocation in literature

### 5.5.1 ALCHEMIST as a DES

ALCHEMYST is a discrete event simulator (DES), since it combines a continuous time base with the description of system dynamics by distinguished state changes [Zei76]. Since it allows for in-simulation modifications of the environment, it can be considered to belong to the fourth generation of DESs according to [BW08]. The work in [Pol89] surveys the classical DES approaches: according to it, ALCHEMIST belongs to the category of simulators featuring “internal clock”, “next-event time advance”, and adhering to the “Event-scheduling World View”, namely, the simulator handles events and is concerned with their effect on system states.

Apart from the meta-model adopted, the main innovative aspect of ALCHEMIST with respect to the general DES approach is its ability of optimising the “compile current event list” stage

[Pol89], which ALCHEMIST quickly executes incrementally by means of the management of dependencies that the adoption of a chemical-like model enables.

As far as the simulator model is concerned, instead, the class of DES more related to our approach are those commonly used to simulate biological-like systems. A recent overview of them is available in [EMRU07], which takes into account: DEVS [Zei84], Petri Nets [Mur89], State Charts [Har87] and stochastic  $\pi$ -calculus [Pri95]. Such an overview however emphasises that all such models require a considerable effort to map biological components into abstractions of the chosen formalism: this is because none of them was specifically developed with biology or bio-inspiration in mind. As described in Section 5.2, our model is meant to overcome such limitations, since all the enhancements to the basic chemical model we support can be seen as a generalisation of the abstractions of the works presented in [EMRU07], and add to them the possibility of customising with much greater flexibility a simulation to the scenarios of bio-inspired computational systems.

We should finally note that the DES approach typically contrasts the use of mathematical techniques (e.g. modelling the system by differential equations). However, the possible different choice of translating a system model to ordinary or partial differential equations – which can be solved numerically in a time considerably shorter than a set of stochastic simulations – is shown to be impractical in our case. This path, explored for example by [MTUG09], can provide good results for dynamics that progress at more or less the same speed, and in which abundance of species (data-items or agents in our case) is so high that it can be relaxed to real-valued variables [EMRU07]. This is not the case for many scenarios even in system biology (see, for example, [Cow00, UP05]), not to mention scenarios like pervasive computing where reaction rates do not change continuously, and where the effect on an even small number of agents can be key to a given system evolution.

### 5.5.2 ALCHEMIST as a MABS

Even though chemical-inspired, the meta-meta-model described in Section 5.2 holds evident relationships with multi-agent-based simulation (MABS) approaches.

According to [BMV09a], agent-based platforms for simulations can be split in three categories: general purpose frameworks with specific languages (such as NetLogo), general purpose frameworks based on general purpose programming languages (such as Repast [NHC07]) and frameworks which provide an higher level linguistic support, often targeted to a very specific application (e.g. [WBH06]). Each approach has clearly its own advantages and weaknesses. Usually, the more general purpose is the language, the wider is the set of possible scenarios, and the wider is also the gap between the language and the simulated model. This means that an higher level language allows the user to create and tune its simulations in an easier way, on the other hand it often restricts the generality of the tool.

ALCHEMIST is meant to provide a set of meta-models, possibly each with its own domain specific language, still maintaining the possibility to extend or re-implement certain abstractions using the general purpose Java language (defining a new “incarnation”, namely a new meta-

model).

### Advantages

There is a set of applications which are better tackled by ALCHEMIST. In particular, ALCHEMIST is suitable for all those simulation scenarios in which agents have relatively simple behaviour, and the notion of agent-based environment plays instead a fundamental role in organising and regulating the agents' activities [BV07] by both enabling local interactions among the proactive entities [HVUM07] and enforcing coordination rules [MOV09], allowing the modeller to shift her focus from the local behaviour of the single agent to a more objective vision of the whole MAS [SGJ07]. The idea of dealing with a strong notion of environment in multi-agent systems has been deeply developed in a series of work: other than its importance in the simulation context [HVUM07], at the infrastructure level [VHR<sup>+</sup>07] and at the methodological level [MOV09], there have been proposals of meta-models such as A&A [ORV08] and infrastructures such as TuCSoN [OZ99] and CArtAgO [RPV11]. A common viewpoint of all these works is that the behaviour of those passive and reactive components structuring the environment (e.g. *artifacts* in A&A) is well defined in terms of rules expressed as declarative conditions-implies-actions fashion. Accordingly, ALCHEMIST is particularly useful either in computing systems following the chemical inspiration [VCMZ11, VZ10, ZCF<sup>+</sup>11], or for agent-oriented systems where the role of the environment is key, up to be a very dynamic part of the whole system—where network nodes (or, in biological terms, compartments) can move, new nodes can be spawn or be removed from the system, and links can appear or break at runtime as happen e.g. in pervasive computing scenarios [ZCF<sup>+</sup>11].

### Limitations

Inevitably, as an attempt to build a hybrid between agent-based simulators and (bio)chemical simulators, some trade-offs had to be accepted, which ultimately makes ALCHEMIST less suited for certain classes of applications. In particular, a limitation is that ALCHEMIST's agent actions have to be mapped onto the concept of reaction. On the one hand, this makes it rather straightforward to create reactive memoryless agents [BMV09a], whose goal is just to perform rather easy computations resulting in the creation, deletion or modification of information items in the environment. In fact, since it is allowed to program different nodes with different reactions, it is easy to code reactive and context-dependent agents. On the other hand, there is neither out-of-the-box facility nor any high level abstraction useful to define intelligent agents [BMV09a]. The simulator is able to run them provided the user manually writes their whole behaviour as a single Java-written reaction—and properly specifies dependencies with other reactions. Although possible in principle, performing this task too frequently is likely breaking the ALCHEMIST abstraction, making the programmer losing the nice declarative approach that chemistry endorses, and possibly hampering the optimisation techniques that motivated ALCHEMIST.

## 5.6 Meta-model agnostic features

This section analyses the main features of ALCHEMIST that can be exploited regardless the specific “incarnation”. They add a great value to the tool, since they can be used immediately after the creation of a novel meta-model, with no need to specifically write any new code besides some minimal glue.

### 5.6.1 Distributed statistical analysis

As we have seen in Section 2.6, simulation represents a key step in the engineering of pervasive systems. Being a simulation basically a sampling over an enormous probability space, obvious questions accompany these procedures:

- how reliable are the obtained values?
- How is the number of performed simulations chosen?
- How many simulations are required in order to state system properties with a certain degree of confidence?

Moreover, there is frequently a lack of decoupling between the model specification and the definition of the system’s properties of interest: they are often embedded in the model, and their values are obtained via logging operations performed during the simulation process.

This section presents a tool meant to face these challenges, obtained chaining ALCHEMIST, with MULTIVESTA<sup>4</sup> [SV13], a recently proposed lightweight tool extending VESTA [SMA05] and PVESTA [AM11] which allows to enrich existing discrete event simulators with automated and distributed statistical analysis capabilities. The result is thus a statistical analysis tool tailored to chemical-inspired pervasive systems. The benefits obtained by chaining the simulator with MULTIVESTA are:

1. a language (MULTIQUATEX) to compactly and cleanly express systems properties, decoupled from the model specification;
2. the automated estimation of the expected values of MULTIQUATEX expressions with respect to  $n$  independent simulations, with  $n$  large enough to respect a user-specified confidence interval;
3. an interactive plot as well as the generation of gnuplot input files to visualize the obtained results;
4. a client-server architecture to distribute simulations.

---

<sup>4</sup><http://code.google.com/p/multivesta/>

MULTIVESTA is independent on the model specification language: it only assumes that discrete event simulations can be performed on the input model. As described in [SV13], the tool offers a clean interface to integrate existing discrete event simulators, enriching them with a property specification language, and with efficient distributed statistical analysis capabilities.

## MULTIVESTA and MULTIQUATEX

MULTIVESTA performs a statistical (Monte Carlo based) evaluation of MULTIQUATEX expressions, allowing to query about expected values of observations performed on simulations of probabilistic models. A MULTIQUATEX expression may regard more measures of a model, in which case the same simulations are reused to estimate them, thus improving the performance of the analysis tasks. Moreover, the tool has a client-server architecture allowing to distribute the simulations on different machines. A detailed description of MULTIQUATEX and of the procedure to estimate its expressions is out of the focus of this work, the interested reader can deepen her knowledge reading [SV13, AMS06]. Before defining a MULTIQUATEX expression, it is necessary to specify the state characteristics to be observed. This model-specific step “connects” MULTIQUATEX with the model. The state observations are offered via the `rval(obs)` predicate which returns a real number for each observation, specified by the string parameter `obs`.

A MULTIQUATEX expression consists of a set of definitions of *parametric recursive temporal operators*, followed by a list of `eval clauses`. Each `eval` clause relies on the defined temporal operators, and specifies a property whose expected value must be evaluated.

In order to evaluate a MULTIQUATEX expression, MULTIVESTA performs several simulations, obtaining from each a list of samples (real numbers). One sample for each `eval` clause is obtained, thus all the queried measures are evaluated using the same simulations, improving performance. Based on the samples obtained from simulations, MULTIVESTA estimates MULTIQUATEX expressions with respect to two user-provided parameters:  $\alpha$  and  $\delta$ . Considering the case of simple expressions with just one `eval` clause, the estimations are computed as the mean value of the  $n$  samples obtained from  $n$  simulations, with  $n$  large enough to grant that the size of the  $(1 - \alpha) * 100\%$  *Confidence Interval* (CI) is bounded by  $\delta$ . In other words, if a simple MULTIQUATEX expression is estimated as  $\bar{x}$ , then, with probability  $(1 - \alpha)$ , its actual expected value belongs to the interval  $[\bar{x} - \delta/2, \bar{x} + \delta/2]$ . The case of expressions with multiple `eval` clauses is similar, with the note that the `eval` clauses may regard values of different orders of magnitude, and thus the user may provide a list of  $\delta$  rather than just one. After having obtained a sample for every `eval` clause from a simulation, these values are used to update the means of the samples obtained from previous simulations (one mean per `eval` clause). If the CIs have been reached for every `eval` clause, the evaluation of the expression is terminated, otherwise further simulations are performed. Note that each `eval` clause may require a different number of simulations to reach the required CI. Once the CI of an `eval` clause has been reached, such `eval` clause is ignored in eventual further simulations performed for other `eval` clauses.

## Integrating MULTIVESTA and ALCHEMIST

In order to chain the tools, some steps, have been tackled once and for all, while others are model-specific. Essentially, in order to allow the interaction with MULTIVESTA, ALCHEMIST has to fulfill two requirements:

1. the ability to advance the simulation in a step-by-step manner (which is provided by the `playSingleStepAndWait` method in `ISimulation` interface);
2. the ability to analyze the model status after each simulation step, providing measures in form of real numbers about properties of interest.

Since both MULTIVESTA and ALCHEMIST are Java-based, their interaction has been easily realized by subclassing the `NewState` class of MULTIVESTA. The obtained `AlchemistState` class is sketched in Listing 5.1, where unnecessary details are omitted. The new class contains some ALCHEMIST-specific code, providing MULTIVESTA with the simulation control and proper entry points for the analysis.

In the constructor (lines 6-10), the superclass initialization is done by a simple `super()` call. The remaining code initializes ALCHEMIST-specific parameters such as the maximum time or steps to simulate. Those are in general not required, since MULTIVESTA is able to detect when the analysis requirements have been met, and consequently stop the simulation flow. The method `setSimulatorForNewSimulation()` is depicted in lines 12-20. The method is invoked by MULTIVESTA before performing a new simulation to (re)initialize the status of the simulator, providing to it a new random seed. In lines 22-24, `performOneStepOfSimulation()` is provided: resorting to the ALCHEMIST method `playSingleStepAndWait()`, it allows MULTIVESTA to order the execution of a single simulation step.

In order to inspect the simulation state, the `rval()` method defined in lines 26-32 is invoked. The argument specifies the observations of interest. This method inspects the simulation state for all aspects common to any ALCHEMIST model, e.g., in the listing are sketched the current simulated time (line 27) and the number of performed simulation steps (line 28). Clearly, each ALCHEMIST model will have its own observations of interest.

Depending on the model at hand, it may be necessary to refine the model-independent observations exposed by `AlchemistState` with a set of model-specific ones. This can be done by simply instantiating the `IStateEvaluator` interface provided by MULTIVESTA, constituted by one method only.

### Example

This section discusses the analysis performed on a crowd steering scenario, similar to the one which will be presented in detail in Section 6.2. In such scenario, two group of people are inside an indoor environment, and are steered towards two point of interest (POI) such that their ideal trajectory intersect, creating a crowd. The steering system is implemented leveraging the SAPERE meta-model, and as such data items in this scenario are called LSAs. This example

Listing 5.1: AlchemistState extending MULTIVESTA's NewState class

```

1  public class AlchemistState<N extends Number, D extends Number, T> extends NewState {
2      private final long maxS;
3      private final ITime maxT;
4      private ISimulation<N, D, T> sim;
5      ...
6      public AlchemistState(final ParametersForState params) throws ...{
7          super(params);
8          final StringTokenizer otherparams = new StringTokenizer(params.getOtherParameters());
9          /* Initialization of Alchemist-specific parameters
10             and execution environment resorting to otherParams */
11      }
12      ...
13      public void setSimulatorForNewSimulation(final int seed) {
14          /* Stop current simulation, create a new one. */
15          ...
16          sim.stop();
17          ...
18          env = getFreshEnvironment(seed);
19          sim = new Simulation<>(env, maxS, maxT);
20          ...
21      }
22      ...
23      public void performOneStepOfSimulation() {
24          sim.playSingleStepAndWait();
25      }
26      ...
27      public double eval(final String obs) {
28          if(obs.equals("time")) return getTime();
29          if(obs.equals("steps")) return getStep();
30          //other model-independent observations
31          ...
32          return getStateEvaluator().getVal(obs, this);
33      }
34  }

```

is not meant to provide significant scientific insights on crowd steering (thoughts about it will be provided in the remainder of this thesis), but rather to act as toy-example to demonstrate the flexibility and possibilities of the chaining of ALCHEMIST and MULTIVESTA.

The outcome of the analysis is summarized in the three charts of Figure 5.6 (obtained using the gnuplot input files provided by MULTIVESTA), showing, at the varying of the simulated time, the expected values of: the number of people which have reached their point of interest (top), the average number of connections of the devices (middle), and the number of LSAs in the system (bottom).

The three charts have been obtained by evaluating *MainMQ* of Listing 5.2, having 5 parametric temporal operators (lines 1-9): `people@RPOI` and `people@LPOI` regard the number of people which have reached, respectively, the POI on the right and the one on the left, while `people@POI` counts instead how many people have reached their destination. The fourth temporal operator (`avgConn`) regards the connectivity degree of the devices. Finally, `LSAs` regards the number of LSA in the system, giving insights on the amount of memory required to sustain the system. The temporal operators are analysed at the varying of the simulated time

Listing 5.2: The evaluated parametric multi-expression (*MainMQ*)

```

1 people@POI(x) = if {s . rval("time") >= x}
2     then s . rval("bPOI")
3     else #people@POI(x)
4     fi ;
5 people@RPOI(x) = // similar to people@POI
6 people@LPOI(x) = // similar to people@POI
7 avgConn(x) = if {s . rval("time") >= x}
8     then s . rval("degree")
9     else #avgConn(x)
10    fi ;
11 LSAs(x) = if {s . rval("time") >= x}
12     then s . rval("LSAs")
13     else #LSAs(x)
14 fi ;
15 eval parametric(E[people@POI(x)] , E[people@RPOI(x)] ,
E[people@LPOI(x)] , E[avgConn(x)] , E[LSAs(x)] , x , 0.0 , 1.0 , 50.0);

```

from 0 to 50 seconds, with step 1 (line 10). Thus the analysis consisted in the estimation of the expected values of the 5 temporal operators instantiated with 50 parameters, for a total of 250 expected values.

A high degree of precision has been required:  $\alpha$  has been set to 0.01, while the  $\delta$  values (the size of the CIs) have been chosen considering the orders of magnitude of the measures: 0.5 for the instances of `people@POI(x)`, `people@RPOI(x)` and `people@LPOI(x)`; 0.05 for `avgConn(x)`; and 3 for `LSAs(x)`. To reach such a level of confidence, the tool ran approximately 2500 simulations, requiring less than a hour. The discussed confidence intervals are depicted in the aforementioned charts: the two lines drawn above and below the central lines represent the obtained CIs of the expected values, thus indicating the intervals in which the actual expected values lie with probability 0.99.

The top chart regards the first 3 temporal operators: the top plot refers to the number of people in total which have reached their target POI, while the two almost over-imposed lower plots regard the number of people at the right POI and at the left POI. All the three measures under analysis produced monotonically increasing sigmoid curves. We can deduce from this behavior that there are no systemic errors in the crowd steering system, such as people with no or wrong suggested final destination, in which case we would have noticed one or more flattened zones flawing the sigmoid curve. We note also that after around 30 simulated seconds all the people have reached their target, and that it takes around 10 seconds for the fastest walkers to get to their destination. Note that, despite that the analysis has been performed for 50 simulated seconds, the charts presented in Figure 5.6 have been cut at time 35 to better show the most

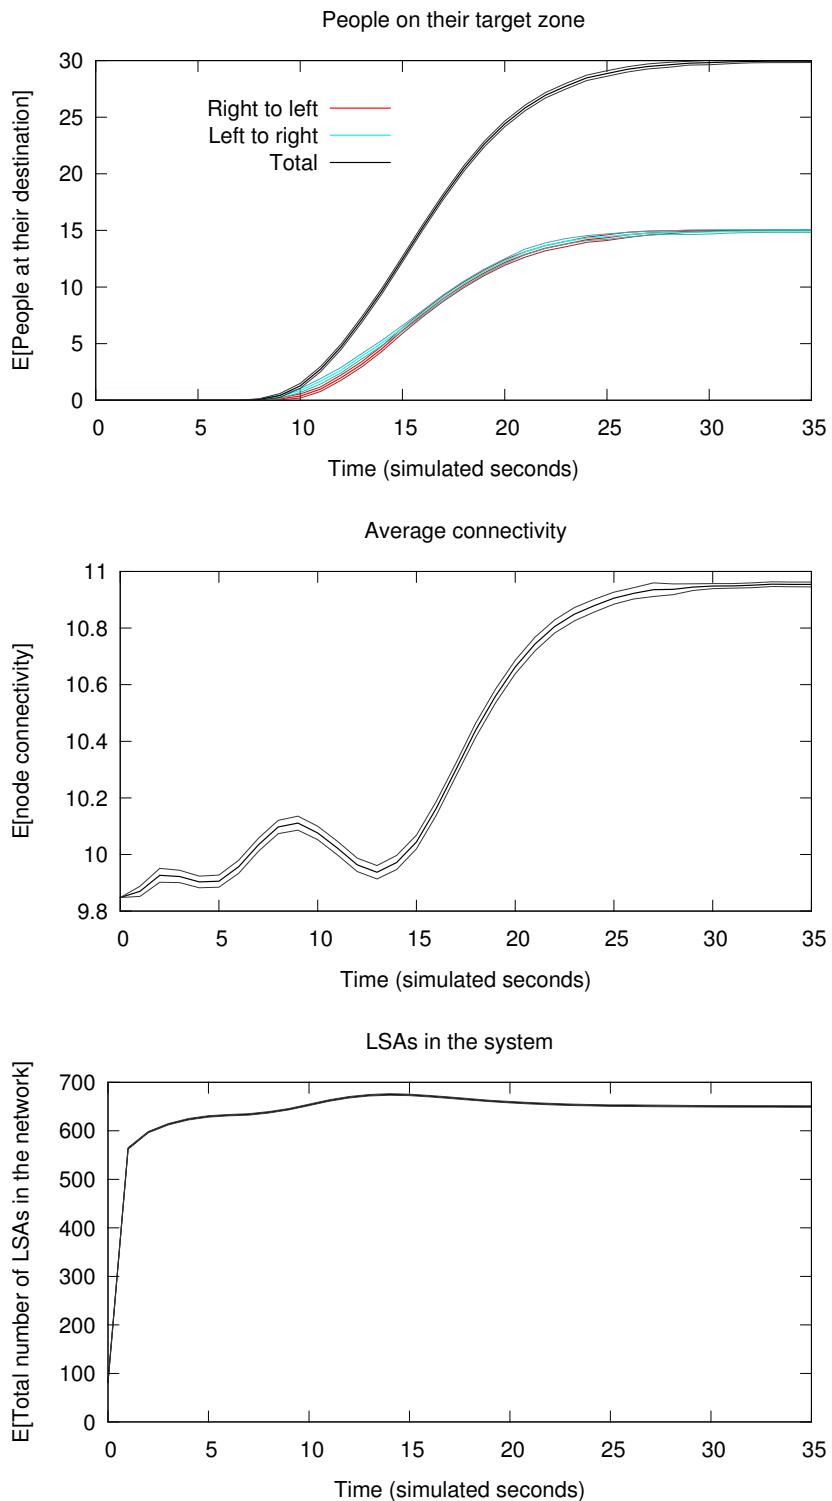


Figure 5.6: Analysis of the crowd steering scenario: (top) number of people at the POIs, (middle) average number of connections per device, (bottom) number of LSAs in the system.

relevant results: in fact, the system reaches stability after this time, and no event of interest happens later. Moreover we can also notice that people going from left to right are slightly faster than the other group. This difference, due to the asymmetric positioning of people in the center of the scenario, is hardly visible and would have been impossible to spot with a lower precision analysis.

The middle chart shows the evolution in time of the average number of connections of the devices (i.e., both sensors and smartphones). Since each device is considered to be connected to all those within a range of 1.5 meters, it also gives us a hint about the crowding level. There is a noticeable peak at around 8 seconds: it is due to a high number of people approaching the central group. After that, the peak disappears when most people overtook the obstacle and are walking towards their POI. Finally, there is a growth: progressively, people reach their destination and tend to create a crowd.

The bottom chart shows the number of LSAs in the whole system, indirectly giving hints on the global memory usage. Once the system is started, there is a very quick growth, due to the gradient being spread from the sources to the whole sensors network and to the LSAs produced by the crowd-sensing. The system reaches a substantial stability after a couple of seconds. From that point on, the number of LSAs has very little variations: the system has no “memory leak”, in the sense that it does not keep on producing new LSAs without properly removing old data.

## Performance Assessment

All the experiments of the previous section have been run on a machine equipped with four Intel® Xeon® E7540 and 64GB of RAM, running Linux 2.6.32 and Java 1.7.0\_04-b20 64bit.

In order to measure the performance scaling of the tool, we ran our analysis multiple times varying the number of MULTIVESTA servers deployed. Results are summarized by the dark line of Figure 5.7, showing that with only 1 server (ALCHEMIST does not have any parallel capability by its own), the analysis required almost 17 hours, while with more than 30 servers (enabled by the chaining with MULTIVESTA) it required less than an hour. For the considered scenario, by distributing simulations we have thus obtained a more than eighteen times faster analysis.

It is worth to note that such comparison of performance is affected by the statistical nature of the analysis procedure. Intuitively, when evaluating a MULTIQUATEX expression resorting to  $x$  or to  $x + y$  servers, in both cases we initialize the first  $x$  servers with the same  $x$  seeds, generating thus the same  $x$  simulations. However, the remaining  $y$  servers are initialized with new seeds, and thus produce new simulations. Clearly, by having different simulations, we may obtain different sample variances, requiring a different number of simulations. For this reason, the bright line of Figure 5.7 depicts the time analysis normalized with respect to the number of simulations (obtained dividing the overall analysis time by the number of performed simulations). The two lines of Figure 5.7 evolve accordingly, thus confirming the great performance gain brought by the distribution of simulations.

MULTIVESTA has a further important feature which dramatically reduces the analysis time: the reuse of simulations to estimate many expected values. The reusage happens on two levels:

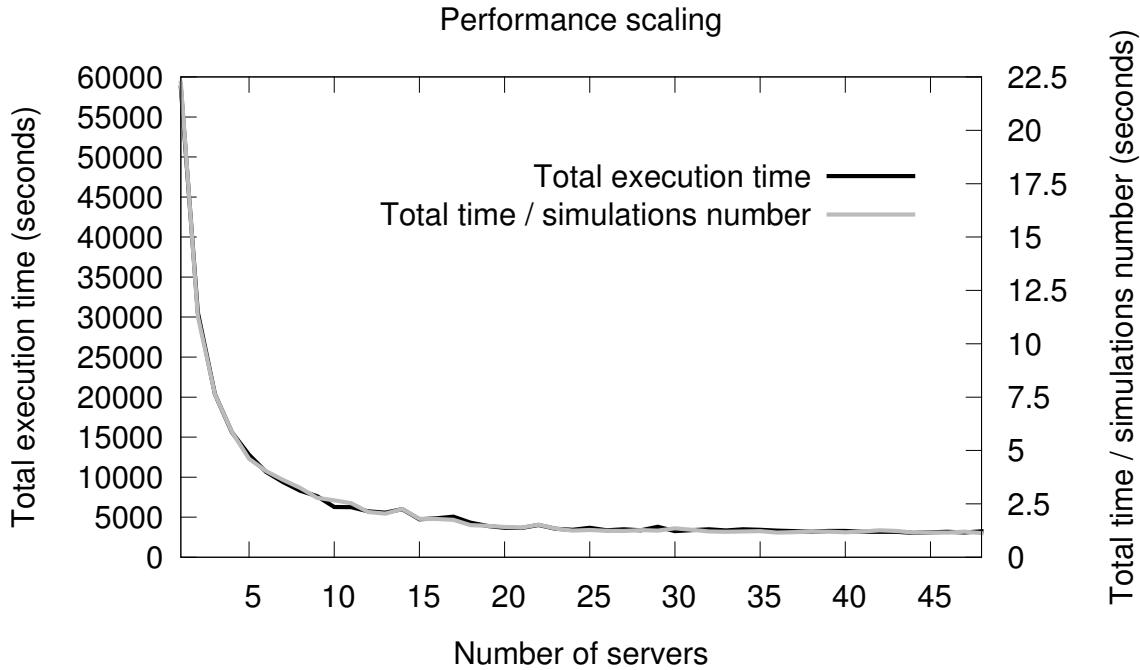


Figure 5.7: Performance scaling with of the number of running servers.

1. expressions can be made parametric, and thus simulations can be reused for computing the same property at the varying of a parameter (e.g., at different time steps);
2. it is possible to write multi-expressions , and thus reuse simulations to evaluate multiple properties.

Moreover, by combining these two features, it is possible to reuse simulations for multiple parametric properties (e.g. *MainMQ* of Listing 5.2). As explicated by Table 5.1 (whose reported analysis have been performed resorting to 48 servers), parametric expressions (second column) allowed us to save a stunning 96% of execution time with respect to the simple expressions case without reuse of simulations (first column). Moreover, the parametric multi-expression feature (third column) allowed us to further cut down this time to a third. We can thus advocate that, by exploiting reuse of simulations, for the considered scenario we obtained a more than 64 times faster analysis. Clearly, the performance gains change depending on the considered scenario or properties. In fact, the more simulations are required, and the longer they take, the greater will be the advantage.

Table 5.1: Time performance improvements reusing simulations (seconds).

	<b>Expressions</b>	<b>Param. expressions</b>	<b>Param. multi-expressions</b>
<b>people@POI</b>	28045.77	2567.26	n.a.
<b>people@RPOI</b>	15891.13	1114.86	n.a.
<b>people@LPOI</b>	15560.83	950.30	n.a.
<b>avgConn</b>	90111.29	1372.64	n.a.
<b>LSAs</b>	58630.45	2504.33	n.a.
<b>Total time</b>	208239.47	8509.39	3215.95

### 5.6.2 Real world maps

ALCHEMIST supports real-world environments. The base functionality lies in the possibility to load maps and simulate within them. Currently, the simulator supports OpenStreetMap[HW08], and it is able to load data in multiple OSM XML, compressed OSM XML and Protocolbuffer Binary Format (PBF). The latter is warmly recommended for both performance and map file size. Various websites provide ready-to-use extracts of the world map in PBF format for cities and whole regions, and arbitrarily sized extracts in OSM XML format can be obtained via public web API.

Once the map has been loaded, ALCHEMIST offers the possibility to enrich the simulations with the maps data, and in particular offers various ways to move nodes within the environment taking into account the characteristics of the map.

The first feature offered is about the initial node displacing. When adding a node to the environment, in fact, it is possible not to displace the node in the exact position indicated, but in the nearest street point. This comes in particularly handy in order to realise the network of static devices: the simulator can be fed with a grid of devices, and it automatically modifies the positions of each node in order to displace it on a street.

The second useful feature is the possibility to rely on the map data to compute routes, as in Figure 5.8(b). ALCHEMIST in fact ships with a module based on GraphHopper<sup>5</sup> which provides the simulator the ability to compute routes between two points of the map. This feature is mainly used to steer nodes correctly along the map, following the allowed ways. It is possible to use such feature also specifying different types of vehicle. Currently, pedestrians, bikes and cars are supported, and can be used together in the same simulation, and even within the same node (for e.g. simulating a pedestrian taking her car, driving and then walking again). Another usage of this system is, for instance, the possibility to use the route distance or expected route time as data items when performing internal calculations, e.g. to compute a spatial gradient where the distance is not the classic euclidean distance but rather the distance computed by the routing subsystem.

The simulator also allows for loading GPS traces. The traces must be in a personalised binary

---

<sup>5</sup><http://graphhopper.com/>

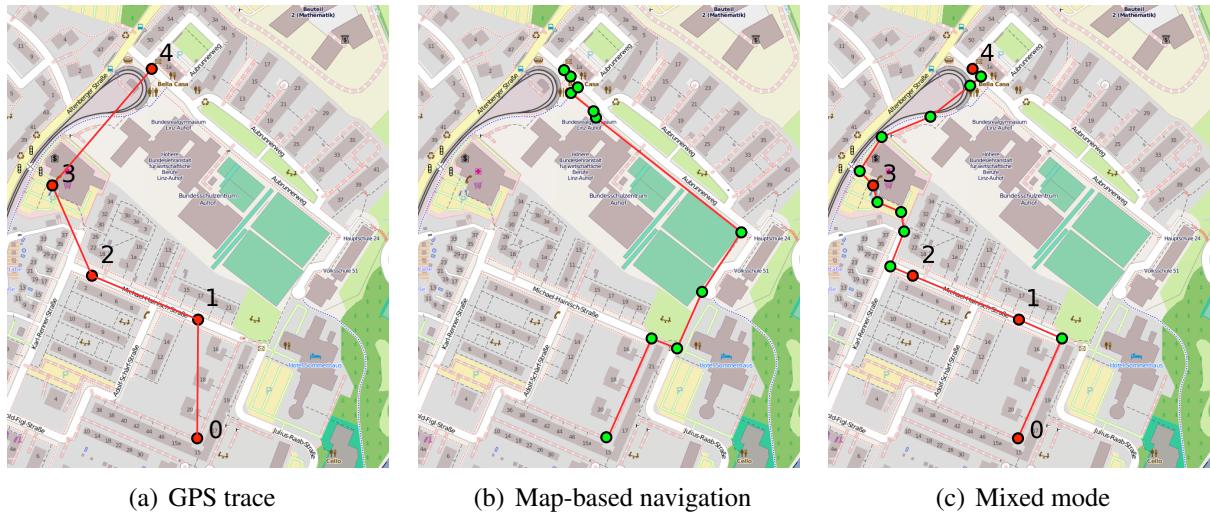


Figure 5.8: Three snapshots showing the different navigation modes available in ALCHEMIST. In Figure 5.8(a) there is a GPS trace consisting in four points. If not specified to behave otherwise, ALCHEMIST computes the average speed between two points, and moves the node along a straight line with constant speed. This, depending on the precision of the trace, may generate paths that cross over buildings. In Figure 5.8(b), the route between the start and end point is computed using the built-in navigator. This could lead the nodes towards paths different from those described by the route. In Figure 5.8(c), the routing subsystem is used to refine the quality of the GPS traces: ALCHEMIST will move the node with constant speed between the two points, using map based routing when necessary to improve precision between two consecutive trace points.

format, fortunately, however, this format is easy to generate. In fact, it is just a Java object stream containing a List of “IGPSPoints”, namely a simple structure with latitude, longitude and time. A converter from JSON to such format is also available in the simulator distribution. Just as the previously mentioned routing system, the GPS traces can be used to move nodes along the scenario, reproducing existing paths, as in Figure 5.8(a).

It is also possible to use a combo of the two techniques: often, the GPS traces could be a bit rough with respect to desired grain of the simulation. In these cases, the navigation subsystem can be used to compute the route between two GPS points, with the assumption that the user followed the paths allowed in the map. In this way, it is possible to arbitrarily refine the grain of a GPS trace without the disadvantage deriving from a simpler interpolation, namely the possible transit over physical obstacles. Such mixed navigation mode is depicted in Figure 5.8(c).

Obviously, it would be rather hard to understand what is going on on the simulation without proper rendering support. In this sense, ALCHEMIST automatically detects when a real-world environment is being used, and renders a map relying on MapsForge<sup>6</sup>.

## 5.7 Chemical meta-model

### 5.7.1 Example scenario: Morphogenesis

As example to test the power of a chemical stochastic simulator enriched a more general model, we propose a case study on morphogenesis. Development of multicellular organisms begins with a single cell – the fertilised egg, or zygote. The egg cell is always asymmetric, *i.e.*, the distribution of maternal proteins inside the cell is not uniform. After fertilisation it divides mitotically to produce all the cells of the body. The resulting blob of cells starts the process of differentiation which is initially caused by the presence of *maternal factors* located in specific areas of the organism’s egg. The ordered spatial organisation of this diversity is then caused by the interactions among cells. These can be direct or mediated by specific diffusing proteins called *morphogens*. After that, tissues have been created, and the formation of organs begins so as to originate the final shape of the organism [AJL<sup>+</sup>02, Gil06].

The main issues of Developmental Biology, nowadays only partially solved, are:

- which are the processes and the genetic mechanisms that control cellular (or nuclear) duplication;
- which are the processes and the genetic mechanisms by which cells differentiate;
- how it is possible that cellular differentiation is spatially organised.

It is clear that the macroscopic emergent results of pattern and shape formation originates from the microscopic mechanisms of intra-cellular reactions and environmental morphogen

---

<sup>6</sup><https://code.google.com/p/mapsforge/>

diffusion. But how these mechanisms coordinate is still under investigation through experimental techniques and theoretical models.

A well known example of multicellular development is given by *Drosophila Melanogaster*, whose spatial organisation results in the segmented pattern of gene expression shown in Figure 5.9. To show the spatial-temporal evolution of the pattern, the organism has been object of several modelling attempts, most of which focus their investigation in gene interactions and protein diffusion mechanisms, bounding the window of analysis at a period of development that does not massively involve nuclear or cellular divisions. To cite few of them, in [RS95] the change of protein concentration is modelled as an Ordinary Differential Equation depending on the process of gene regulation, diffusion over a discretised space and decay. Mitosis are also modelled as a synchronous atomic process that suspends the synthesis of protein and causes a doubled number of cell nuclei. In [GJK<sup>+</sup>04] it is presented a continuous mathematical model based on a set of coupled non linear reaction-diffusion Partial Differential Equations that describe how protein concentrations change over time and space as a result of the mechanisms of protein synthesis and degradation, gene inhibition and activation, protein diffusion. They replace embryo's cellular structure with a continuum and do not model nuclear divisions. In [MV10] it is described a stochastic and discrete version of [GJK<sup>+</sup>04] based on a set of interacting compartments inside which chemical reactions implementing the gene regulatory network can occur. It does not count for changes in the network topology due to compartment movement and mitosis. These models run over a simulation window that start from cleavage cycle 11 to cleavage cycle 14. In [LIDP10] the gradient formation of the *Drosophila bicoid* protein, mainly caused by the morphogen diffusion, is simulated with an innovative stochastic model of reaction-diffusion systems implemented into a Gillespie-like stochastic simulation algorithm. They do not model neither gene interactions nor cellular / nuclear divisions.

To capture a bigger window of embryo development we built on top of ALCHEMIST's chemical meta-model a model of *Drosophila* development so as to capture both nuclear division and molecular processes such as morphogen diffusion and gene regulatory networks.

### Development of *Drosophila Melanogaster*

*Drosophila* is one of the best known multicellular organism. The egg of *Drosophila* is already polarised by differently localised mRNA molecules which are called *maternal effects*. The early nuclear divisions are synchronous and fast (about every 8 minutes): the first nine divisions generate a set of nuclei forming the *syncytial blastoderm*. All the dividing nuclei share a common cytoplasm, and material can diffuse throughout the embryo. After other four nuclear divisions, during the fourteenth nuclear division, plasma membranes grow to enclose each nucleus, converting the syncytial blastoderm into a *cellular blastoderm* consisting of about 6000 separate cells. After the first ten divisions the time required to complete each of the next four divisions becomes progressively longer: cycle 13 takes for instance 25 minutes. During cycle 14 cells conclude their mitosis in rather different time: some cells take 75 minutes, other 175 minutes to complete this cycle. The rate of division is then constant in the first hours of

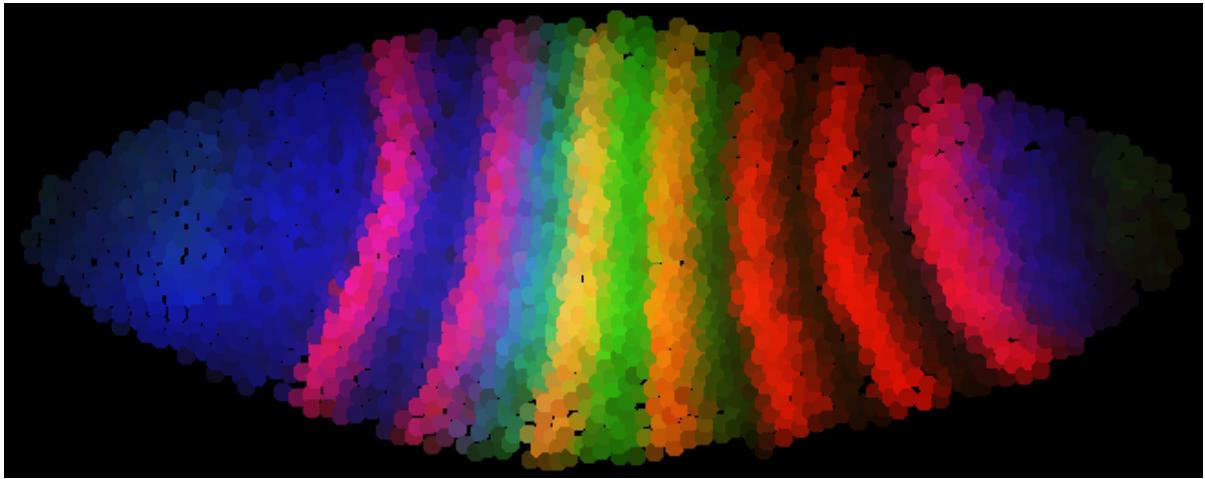


Figure 5.9: The pair-rule gene *even-skipped* (red) together with *hb* (blue) and *Kr* (green) in *Drosophila* embryo at the cleavage cycle 14A temporal class 8. Reconstructed image from [PPSR09]. Embryo name: ba3.

development ( $9.05 \text{ min}^{-1}$ ), then decreases until a low value ( $0.2 \text{ min}^{-1}$ ). The transcription of RNA massively begins during cleavage cycle 14 so that the embryo enters in the mid-blastula transition.

The studies on the genetics at the basis of the segmentation in the anterior-posterior body plan identified a hierarchy of genes that controls segment determination: maternal effects, gap genes, pair-rule genes and segment polarity genes. In Figure 5.10 an illustrative subset of this hierarchy is shown —in the images of this paper we hereafter refer at the anterior pole as their left side, and the posterior as their right side. Maternal effect genes are the building blocks of the anterior-posterior pattern. The most important are *bicoid* (*bcd*)—forming an anterior-to-posterior gradient—and *caudal* (*cad*)—forming a posterior-to anterior gradient. The mRNAs of such genes are placed in different regions of the egg and initiate the hierarchy of transcription, driving the expression of *gap genes*, which are the first zygotic genes to be expressed. The first is *hunchback* (*hb*), that appears early in the embryo development such that sometimes is classified among maternal genes. The basic others are *Krüppel* (*Kr*), *knirps* (*kni*) and *giant* (*gt*). These genes are expressed in specific and partially overlapping domains. Differing combinations and concentrations of the gap gene proteins then regulate the expression of downstream targets, *i.e.*, the *pair-rule genes*, which divide the embryo into a striped pattern of seven segments. The most important pair-rule genes are *even-skipped* (*eve*) and *fushi-tarazu* (*ftz*). The pair-rule gene proteins activate the transcription of the *segment polarity genes*, whose protein products specify 14 parasegments that are closely related to the final anatomical segments [AJL<sup>+</sup>02].

We developed a model that reproduce the process of *Drosophila* development from the fertilised egg until the pattern formation of the *gap genes* during cleavage cycle 14. The whole embryo is modelled as a big cell where nuclei grow and move over a 2D continuous environment filled with diffusing proteins. The mechanisms we explicitly model and we reproduce with

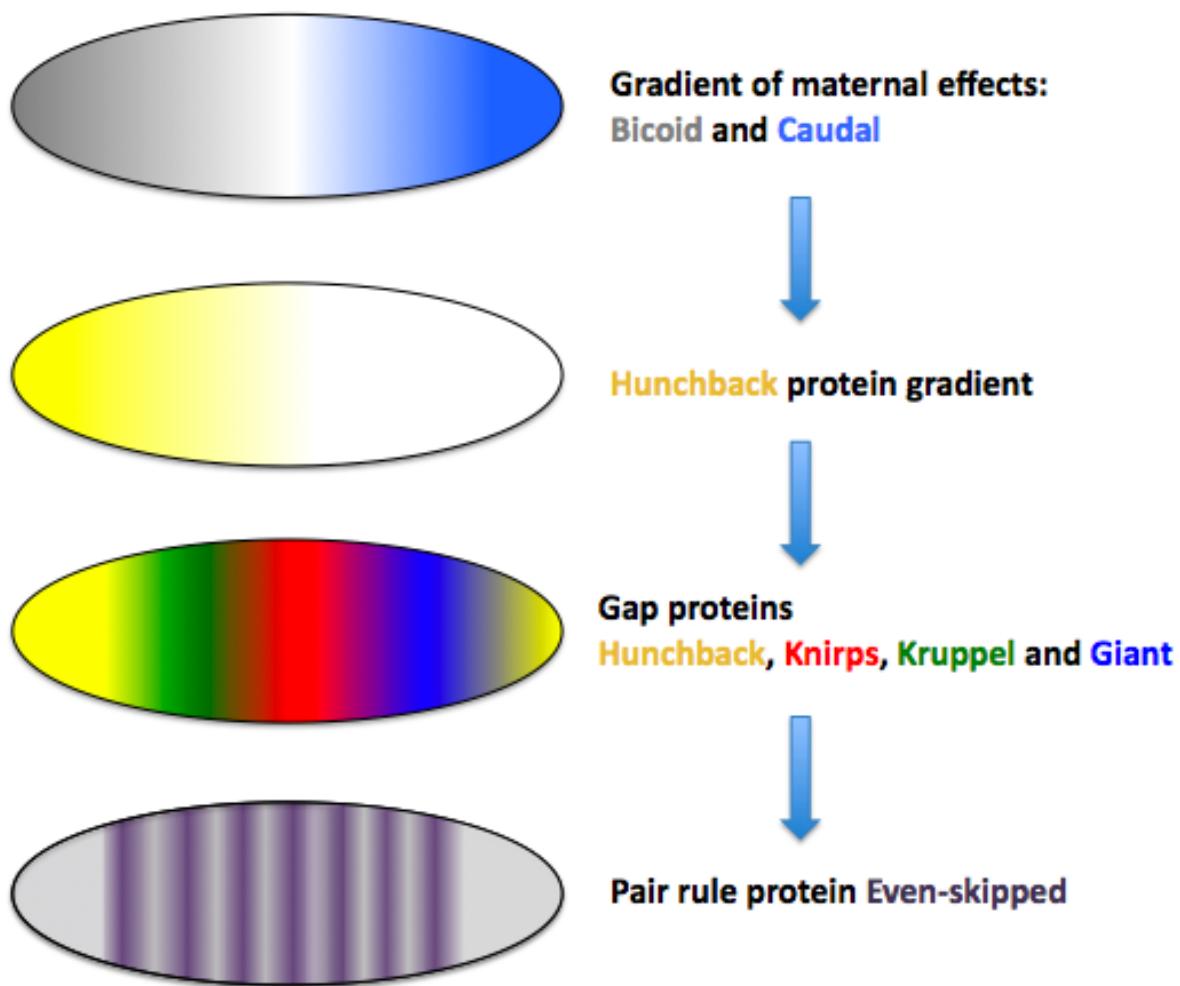


Figure 5.10: Hierarchy of genes establishing the anterior-posterior body plan

simulation are:

- nuclear divisions;
- nuclear migration;
- morphogen diffusion;
- gene interactions.

We model the process of nuclear divisions as a single chemical like reaction, whose precondition is given by the maximum number of other nuclei in the neighbourhood and whose product is a new nucleus. The new nucleus is created close to the dividing one in a casual direction and owns half of its molecular content. The rate of nuclear division is determined according to the rate observed in the real system. Since experimental data show a certain synchronisation among dividing cells, this phenomena has been modelled through a non-Markovian reaction, relying instead on a “drifting” Dirac Comb, namely on a distribution whose events happen every increasing time interval.

Movement of nuclei is based on biomechanical forces of repulsion among neighbouring nuclei. They are in fact constrained to remain within the membrane-delimited area so as to filling pretty homogeneously the available space. If two nuclei are closer than a distance given as a parameter, a new position for them is computed. The nearer two nuclei are, the stronger is the repulsion.

We support diffusion from / to nuclei and inside the environment along the  $x$  and  $y$  axis. For this purpose environment is discretised into a grid of locations. Molecules move from one nucleus into a neighbour location of the environment picked up probabilistically among the whole neighbourhood, or, the other way round, from one location of the environment into a neighbouring nucleus. Diffusion among locations of the environment follows the same law, implementing a Brownian motion.

In Figure 5.11 it is shown the network of interactions among maternal effectors and gap genes we considered. As in [PJRG06], an other gap gene, *tailless* (*tll*), also appears as input of the network whose regulation is not clear and we do not represent in the model.

### Simulation in ALCHEMIST

Simulations results are shown in Figure 5.12. They are evaluated observing the time evolution of the compartment number and of the gene expression pattern. The time evolution of cells number is compared with data we have from literature and described above. Simulations results for gap genes expression are evaluated according to experimental data available online in the FlyEx database [SKK<sup>+</sup>08]<sup>7</sup> and reported in Figure 5.13. They provide only qualitative information and their evaluation is based on the expression area of the different genes.

---

<sup>7</sup><http://urchin.spbcas.ru/flyex/> – last visited on December 2014

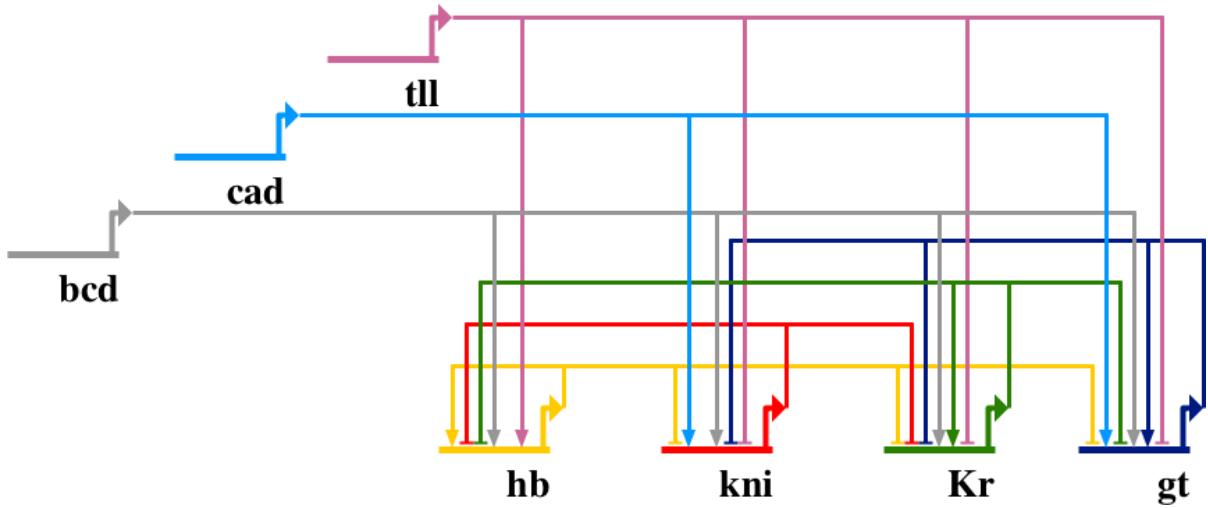


Figure 5.11: Gene regulatory relationships as in the model presented in [PJRG06, RPJ96, GJK<sup>+</sup>04]. The diagram is realised with BioTapestry [LDB09]. The type of link is pointed out by its shape: links with arrowhead are enhancers, while the repressor links have a foot.

The snapshots in Figure 5.12 show one side of the embryo along the anterior-posterior (A-P) axis where almost half of the nuclei are located. The cell number is shown in the label on top of each snapshot. The horizontal axis represents the A-P position and is labelled as % of embryo length.

The first snapshot shows the initial condition with only one nucleus and the egg polarised by maternal effects localised in the extreme pole: *bcd* on the left and *cad* on the right. During the first minutes of simulation only nuclear divisions occur so as to fill the whole maternal cell at the end of cleavage cycle 9 with around 250 nuclei (half of the total 500), as shown in the second snapshot. Finally in the third snapshot it is reproduced the expression pattern of the four gap genes, whose spatial organisation is compared with experimental data of Figure 5.13. Cells are coloured of yellow, red, blue and green if, in order, they express *hb*, *kni*, *gt* and *Kr*, and their size is proportional to the protein concentration. Gene *hb* is massively expressed in the anterior half of the embryo and a small segment appears in the extreme left: simulations correctly reproduce the main expression domain of *hb* while, even if its expression is observable in the posterior pole, it does not form a clear segment. Gene *kni* is mainly expressed between 60% and 80% of the embryo length, either in real embryo and in simulations, as well as gene *Kr*, which is expressed between 40% and 60% of the A-P axis. The expression of *gt* is finally observable between 10% and 30% and 80% and 90%. The qualitative results presented here are not sufficient for observing the expression in the anterior half of the embryo as soon as it totally overlaps the *hb* expression, while the posterior segment clearly appears.

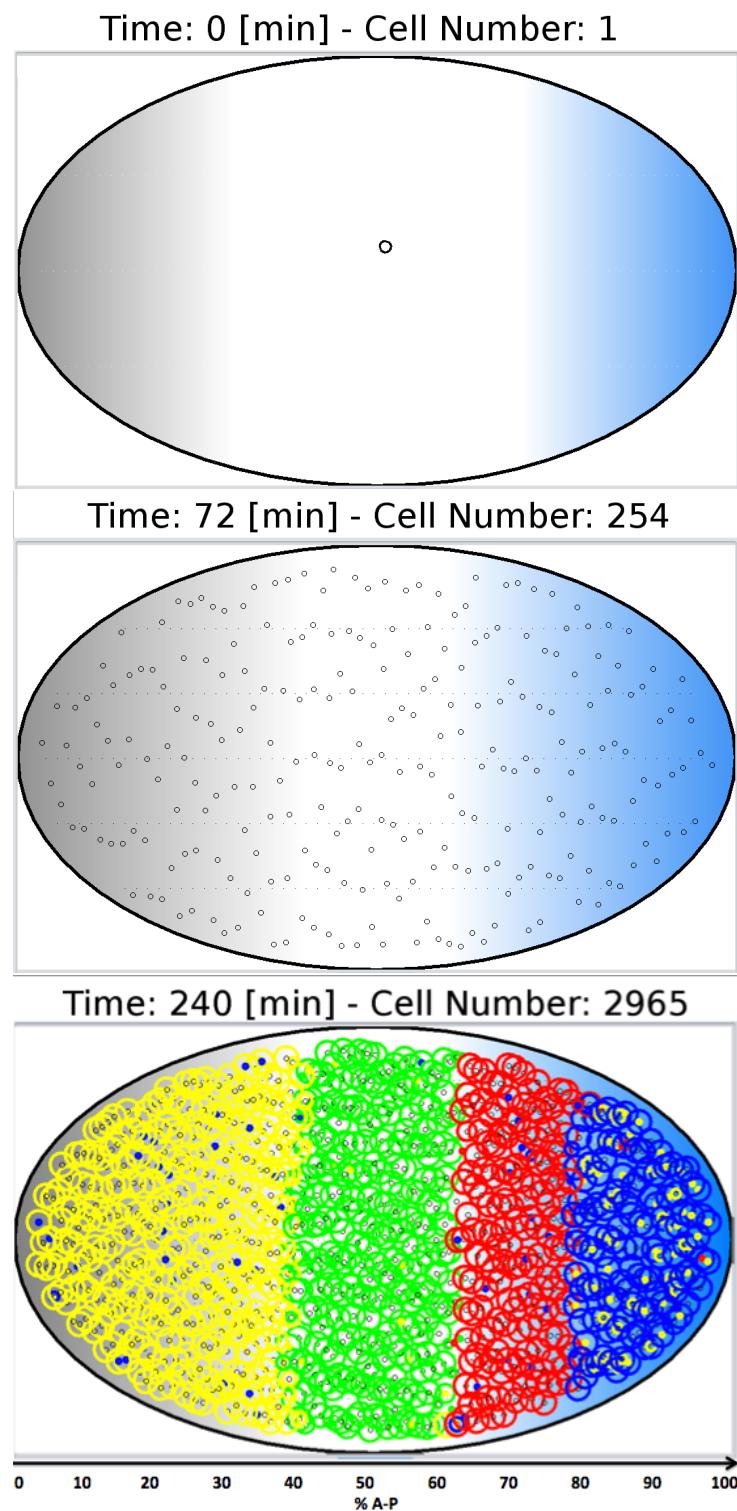


Figure 5.12: Simulation results for the four gap genes *hb* (yellow), *kni* (red), *gt* (blue), *Kr* (green) at a simulation time equivalent to the eighth time step of cleavage cycle 14A

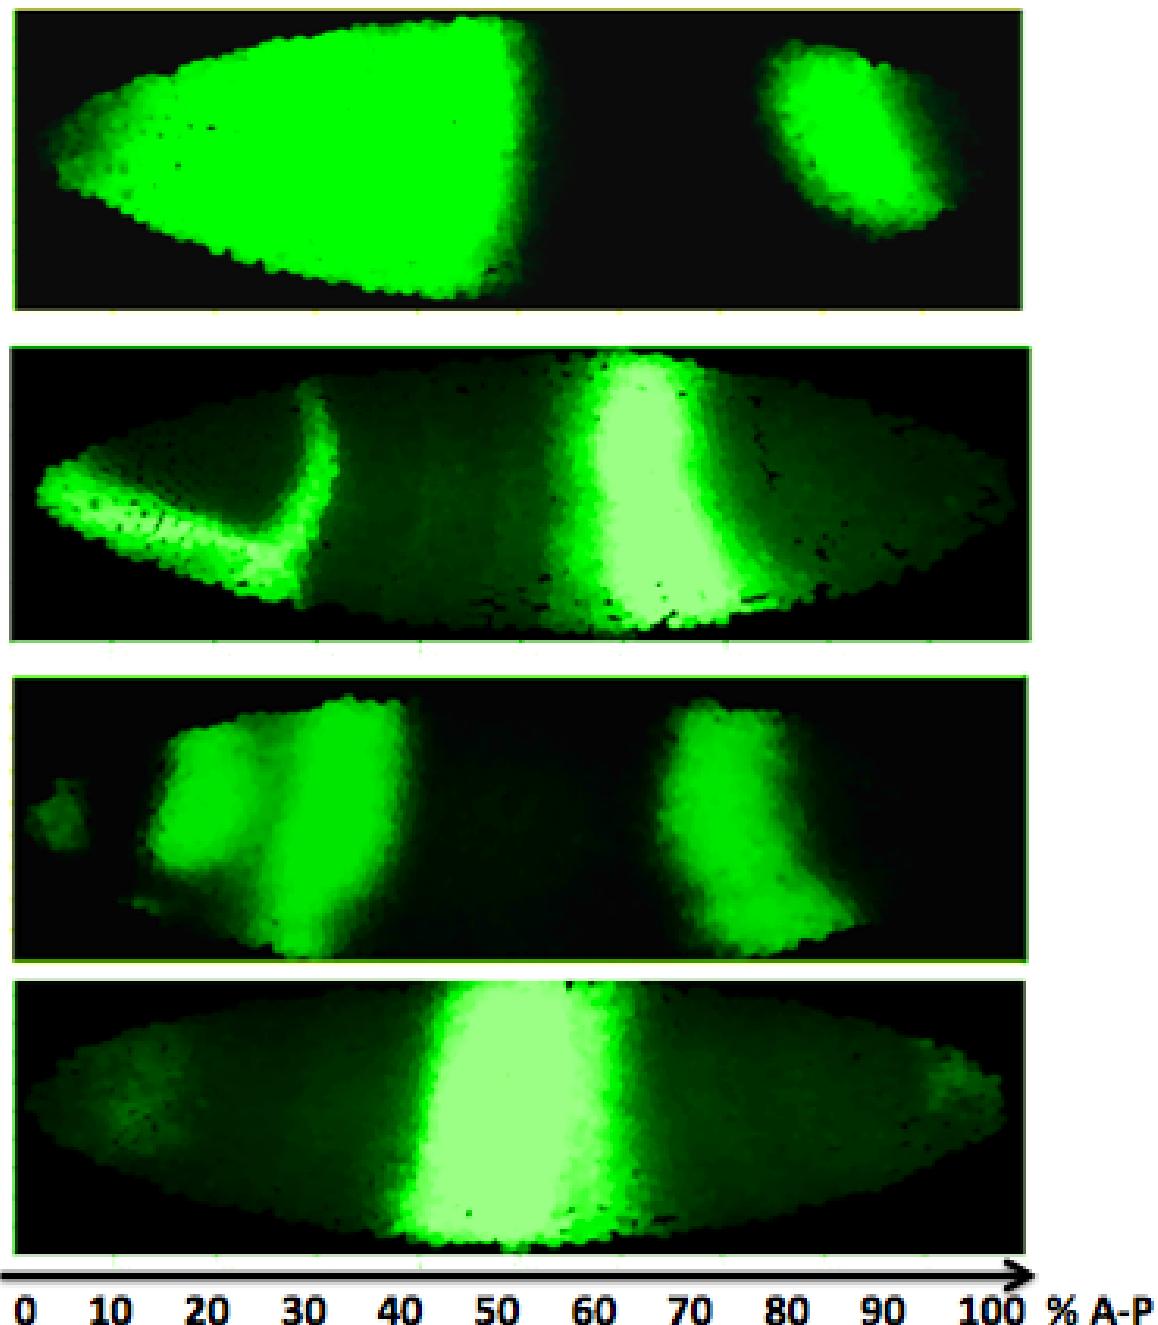


Figure 5.13: The experimental data for the expression of (from the top) *hb*, *kni*, *gt*, *Kr* at the eighth time step of cleavage cycle 14A [SKK<sup>+</sup>08], ©Maria Samsonova and John Reinitz

## **5.8 SAPERE meta-model**



# 6

## Self-organisation patterns

SAPERE

### 6.1 Crowd evacuation

### 6.2 Crowd steering

In this section, we show a crowd steering case study in which a middleware has the goal of leading people in the desired location within a complex environment in short time, avoiding obstacles such as crowded regions and without global supervisioning.

#### 6.2.1 Reference scenario

Consider a museum with a set of rooms, whose floor is covered with a network of computational devices (infrastructure nodes). These devices can exchange information with each other based on proximity, sense the presence of visitors, and hold information about expositions currently active in the museum. Each room has four exits and they are connected via external corridors. Visitors wandering the museum are equipped with a hand-held device that holds the visitor's preferences. By interaction with infrastructure nodes, a visitor can be guided towards rooms with a target matching their interest, thanks to signs dynamically appearing on his smartphone or on public displays. This is done using techniques suggested in the field of spatial computing [VCMZ11]—namely, computational gradients injected in a source and diffusing around such that each node holds the minimum distance value from the source.

The environment is a continuous bidimensional space with walls. Smartphones (or public displays) are agents dynamically linked with the nearest infrastructure node – the neighbours are the sensors inside a certain radius  $r$ , parameter of the model – from which they can retrieve data in order to suggest visitors where to go. Visitors are agents which follow the advices of their hand-held device (or public displays). It is defined a minimum possible distance between them, so as to model the physical limit and the fact that two visitors can't be in the same place at the

same time. Visitors can move of fixed size steps inside the environment. If an obstacle is on their path, their movement is shortened to the allowed position nearest to the desired place.

### 6.2.2 Steering strategy

All the information exchanged is in form of annotations, simply modelled as tuples  $\langle v_1, \dots, v_n \rangle$  (ordered sequence) of typed values, which could be for example numbers, strings, structured types, or function names. There are three forms of annotations used in this scenario:

$$\begin{aligned} & \langle \text{source}, id, type, N_{max} \rangle \\ & \langle \text{field}, id, type, value, tstamp \rangle \\ & \langle \text{info}, id, crowd, M, t' \rangle \end{aligned}$$

A **source** annotation is used as a source with the goal of generating a field:  $id$  labels the source so as to distinguish sources of the same type;  $type$  indicates the type of fields in order to distinguish different expositions;  $N_{max}$  is the field's maximum value. A **field** annotation is used for individual values in a gradient:  $value$  indicates the individual value; the  $tstamp$  reflects the time of creation of the annotation; the other parameters are like in the source annotation. An **info** annotation is supposed to be created and kept up to date by each sensor.  $M$  represents the number of smartphones the sensor is perceiving as neighbours.

The rules are expressed in form of chemical-resembling laws, working over patterns of annotations. One such pattern  $P$  is basically an annotation which may have some variable in place of one or more arguments of a tuple, and an annotation  $L$  is matched to the pattern  $P$  if there exists a substitution of variables which applied to  $P$  gives  $L$ . A law is hence of the kind  $P_1, \dots, P_n \xrightarrow{r} P'_1, \dots, P'_m$ , where: (i) the left-hand side (reagents) specifies patterns that should match annotations  $L_1, \dots, L_n$  to be extracted from the local annotation space; (ii) the right-hand side (products) specifies patterns of annotations which are accordingly to be inserted back in the space (after applying substitutions found when extracting reagents, as in standard logic-based rule approaches); and (iii) rate  $r$  is a numerical positive value indicating the average frequency at which the law is to be fired—namely, we model execution of the law as a CTMC transition with Markovian rate (average frequency)  $r$ . If no rate is given the reaction is meant to be executed “as soon as possible”, which means that the rate that associated with the reaction tends to infinite. To allow interaction between different nodes (hence, annotation spaces), we introduce the concept of *remote pattern*, written  $+P$ , which is a pattern that will be matched with an annotation occurring in a neighbouring space.

As sources annotations are injected in nodes, gradients are built by the first three rules in Figure 6.1. The first one, given a source, initiates the field with its possible maximum value. The second one, when a node contains a field annotation, spreads a copy of it to a neighbouring node picked up randomly with a new value computed considering the crowding around the sensor. The parameter  $k$  allows to tune how much crowding should influence the field, while  $\#D$  is the measure of the distance between the two involved sensors. As a consequence of these laws,

$$\begin{aligned}
\langle \text{source}, id, type, N_{max} \rangle &\xrightarrow{r_{init}} \langle \text{source}, id, type, N_{max} \rangle, \langle \text{field}, id, type, N_{max}, \#T \rangle \\
\langle \text{field}, id, type, N, t \rangle &\xrightarrow{r_{diff}} \langle \text{field}, id, type, N, t \rangle, +\langle \text{pre\_field}, id, type, N - \#D, t \rangle \\
\langle \text{pre\_field}, id, type, N, t \rangle, \langle \text{info}, id, crowd, M, t' \rangle &\mapsto \langle \text{field}, id, type, N - k * C, t \rangle, \langle \text{info}, id, crowd, M, t' \rangle \\
\langle \text{field}, id, type, N, t \rangle, \langle \text{field}, id, type, M, t' + t \rangle &\mapsto \langle \text{field}, id, type, M, t' + t \rangle \\
\langle \text{field}, id, type, N, t \rangle, \langle \text{field}, id, type, M, t \rangle &\mapsto \langle \text{field}, id, type, \max(M, N), t \rangle \\
\langle \text{field}, id, type, N, t \rangle, \langle \text{field}, id', type, N + M, t' \rangle &\mapsto \langle \text{field}, id', type, N + M, t' \rangle
\end{aligned}$$

Figure 6.1: Laws describing the museum application.

each node will carry a field annotation indicating the topological distance from the source. The closest is the field value to  $N_{max}$ , the nearest is the field source. When the spread values reach the minimum value 0, the gradient has to become a plateau.

To address the dynamism of the scenario where people move, targets being possibly shifted, and crowds forming and dissolving, we introduced the following mechanism. We expect that if a gradient source moves the diffused value has to change according to the new position. This is the purpose of the *tstamp* parameter which is used in the fourth law, continuously updating old values by more recent ones (*youngest* law). In this way we ensure that the system is able to adapt to changes of the source states. Finally, the spreading law above will produce duplicate values in locations, due to multiple sources of the same type (indicated by different ids), multiple paths to a source, or even diffusion of multiple annotations over time. For this reason we introduced the last two laws. They retain only the maximum value, i.e. the minimum distance, the former when there are two identical annotations with only a different value, the latter when the id is different (*shortest* laws). The proposed solution is intrinsically able to dynamically adapt to unexpected events (like node failures, network isolation, crowd formation, and so on) while maintaining its functionality.

People are modelled as nodes programmed with a single reaction with no conditions and having, as action, the ability to follow the highest field value among neighbouring sensors. For the crowding annotation, we may assume that sensors are calibrated so as to locally inject and periodically update it by setting the current level of crowding, *i.e.* the number of persons. The reaction rates are identified by hand performing different simulations with different parameters. The results reported are obtained with  $r_{init} = 1$  and  $r_{diff} = 50$ . The other laws show no rate because it is assumed to be infinite (as-soon-as-possible semantics).

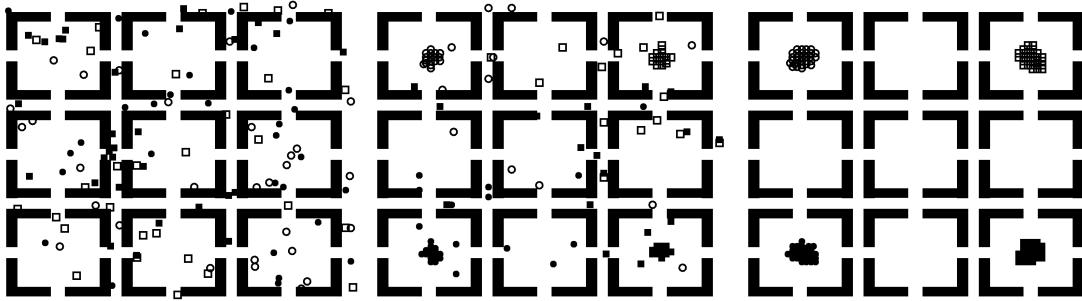


Figure 6.2: A simulation run of the reference exposition: three snapshots of ALCHEMIST’s graphic reporting module with this simulation.

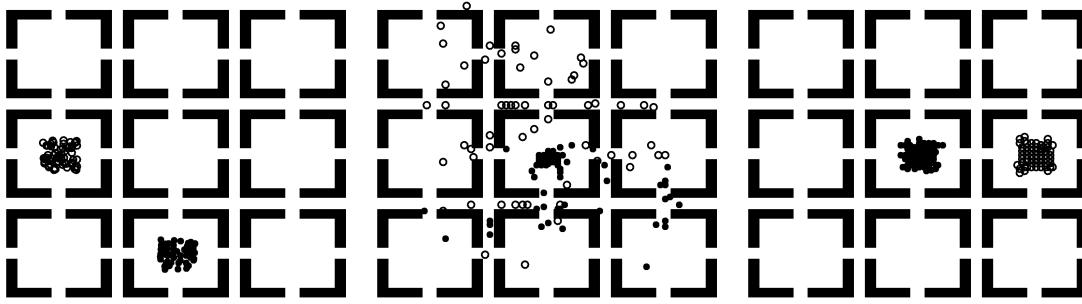


Figure 6.3: A run showing the effect of crowding: dark visitors occupy a central room, making other visitors moving left to right by a longer, less crowded path

### 6.2.3 Simulation in ALCHEMIST

We here present simulations conducted over an exposition, where nine rooms are connected via corridors. People can express different preferences represented by their shape.

Three snapshots of a first simulation run are reported in Figure 6.2. We here consider four different targets that are located in the four rooms near environment angles. People are initially spread randomly in the museum, as shown in the first snapshot, and they eventually reach the room in which the desired target is hosted, as shown in the last snapshot.

Figure 6.3 shows a simulation experimenting with the effect of crowding in the movement of people. Two groups of people – denoted with empty and full circles – with common interests are initially located in two different rooms, as shown in the first snapshot. The target for the dark visitors is located in the central room of the second row, while the others’ is in the right room of the second row. In the simulation, dark visitors reach their target soon because it is nearer, thus forming a crowded area intersecting the shortest path towards the target for the other visitors. Due to this jam the latter visitors choose a different path that is longer but less crowded.

Both tests show qualitative effectiveness of the proposed laws, and suggest that our simulation

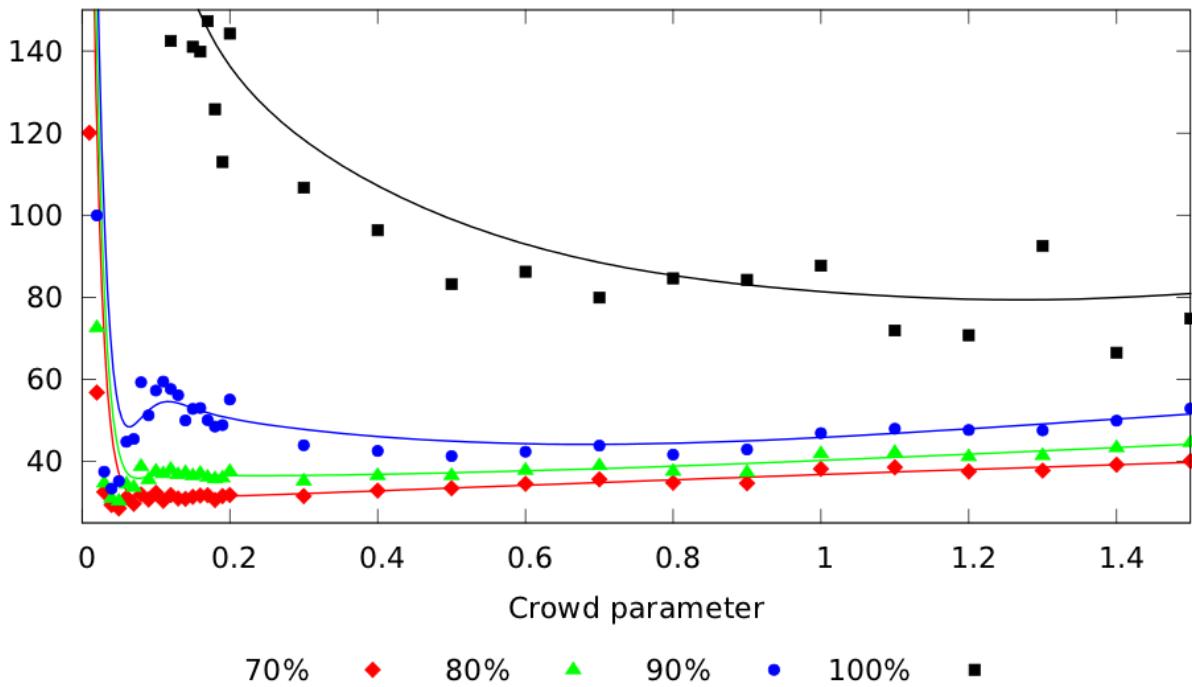


Figure 6.4: Time units of convergence time with different values of crowd parameter and different percentages of people

approach can be used for additional experiments focussing on tuning system parameters (factor  $k$ ) or alternative strategies (e.g., diffusing crowd information) to optimise paths to destinations. For instance, in the context of the second case, Figure 6.4 shows how factor  $k$  can influence the time for (sub)groups of (light) people to reach the destination, by which we can see that even small values of  $k$  lead to a significant improvement—which slowly decreases as  $k$  grows.

## **6.3 Self composition**

## **6.4 Semantic resource discovery**

## **6.5 Anticipative adaptation**

## **6.6 Crowd disasters prediction**

## **6.7 Crowd steering at the urban scale**

In section, we focus on the problem of crowd steering at the urban scale. The problem is noticeably different from the classic per-user navigation: we want to consider the current position of people to make the user avoid congested areas, thus reducing the overall trip time and improving the security level. Also, this differs from the experiment in Section 6.2 for the complexity of the environment and the number of devices involved. Crowds of people vary with time, and consequently we need a system able to dynamically adapt to such changes. As per our line of research, we obviously want the computation to happen in a completely distributed fashion, with no centralised computing system involved.

### **6.7.1 Devices and physical configuration**

We suppose users to be equipped with smart devices, also able to communicate with other devices within a certain range with a wireless technology, and must be able to be aware of their current position.

We also suppose the organisers to have spread around the city some “static” network nodes, for instance on public illumination poles, traffic lights, or signals. Such nodes must be equipped with network capabilities similar to those of a Wi-Fi access point. In particular, they must allow nearby mobile devices to connect to all the static nodes within their communication range. There is no upper or lower number or density (devices per square meter) limit to the number of static devices to deploy. The only strict requirement involves both the distance among devices and their communication range: the devices must be placed at a distance and have a communication range large enough such as there is no network segmentation. Besides strict requirements, some other guidelines apply:

- it is preferable to displace static nodes in points where there are people, e.g. on streets and crossings;
- the wider the number of static nodes that can communicate, the more precise will be the results;
- the higher the density of static devices, the more precise will be the results.

The last assumption we make is that devices are able to estimate the time needed for walking towards any of the connected device in normal conditions. This last assumption is again perfectly acceptable considering the current technology, where smart devices are always connected to the Internet and can access public navigation services.

### 6.7.2 Distributed crowd steering

Given the structure described above and the mobile smart devices able to communicate with the static nodes, we now outline a possible software solution. The device wanting to be steered publishes a gradient including the information about the wanted destination. The information about destination can be either expressed as physical location (e.g. latitude and longitude) or as a description: the system works in a fashion totally orthogonal to this choice. One or more static devices located near a potential point of interest (POI) receive such gradient, and react becoming sources of a gradient which measures the distances between nodes. This response gradient diffuses on static nodes until it reaches the source, carrying within itself information about the chain of static devices forming the shortest path. The requester will receive a gradient pointing towards the nearest POI, and can navigate step-by-step by reaching in order all the nodes of the path. Although it is not part of this very experiment, the reply could be tuned to be different depending on the request and on the characteristics of the point of interest. For instance, a value representing the “level of matching” between the request and the response could be embedded in the gradient, and used to modify its spatial structure in such a way that POIs more affine with the request are preferred even if farther. Some example of similar usage of gradients is available in [SYD<sup>+</sup>13].

Now that a basic form of steering is in place, we can improve it by adding contextual information. We assume that static nodes are able to detect the number of people in their surroundings. Again, this is not critical even in a real deployment, and can be achieved in numerous ways with a different level of precision, which spans from just keeping track of which mobile nodes are connected (and, as a consequence, within the communication range) to the usage of dedicated sensors and techniques (e.g. cameras and computer vision). Let’s call  $C$  the perceived number of mobile devices surrounding a static node. We can alter the spatial shape of the gradient by acting on the function that outputs the actual distance between devices. For instance, adopting the notation in Section 2.3.2, we can use  $f(n) = \Gamma_n + d(n) + K \cdot C$  where  $K$  is a system parameter. What happens? Basically, areas densely populated with mobile devices will appear as more distant, and consequently the requester will be steered towards alternative routes. The whole system works in a totally emergent and distributed way, with no global knowledge of what’s going on, no central control involved, and complete distribution of the computation among its components.

A further refinement of these steering mechanisms involve the ability to react to events that will happen in future (given that there is information about that). The patterns required to do so have already been presented in Section 6.5, and the interested reader can deepen reading [MPV12].

### 6.7.3 Simulation in ALCHEMIST

We decided to focus on the case of mass urban sports events, and in particular on the Vienna City Marathon 2013, an event that every year involves about 40.000 actives and 300.000 spectators. During such event, a smartphone application based on SAPERE [ZCF<sup>+</sup>11] concepts was deployed, and gathered 1503 high quality GPS traces [APNF13]. We relied on such data set to build simulations demonstrating our concepts of crowd steering at a urban level.

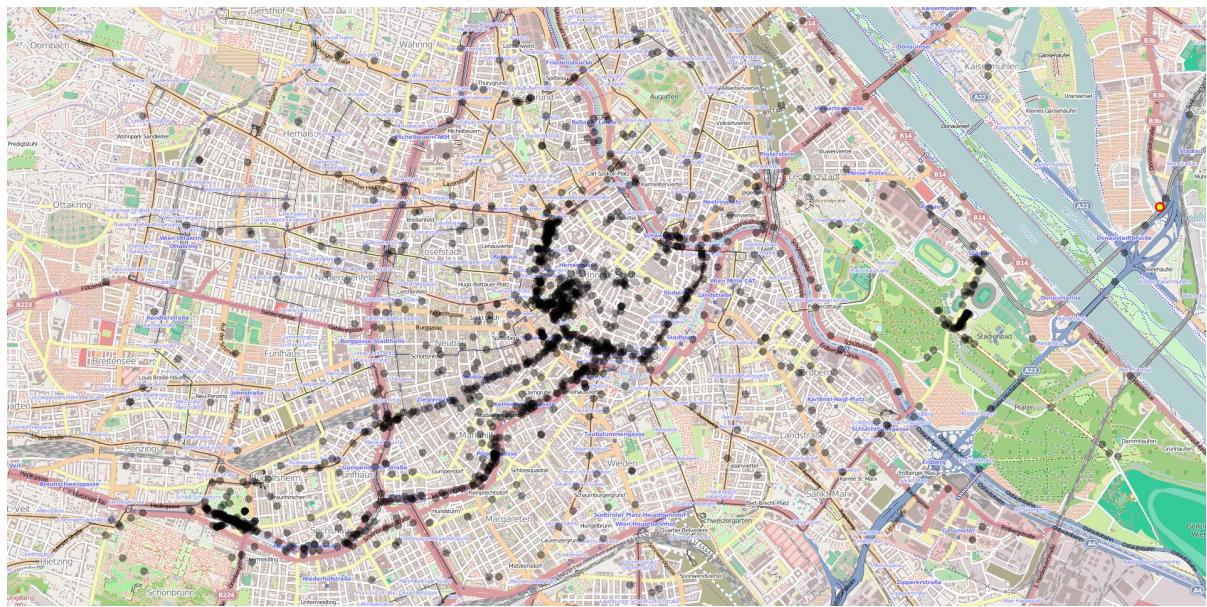


Figure 6.5: A snapshot of the whole city of Vienna as simulated in ALCHEMIST. This snapshot is taken while simulating the city at 10am, each black point corresponds to a GPS trace. The more an area is crowded, the blacker it appears in the image.

As ALCHEMIST environment, we obviously used the map of Vienna. Another rather straightforward choice was to map each device to a node, supposing the users to carry one and one only device which participates to the system. The network of static nodes was built by creating a grid of 572 devices spread around the city and enabling the positioning in the nearest street point. After that, we positioned a node at Burggarten (48.203926,16.365765), representing our POI, and a node at (48.21441,16.35825) representing our navigation system user. Finally, we positioned the nodes of the users which follow the traces. As result, we had in total 2077 simulated nodes, of which 1504 mobile and moving during the actual simulation, 1503 by following the traces using the mixed mode (see Figure 5.8(c)) and one following the steering system.

As linking rule we decided to let the infrastructure nodes connect to every other node within a communication range of 150 meters. We used the same set of reactions for each node in the system, but on the mobile nodes, in which also the reactions needed to program the mobility were included. In Figure 6.5 a graphical output of the simulator is proposed. We relied on the

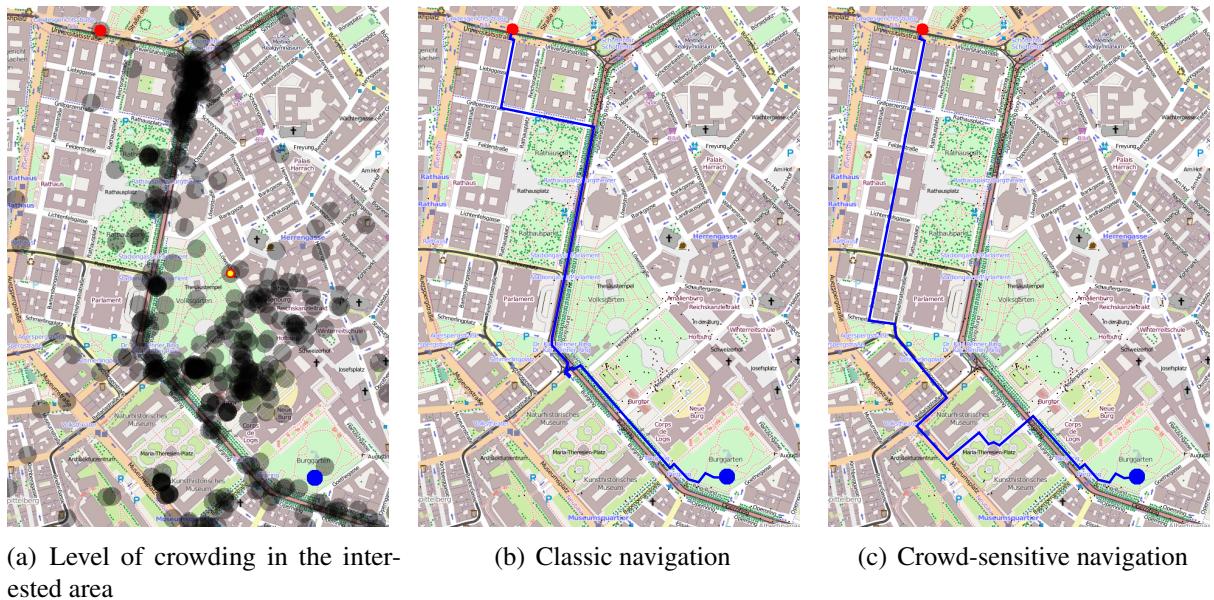


Figure 6.6: In these three snapshots a qualitative evaluation of the crowd-sensitive steering system benefits is offered. The requester is painted in red, while the nearest POI is in blue. In Figure 6.6(a), the density of people in the area is shown (at the time at which the red node starts): a more intense black mean a more crowded area. In Figure 6.6(b) and Figure 6.6(c) is possible to see the suggested routes of, respectively, the classic navigation and the crowd-sensitive navigation. The second suggests a longer but much less jammed path.

SAPERE meta-model. This way, we were able to simulate a network of programmable tuple spaces.

A complete evaluation of such an emergent environment is all but trivial. The best way to measure the effectiveness of a crowd-sensitive user steering system would probably be the measure of the average walking time on routes generated probabilistically according to the most walked paths. This kind of evaluation has two strong requirements to prove itself scientifically relevant: first, it requires to identify how popular are the routes; second, and most important, it would require a realistic model of pedestrian, including the physical interaction with other people. At the time of writing, we do not have such model, although we are working to find a decent approximation.

As proof of concept, we chose to monitor the path of a single user steered from Universitätstrasse to Burggarten. We chose this path because the shortest path connecting the user to her destination includes a walk in between one of the most crowded areas during the sport event, and as such is the perfect test bed for an approach whose goal is to improve the trip time in such conditions.

In Figure 6.6 a qualitative impact is given. It is immediately clear that the suggested path in Figure 6.6(a) is longer but successfully reduces the space walked within crowded areas, with

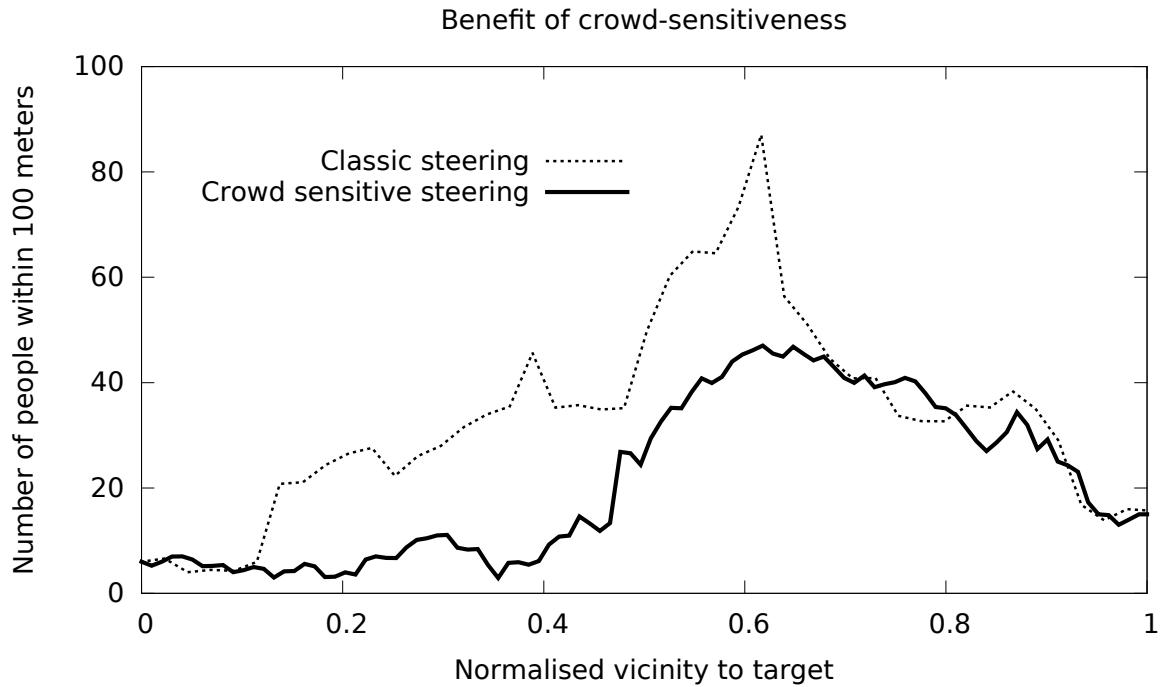


Figure 6.7: This chart shows how the number of users surrounding the user (within 100 meters from her) vary with the proximity to the target. With “normalised distance”, we mean that we divided the distance the user still has to walk to reach the target by the total length of the suggested path.

respect to the one in Figure 6.6(b).

The chart in Figure 6.7 confirms the qualitative evaluation. Towards the beginning and the end, the two algorithms have a similar performance, due to the fact that these parts of the walk (as the initial) are common. In the central part of the walk, however, the path suggested by the crowd-sensitive algorithm tends to avoid a considerable amount of jammed areas. We claim that, assuming the traces to be a representative sampling of the actual population of the area and the more crowded areas to be slower to walk (both assumptions sound rather straightforward to the authors), then the crowd sensitive algorithm guarantees an average lower time to destination.

# 7

## Aggregate programming languages

When we introduced Field Calculus in Section 2.5.2, we said that it is more a theoretical than a practical framework. Building on the foundation of field calculus, the subsequent step is to provide a framework that can be practically leveraged to express a collection of higher-level “building block” algorithms, each a simple and generalized basis element of an “algebra” of programs with desirable resilience properties (e.g., the operators presented in [BV14]). On top of this, higher-level library APIs can be built, enabling simple and transparent construction of robust distributed systems.

Any practical implementation of field calculus must embed a field calculus interpreter within an architecture that handles the pragmatics of communication, execution, and interfacing with hardware, operating system, and other software. At the same time, it is important that this system be readily portable across both simulation environments and real networked devices. Finally, both system development and maintainability are greatly enhanced if the exact same code is used for execution in all contexts.

Here is where Protelis comes into play: it represents our effort to bring Field Calculus to practice, satisfying the requirements stated above.

### 7.1 Higher order functions in field calculus

#### 7.1.1 Impact on alignment

### 7.2 Protelis: practical aggregate programming

We have designed the Protelis language as an implementation of the field calculus [VDB13] closely related to Proto [BB06]. On the one hand, it incorporates the main spatial computing features of the field calculus, hence enjoying its universality, consistency, and self-stabilization properties [BVD14, VD14]. On the other hand, it turns the field calculus into a modern specification language, improving over Proto by providing

- access to a richer API through Java integration;
- support for code mobility through first-order functions;

$P ::= \bar{I} \bar{F} \bar{S};$	;; Program
$I ::= \text{import } m \mid \text{import } m.*$	;; Java import
$F ::= \text{def } f(\bar{x}) \{ \bar{s}; \}$	;; Function definition
$s ::= e \mid \text{let } x = e \mid x = e$	;; Statement
$w ::= x \mid l \mid [\bar{w}] \mid f \mid (\bar{x}) \rightarrow \{ \bar{s}; \}$	;; Variable/Value
$e ::= w$	;; Expression
$\mid b(\bar{e}) \mid f(\bar{e}) \mid e.\text{apply}(\bar{e})$	;; Fun/Op Calls
$\mid e.m(\bar{e}) \mid \#a(\bar{e})$	;; Method Calls
$\mid \text{rep } (x < -w) \{ \bar{s}; \}$	;; Persistent state
$\mid \text{if}(e) \{ \bar{s}; \} \text{ else } \{ \bar{s}'; \}$	;; Exclusive branch
$\mid \text{mux}(e) \{ \bar{s}; \} \text{ else } \{ \bar{s}'; \}$	;; Inclusive branch
$\mid \text{nbr} \{ \bar{s}; \}$	;; Neighborhood values

Figure 7.1: Protelis abstract syntax, colored to emphasize definition and application (red), functions (blue), variables (green), and special field calculus operators (purple).

- a novel syntax inspired by the more widely adopted C-family languages.

Protelis is freely available and open source, and can be downloaded as part of the ALCHEMIST distribution.

### 7.2.1 Syntax

We present the Protelis language in terms of its abstract syntax, provided in Figure 7.1 as a means to guide the discussion of the language’s features. This syntax uses similar conventions to well-known core languages like Featherweight Java [IPW01]. We let meta-variable  $f$  range over names of user-defined functions,  $x$  over names of variables and function arguments,  $l$  over literal values (Booleans, numbers, strings),  $b$  over names of built-in functions and operators (including the “hood” functions described in Section 7.2.3),  $m$  over Java method names, and  $\#a$  over aliases of static Java methods. All such meta-variables are used as non-terminal symbols in Figure 7.1. Overbar notation  $\bar{y}$  generally means a comma-separated list  $y_1, \dots, y_n$  of elements of kind  $y$ , with the two exceptions that in  $\bar{F}$  we use no comma separator, and in  $\bar{s};$  semi-colon is used as separator instead.

### 7.2.2 Ordinary Language Features

One of the distinctive elements of Protelis when compared to other aggregate programming languages (particularly Proto), is the adoption of a familiar C- or Java-like syntax, which

can significantly reduce barriers to adoption. Despite this syntactic similarity, Protelis is a purely functional language: a program is made of a sequence of function definitions ( $F_1 \dots F_n$ ), modularly specifying reusable parts of system behavior, followed by a main block of statements. Following the style of C-family languages, a function's body is a sequence of statements surrounded by curly brackets. As in the Scala programming language<sup>1</sup>, however, statements can also be just expressions, and a statement sequence evaluates to the result of the last statement. Each statement is an expression to be evaluated ( $e$ ), possibly in the context of the creation of a new variable (`let x = e`) or a re-assignment (`x = e`)<sup>2</sup>. As an example, consider the following function taking four fields as parameters, after the “channel” pattern from [But02]:

```
def channel(distA, distB, distAB, width) {
    let d = distA + distB;
    d = d - distAB;
    d < width
}
```

This function assumes that its input `distA` maps each device to its distance to a region  $A$ , `distB` maps each device to its distance to a region  $B$ , `distAB` is a constant field holding at each device the minimum distance between regions  $A$  and  $B$ , and `width` is a constant field holding at each device the same positive number. The function then computes a Boolean field, mapping each device to `true` only if it belongs to a “channel” area around the shortest path connecting regions  $A$  and  $B$  and approximately `width` units wide. All devices elsewhere map to `false`.

Atomic expressions  $w$  can be literal values (`l`), variables (`x`), tuples (`[w]`), function names (`f`) or lambdas (`(x) -> {s;}`). Structured expressions include three kinds of “calls”: (i)  $b(\bar{e})$  is application to arguments  $\bar{e}$  of a built-in operation  $b$ , which could be any (infix- or prefix-style) mathematical, logical or purely algorithmic function<sup>3</sup>; (ii) `f(\bar{e})` is application of a user-defined function; and (iii) `e.apply(\bar{e})` is application of arguments to an expression  $e$  evaluating to a lambda or function name. The following shows examples of such calls:

```
def square(x) {
    x * x;
}
let f = square;
let g = (x) -> {
    square(x) + 1
};
f.apply(g.apply(2)) // gives 25 on all devices
```

<sup>1</sup><http://www.scala-lang.org>.

<sup>2</sup>Technically, “re-assignment” is actually the creation of a new variable that shadows the old.

<sup>3</sup>For simplicity of presentation, we omit the syntax for infix operations and order of operations, which is closely patterned after Java.

In addition, arbitrary Java method calls can be imported and used by Protelis: (i) `e.m(Ē)` is method call on object `e` and (ii) `#a(Ē)` is invocation of a static method, via an alias `#a` (always starting with '#') defined by an `import` clause. The alias is created automatically as the bare method name for single imports or imports of all methods in a class with `*`. Protelis can thus interact with Java reflection to support dynamic invocation of arbitrary Java code, as shown in the following example:

```
import java.lang.Class.forName
let c = #forName("String");
let m = c.getMethod("length");
m.invoke("Lorem ipsum dolor sit amet")// gives 26 on all devices
```

### 7.2.3 Special Field Calculus Operators

The remaining constructs of Protelis are the special operations specific to field calculus, dealing with the movement of information across space and time:

- Construct `rep (x<-w) {S;}` defines a locally-visible variable `x` initialized with `w` and updated at each computation round with the result of executing body `{S;}`: it provides a means to define a field evolving over time according to the update policy specified by `{S;}`.
- Construct `nbr{S;}` executed in a device gathers a map (actually, a field) from all neighbors (including the device itself) to their latest value from computing `S`. A special set of built-in “hood” functions can then be used to summarize such maps back to ordinary expressions. For example, `minHood` finds the minimum value in the map.
- The branching constructs `mux(e){S;} else {S';}` and `if(e){S;} else {S';}` perform two critically different forms of branching. The `mux` construct is an inclusive “multiplexing” branch: the two fields obtained by computing `S` and `S'` are superimposed, using the former where `e` evaluates to `true`, and the second where `e` evaluates to `false`. Complementarily, `if` performs an exclusive branch: it partitions the network into two regions: where `e` evaluates to `true` `S` is computed, and elsewhere `S'` is computed instead.

The following code shows some example uses of these constructs:

```

def count() {
    rep(x<-0) { x + 1 }
}

def maxh(field) {
    maxHood(nbr{field})
}

def distanceTo(source) {
    rep(d <- Infinity) {
        mux (source) {
            0
        } else {
            minHood(nbr{d} + nbrRange)
        }
    }
}

def distanceToWithObstacle(source,obstacle) {
    if (obstacle) {
        Infinity
    } else {
        distanceTo(source)
    }
}

```

Function `count` yields an evolving field, counting how many computation rounds have been executed in each device. Function `maxh` yields a field mapping each device the maximum value of `field` across its neighborhood—note a `nbr` construct should always be eventually wrapped inside a “hood” function. Function `distanceTo` nests `nbr` inside `rep` to create a chain of interactions across many hops in the network, computing minimum distance from any device to the nearest “source device” (i.e., where `source` holds true). It does so by a field `d` initially `Infinity` everywhere, and evolving as follows: `d` is set to 0 on sources by `mux`, and elsewhere takes the minimum across neighbors of the values obtained by adding to `d` the estimated distance to the current device—a triangle inequality relaxation computing a distance field also often termed *gradient* [LK87, BBVT08, VCMZ11]. Finally, function `distanceToWithObstacle` shows exclusive branch at work; `distanceTo(source)` is computed in the sub-region where there is no obstacle, which causes the computation of distances to implicitly circumvent such obstacles.

### 7.3 Architecture of Protelis

In Protelis, we designed an architecture, subsumed in Figure 7.2(a), following the same general pattern as was used for the Proto Virtual Machine [BB07]. First, a parser translates a text Protelis program into a valid representation of field calculus semantics. This is then executed by a Protelis interpreter at regular intervals, communicating with other devices and drawing contextual information from environment variables implemented as a tuple store of  $(token, value)$  pairs. This abstraction is instantiated for use on particular devices or simulations by setting when executions occur, how communication is implemented and the contents of the environment.

We have chosen to implement this architecture in Java. One key reason for this choice is that Java is highly portable across systems and devices. Another key reason (discussed further in the next section) is that Java’s reflection mechanisms make it easy to import a large variety of useful libraries and APIs for use in Protelis. Finally, the pragmatics of execution on embedded devices have also changed significantly since the publication of [BB07]: a much wider variety of low cost embedded devices are now capable of supporting Java, while at the same time improvements in Java implementations have made it much more competitive in speed and resource cost with low-level languages like C [BSB<sup>+</sup>03, ORAI11].

In particular, we have chosen to implement Protelis and its architecture via the Xtext language generator [EB10] and within the ALCHEMIST framework [PMV13]. Usefully, Xtext also features support for generating a language-specific Eclipse plug-in, which provides developer assistance through code highlighting, completion suggestions, and compile-time error detection.

For an initial validation, we have exercised this architecture by construction of two instantiations: one in the ALCHEMIST framework for simulation of large-scale spatially-embedded systems; the other as a daemon for coordinating management of networked services. Figure 7.2(b) shows the ALCHEMIST instantiation: simulations are configured using a simple scripting language, which specifies a Protelis program as well as the collection of devices that will execute it, communication between those devices, and other aspects of the environment to be simulated. The ALCHEMIST event-driven simulation engine then handles execution scheduling, message delivery, and updates to the environment tuple store. Figure 7.2(c) shows the network service management instantiation. Here, each Protelis device lives on a separate server in an enterprise network, and is tethered to the networked service it is intended to manage by a service manager daemon. This daemon monitors the service, injecting information about its status and known dependencies into the environment and maintaining a neighborhood by opening parallel communication links to the corresponding daemons on any other servers that the monitored service communicates with. Examples using each of these implementations are shown in Section 7.4.

### 7.4 Application examples

To demonstrate how these features combine to offer simple programming of complex distributed algorithms across a potentially broad range of applications domains, we now present two example

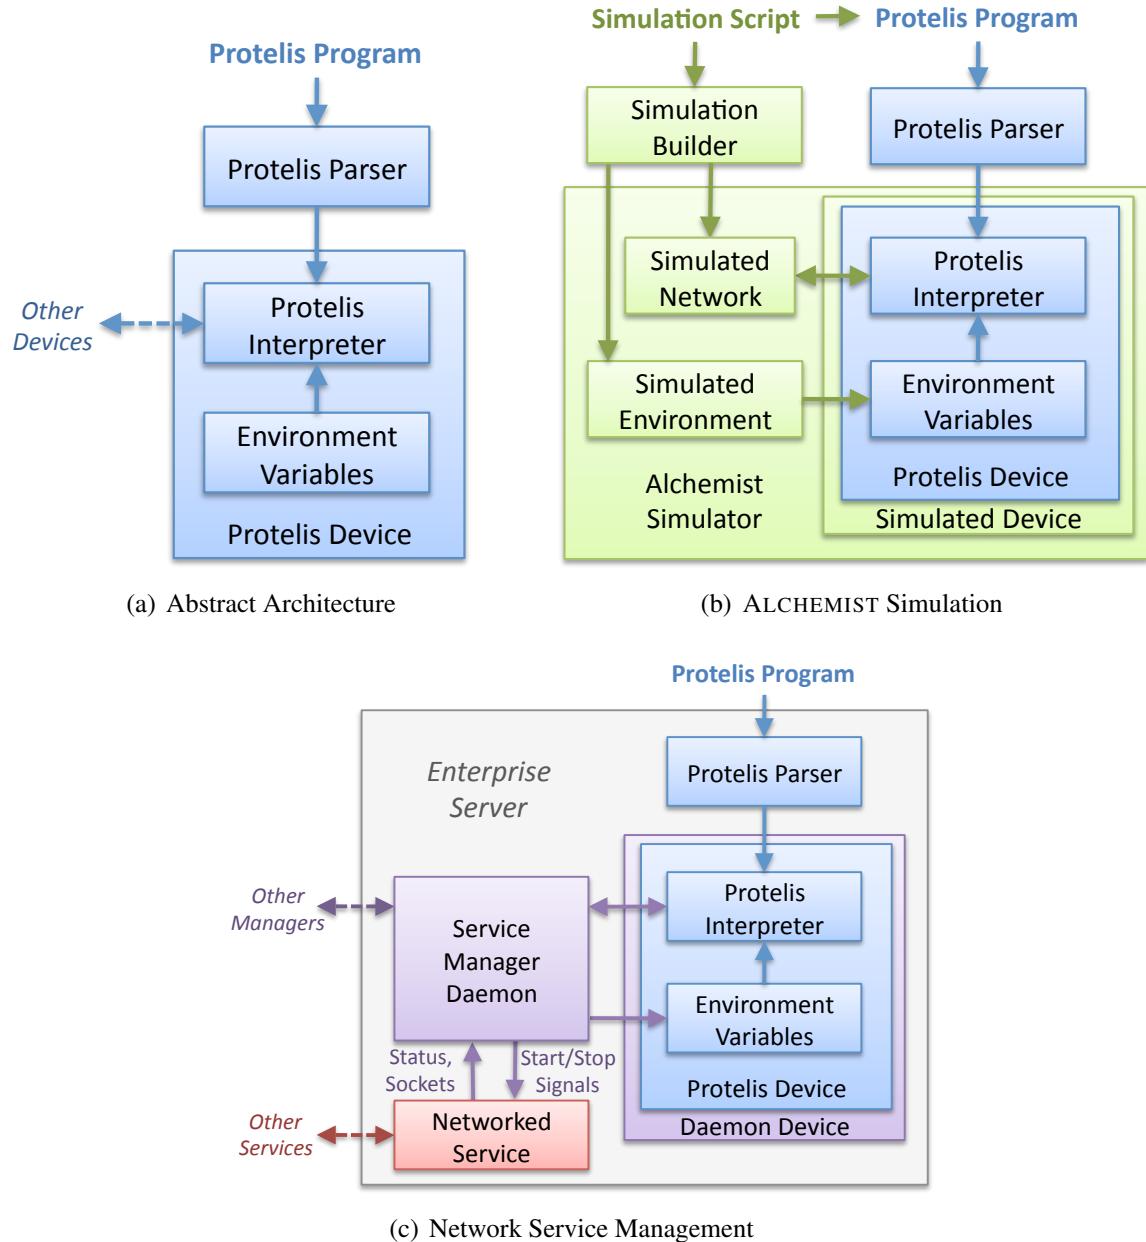


Figure 7.2: In the abstract Protelis architecture (a), an interpreter executes a pre-parsed Protelis program at regular intervals, communicating with other devices and drawing contextual information from a store of environment variables. This is instantiated by setting when executions occur, how communication is implemented and the contents of the environment. Two such instantiations are presented in this thesis: as a simulation in the ALCHEMIST framework (b) and as a daemon for coordinating management of networked services (c).

applications. The first aims at a pervasive computing scenario and is executed in simulation using Alchemist [PMV13], the second aims at enterprise network management and is executed on a collection of EmuLab [WLS<sup>+</sup>02] servers.

### 7.4.1 Rendezvous at a Mass Event

A common problem in large public events is to rendezvous with other companions attending the same large public event. At mass events, access to external cloud-based services may be difficult or impossible, and pre-arranged rendezvous points may be inaccessible or inconveniently distant. Simple peer-to-peer geometric calculations across the network, however, can readily compute a route that will allow two people to rendezvous:

```
// Follow the gradient of a potential field down from a source
def descend(source, potential) {
    rep(path <- source) {
        let nextStep = minHood(nbr([potential, self.getId()]));
        if (nextStep.size() > 1) {
            let candidates = nbr([nextStep.get(1), path]);
            source || anyHood([self.getId(), true] == candidates)
        } else {
            source
        }
    }
}
def rendezvous(person1, person2) {
    descend (person1 == owner, distanceTo(person2 == owner))
}
// Example of using rendezvous
rendezvous ("Alice", "Bob");
```

Figure 7.3 shows an example of running this rendezvous process in a simulated city center. We chose London as a simulation environment, using Alchemist's capability for importing OpenStreetMap data. We displaced 1000 devices randomly across the city streets (represented by pale blue dots), with a communication range of 475 meters (this range chosen to ensure no network segmentation). We then picked two devices whose owners want to meet: one device on Lambeth Bridge (lower left of the image) and one device on Tower Bridge (upper right), each marked with a yellow square. To mark the devices for Protelis, we injected their environments with a property `owner`, assigning the strings "Alice" and "Bob" as values for the first and the second device respectively.

Implementing this application requires only 21 lines of code: the listing above and the function `distanceTo` that can be found in Section 7.2.3. This implementation measures

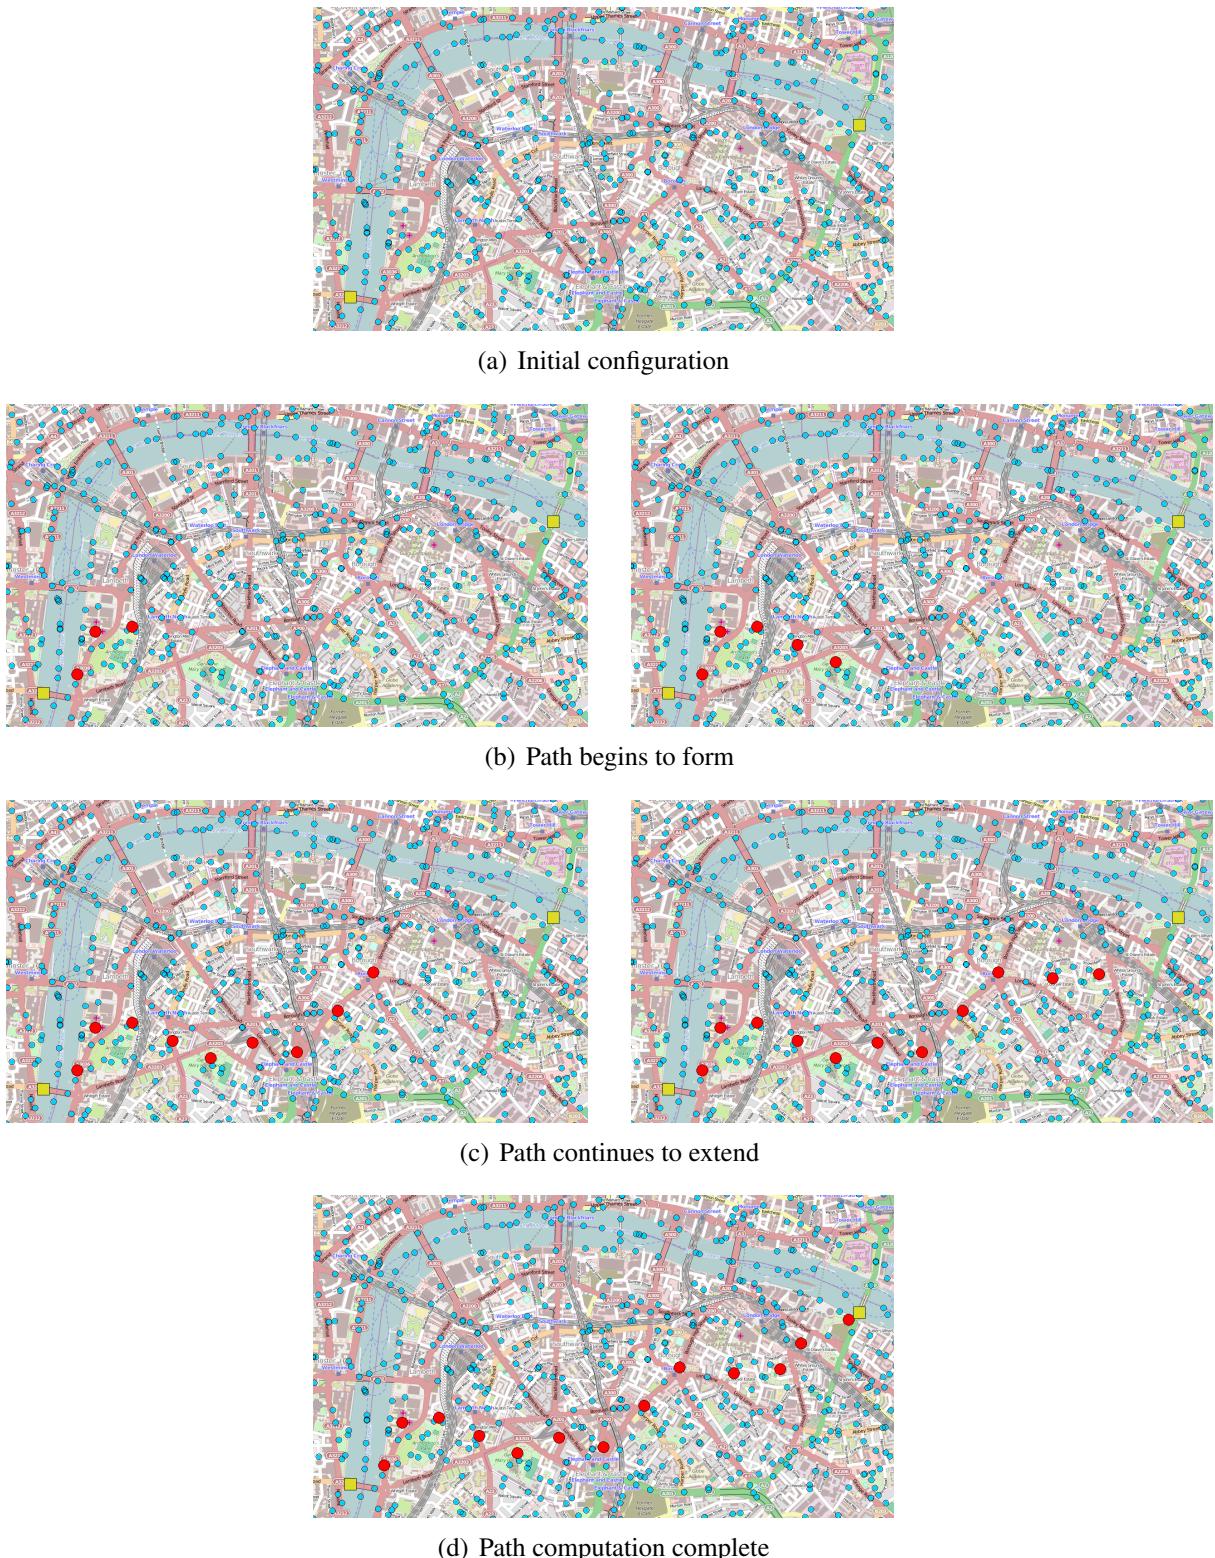


Figure 7.3: Example of computing a rendezvous route for two people in a crowded urban environment.

distance to one of the participants, creating a potential field, then, starting from the other one, builds an optimal path descending the distance potential field to return to the first participant at distance zero. The first half of the algorithm has already been described, and relies on `distanceTo`, while the second half is implemented by the function `descend`. This function, given a device and a potential field, builds a path of devices connecting the former with the source of the latter. The strategy is to mark the device we want to connect to the potential field's source as part of the path, and then, in every device, compute which of the neighbors is closest to the destination. Given this information, a device is in the path if one of the neighbors is in the path already and has marked this device as the closest of its neighbors towards the destination. Note how the whole algorithm can be elegantly compressed into just a few lines of code, and how there is no need to explicitly declare any communication protocol for exchanging the required information, thanks to the repeated use of the `nbr` operator.

As Figure 7.3 shows, once the simulation starts, a chain of devices is rapidly identified (red dots), marking a sequence of way-points for both device owners to walk in order to meet in the middle. Note also that, due to the ongoing nature of the computation, if one of the device owners moves in a different direction instead, the path will automatically adjust so that it continues to recommend the best path for rendezvous.

### 7.4.2 Network Service Management

One of the common problems in managing complex enterprise services is that there are often many dependencies between different servers and services. Frequently, some of these services are legacy or poorly coded, such that they do not respond gracefully to the failure of their dependencies. These services may continue to attempt to operate for some time, creating inconsistent state, or may be unable to resume service correctly after the server they depend on is brought back on line.

Thus, responding to a service failure often requires a coordinated shutdown and restart of services in an order dictated by service dependencies. This type of service management can be automated by attaching a daemon that watches the state of each service, then communicates with the daemons of other services to coordinate shutdown and restart in accordance with their dependencies.

Figure 7.4(a) shows an example scenario of an enterprise network for a small manufacturing and supply company, with dependencies between two key databases and the internal and external servers running web applications. This scenario was implemented on a network of EmuLab [WLS<sup>+</sup>02] servers. The services were emulated as simple query-response networking programs in Java that entered a “hung” state either upon being externally triggered to crash or after their queries began to consistently fail.

Each service was wrapped with an embedded Protelis execution engine, which was interfaced with the services by a small piece of monitoring glue code that inserted environment variables containing an identifier for the `serviceID` running on that server, a tuple of identifiers for dependencies, and the current `managedServiceStatus` of stop, starting, run,

stopping, or hung. The glue code also provides `stopService` and `startService` methods to send signals to the service, tracks interactions between the services in order to maintain the set of neighbors for Protelis, and allows an external monitoring application to attach and receive status reports.

Dependency-directed coordination of service starting and stopping was then implemented as follows:

```

import it.unibo.alchemist.language.protelis.datatype.Tuple. *
import com.bbn.a3.distributedrestart.DaemonNode. *

// Compare required and available services
let nbr_set = unionHood(nbr([serviceID]));
let nbr_missing = dependencies.subtract(nbr_set);
let nbr_required = #contains(dependencies, nbr(serviceID));
let nbr_down = nbr(managedServiceStatus=="hung" ||
                   managedServiceStatus=="stop");

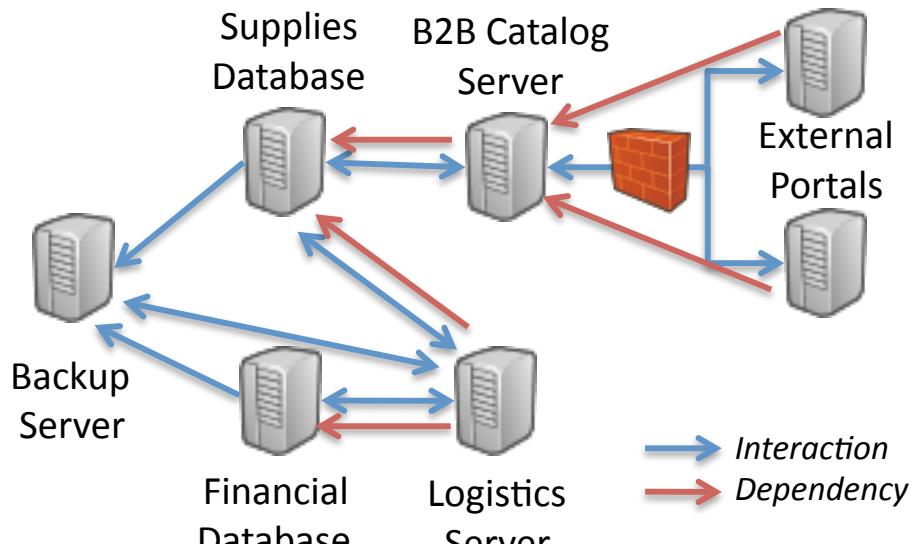
// Is service currently safe to run?
let problem = anyHood(nbr_down && nbr_required) ||
              !nbr_missing.isEmpty();

// Take managed service up and down accordingly
if (managedServiceStatus=="run" && problem) {
    #stopProcess(managedService);
} else {
    if (managedServiceStatus=="stop" && !problem) {
        #startProcess(managedService);
    } else {
        managedServiceStatus
    }
}

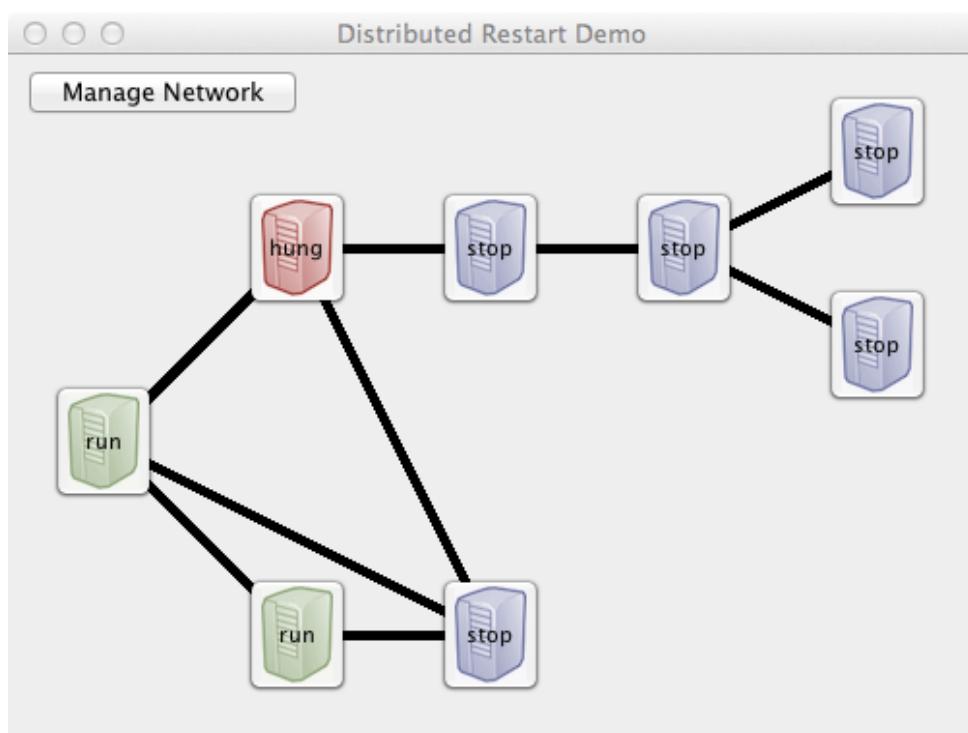
```

In this program, each device shares information about its service ID and status with its neighbors, enabling them to track which dependencies are currently down or missing. When there is a problem with dependencies, the device invokes `stopProcess` to shut its service down, when dependencies are good, it brings it up again with `startProcess`, and when it is hung it waits for a human to sort out the problem.

Figure 7.4(b) shows a typical screenshot of the network of services in operation on an EmuLab network of Ubuntu machines, one service per machine, as visualized by the monitoring application. In this screenshot, the supplies database has crashed, causing many of the other services to gracefully shut themselves down. As soon as the supplies database is restarted, however, the rest of the services automatically bring themselves up in dependency order.



(a) Example Dependent Services Scenario



(b) Example of Coordinated Restart Execution

Figure 7.4: (a) An example scenario of an enterprise network for a small manufacturing and supply company. (b) Example of execution on a network of 8 EmuLab [WLS<sup>+</sup>02] machines: the supplies database has crashed (red), and so all dependent services have shut themselves down (blue), while other services continue to run normally (green).

## **Part III**

## **Conclusion**



# 8

## Results achieved

### 8.1 Integrated toolchain for pervasive ecosystems

### 8.2 Aggregate programming languages

#### 8.2.1 Higher Order Functions in Field Calculus

#### 8.2.2 Protelis

Protelis ensures universality and coherence between aggregate specification and local execution by building atop the field calculus introduced in [VDB13]. At the same time, accessibility, portability, and ease of integration are ensured by embedding Protelis within Java. This enables Protelis programs to draw on the full breadth of available Java APIs and to readily integrate with a wide range of devices and applications, as illustrated by our examples of pervasive computing simulation and networked service management. This implementation of Protelis thus forms an important component of the toolchain necessary for practical application of aggregate programming principles and methods to address real-world problems.



# 9

## Future and ongoing work

### 9.1 Aggregate programming

The Protelis framework continues to be actively developed: we plan to enrich it in the future by adding higher-level abstractions for aggregate programming grounded on the mechanisms discussed in this work.

Also, we observe that the current situation is much like that of cooperative multitasking in the operating systems world, before the development of the current kernel-user distinction and pre-emptive multitasking that dominates the world now. We think that Protelis could be a new architecture for decentralized pre-emptive management of distributed processes, which provides these same sort of capabilities for distributed systems. In particular, we want to exploit the field calculus model for aggregate programming of distributed systems, extending it to handle first-class functions, which enables its alignment mechanism for pre-emptive control of the scope of a distributed computation to be used for process management. These mechanisms enable the construction of a prototype “thread manager” for pervasive and distributed systems, including the ability to preempt a process, to kill a rogue process, or to dynamically manage the scope of a process. This point of view will be deeply analysed a paper which is an ongoing work.

### 9.2 Biochemical meta model for ALCHEMIST

ciao



# Bibliography

- [AAC<sup>+</sup>00] Harold Abelson, Don Allen, Daniel Coore, Chris Hanson, George Homsy, Thomas F. Knight, Jr., Radhika Nagpal, Erik Rauch, Gerald Jay Sussman, and Ron Weiss. Amorphous computing. *Communications of the ACM*, 43:74–82, 2000.
- [AARC<sup>+</sup>12] Seyedreza Abdollahi, Hamed Al-Raweshidy, Erik Calvo, Luis Manuel Dahlman, Klaus David, Marian Codreanu, Panagiotis Demestichas, Markus Dillinger, Miguel Egido, Markus Gruber, Stefan Kaiser, Matti Latva-aho, Peter Merz, Xavier Mestre, Emmanouil Panaousis, Stephan Pfletschinger, Mahdi Pirmoradian, Christos Politis, Bernhard Raaf, Frank Schaich, Christoph Schmelz, Rahim Tafazolli, Antti Tölli, Nikola Vucic, and Rudolf Winkelmann. Spectrum Crunch. Technical report, Net!Works European Technology Platform, 06 2012.
- [ABI09] Marco Autili, Paolo Benedetto, and Paola Inverardi. Context-aware adaptive services: The plastic approach. In *FASE '09 Proceedings*, pages 124–139, Berlin, Heidelberg, 2009. Springer-Verlag.
- [AGS11] Michael Adeyeye and Paul Gardner-Stephen. The village telco project: a reliable and practical wireless mesh telephony infrastructure. *EURASIP J. Wireless Comm. and Networking*, 2011:78, 2011.
- [AJL<sup>+</sup>02] Bruce Alberts, Alexander Johnson, Julian Lewis, Martin Raff, Keith Roberts, and Peter Walter. *Molecular Biology of the Cell*. Garland Science Textbooks. Garland Science, Fourth edition, June 2002.
- [AM11] Musab AlTurki and José Meseguer. Pvesta: A parallel statistical model checking and quantitative analysis tool. In *CALCO*, 2011.
- [AMS06] Gul A. Agha, José Meseguer, and Koushik Sen. PMaude: Rewrite-based specification language for probabilistic object systems. In Antonio Cerone and Herbert Wiklicky, editors, *QAPL 2005*, volume 153(2) of *ENTCS*, pages 213–239. Elsevier, 2006.
- [Apa13] Tatjana Apanasevic. Factors influencing the slow rate of penetration of NFC mobile payment in western europe. In *12th International Conference on Mobile Business, ICMB 2013, Berlin, Germany, 10-13 June 2013*, page 8, 2013.
- [APNF13] Bernhard Anzengruber, Danilo Pianini, Jussi Nieminen, and Alois Ferscha. Predicting social density in mass events to prevent crowd disasters. In Adam Jatowt, Ee-Peng Lim, Ying Ding, Asako Miura, Taro Tezuka, Gaël Dias, Katsumi Tanaka, Andrew J. Flanagan, and Bing Tian Dai, editors, *SocInfo*, volume 8238 of *Lecture Notes in Computer Science*, pages 206–215. Springer, 2013.
- [ARGL<sup>+</sup>07] Michael P. Ashley-Rollman, Seth Copen Goldstein, Peter Lee, Todd C. Mowry, and Padmanabhan Pillai. Meld: A declarative approach to programming ensembles. In *IEEE International Conference on Intelligent Robots and Systems (IROS '07)*, pages 2794–2800, 2007.

- [ASS<sup>+</sup>11] Henrik Abramowicz, Ernst-Dieter Schmidt, Lars Christoph Schmelz, Cornel Pampu, Cornelia Kappler, Mirko Schramm, Konstantinos Pentikousis, Dongming Zhou, Savo Glisic, Juan Palacios, Josep Mangues, Panagiotis Demestichas, Ralf Toenjes, Tipu Arvind Ramrekha, Grant Millar, Christos Politis, Sergi Figuerola, Dimitra Simeonidou, Alex Galis, and Rahim Tafazolli. Future Networks and Management . Technical report, Net!Works European Technology Platform, 06 2011.
- [BB06] Jacob Beal and Jonathan Bachrach. Infrastructure for engineered emergence on sensor/actuator networks. *IEEE Intelligent Systems*, 21(2):10–19, 2006.
- [BB07] Jonathan Bachrach and Jacob Beal. Building spatial computers. Technical Report MIT-CSAIL-TR-2007-017, MIT, March 2007.
- [BBL<sup>+</sup>11] Kevin Burrage, PamelaM. Burrage, André Leier, Tatiana Marquez-Lago, and Jr Nicolau, DanV. Stochastic simulation for spatial modelling of dynamic processes in a living cell. In Heinz Koeppl, Gianluca Setti, Mario di Bernardo, and Douglas Densmore, editors, *Design and Analysis of Biomolecular Circuits*, pages 43–62. Springer New York, 2011.
- [BBVT08] Jacob Beal, Jonathan Bachrach, Daniel Vickery, and Mark Tobenkin. Fast self-healing gradients. In *Proceedings of ACM SAC 2008*, pages 1969–1975, 2008.
- [BCD<sup>+</sup>06] Özalp Babaoglu, Geoffrey Canright, Andreas Deutsch, Gianni Di Caro, Frederick Ducatelle, Luca Maria Gambardella, Niloy Ganguly, Márk Jelasity, Roberto Montemanni, Alberto Montresor, and Tore Urnes. Design patterns from biology for distributed computing. *TAAS*, 1(1):26–66, 2006.
- [BDT99] Eric Bonabeau, Marco Dorigo, and Guy Theraulaz. *Swarm Intelligence: From Natural to Artificial Systems*. Oxford University Press, 1999.
- [BDU<sup>+</sup>13] Jacob Beal, Stefan Dulman, Kyle Usbeck, Mirko Viroli, and Nikolaus Correll. Organizing the aggregate: Languages for spatial computing. In Marjan Mernik, editor, *Formal and Practical Aspects of Domain-Specific Languages: Recent Developments*, chapter 16, pages 436–501. IGI Global, 2013. A longer version available at: <http://arxiv.org/abs/1202.5509>.
- [Ber92] Gérard Berry. The chemical abstract machine. *Theoretical Computer Science*, 96(1):217–248, April 1992.
- [BFLM01] Jean-Pierre Banâtre, Pascal Fradet, and Daniel Le Métayer. Gamma and the chemical reaction model: Fifteen years after. In Cristian S. Calude, Gheorghe Păun, Grzegorz Rozenberg, and Arto Salomaa, editors, *Multiset Processing. Mathematical, Computer Science, and Molecular Computing Points of View*, volume 2235 of *LNCS*, pages 17–44. Springer, 2001.
- [BLM90] Jean-Pierre Banâtre and Daniel Le Métayer. The GAMMA model and its discipline of programming. *Science of Computer Programming*, 15(1):55–77, November 1990.

- [BMM02] Özalp Babaoglu, Hein Meling, and Alberto Montresor. Anthill: A framework for the development of agent-based peer-to-peer systems. In *Proceedings of the 22 Nd International Conference on Distributed Computing Systems (ICDCS'02)*, ICDCS '02, pages 15–, Washington, DC, USA, 2002. IEEE Computer Society.
- [BMP<sup>+</sup>11] Chatschik Bisdikian, Bernhard Mitschang, Dino Pedreschi, Vincent S. Tseng, and Claudio Bettini. Challenges for mobile data management in the era of cloud and social computing. In *12th IEEE International Conference on Mobile Data Management*, page 6, 2011.
- [BMV09a] Stefania Bandini, Sara Manzoni, and Giuseppe Vizzari. Agent based modeling and simulation: An informatics perspective. *Journal of Artificial Societies and Social Simulation*, 12:4, 2009.
- [BMV09b] Stefania Bandini, Sara Manzoni, and Giuseppe Vizzari. Crowd Behavior Modeling: From Cellular Automata to Multi-Agent Systems. In Adelinde M. Uhrmacher and Danny Weyns, editors, *Multi-Agent Systems: Simulation and Applications*, Computational Analysis, Synthesis, and Design of Dynamic Systems, chapter 13, pages 389–418. CRC Press, June 2009.
- [Bon99] Eric Bonabeau. Editor’s introduction: Stigmergy. *Artificial Life*, 5(2):95–96, Spring 1999.
- [BP09] Jean-Pierre Banâtre and Thierry Priol. Chemical programming of future service-oriented architectures. *JSW*, 4(7):738–746, 2009.
- [BSB<sup>+</sup>03] J. Mark Bull, L. A. Smith, C. Ball, L. Pottage, and R. Freeman. Benchmarking java against c and fortran for scientific applications. *Concurrency and Computation: Practice and Experience*, 15(3-5):417–430, 2003.
- [But02] William Butera. *Programming a Paintable Computer*. PhD thesis, MIT, Cambridge, MA, USA, 2002.
- [BV07] Stefania Bandini and Giuseppe Vizzari. Regulation function of the environment in agent-based simulation. In *Proceedings of the 3rd international conference on Environments for multi-agent systems III*, E4MAS’06, pages 157–169, Berlin, Heidelberg, 2007. Springer-Verlag.
- [BV14] Jake Beal and Mirko Viroli. Building blocks for aggregate programming of self-organising applications. In *2nd FoCAS Workshop on Fundamentals of Collective Systems*, pages 1–6. IEEE CS, to appear, 2014.
- [BVD14] Jacob Beal, Mirko Viroli, and Ferruccio Damiani. Towards a unified model of spatial computing. In *7th Spatial Computing Workshop (SCW 2014)*, AAMAS 2014, Paris, France, May 2014.

- [BW08] Eduard Babulak and Ming Wang. Discrete event simulation: State of the art. *iJOE*, 4(2):60–63, 2008.
- [CDB<sup>+</sup>12] Marco Conti, Sajal K. Das, Chatschik Bisikian, Mohan Kumar, Lionel M. Ni, Andrea Passarella, George Roussos, Gerhard Tröster, Gene Tsudik, and Franco Zambonelli. Looking ahead in pervasive computing: Challenges and opportunities in the era of cyber-physical convergence. *Pervasive and Mobile Computing*, 8(1):2–21, 2012.
- [CG09] Federica Ciocchetta and Maria Luisa Guerriero. Modelling biological compartments in bio-pepa. *Electr. Notes Theor. Comput. Sci.*, 227:77–95, 2009.
- [Chu32] Alonzo Church. A set of postulates for the foundation of logic. *Annals of mathematics*, 33(2):346–366, 1932.
- [CLMZ03] Giacomo Cabri, Letizia Leonardi, Marco Mamei, and Franco Zambonelli. Location-dependent services for mobile users. *IEEE Transactions on Systems, Man, and Cybernetics, Part A*, 33(6):667–681, 2003.
- [Cow00] Richard Cowan. Stochastic models for dna replication. In *In Handbook of Statistics*. Elsevier, 2000.
- [CVG09] Matteo Casadei, Mirko Viroli, and Luca Gardelli. On the collective sort problem for distributed tuple spaces. *Science of Computer Programming*, 74(9):702–722, 2009.
- [DDF<sup>+</sup>06] Simon Dobson, Spyros Denazis, Antonio Fernández, Dominique Gaïti, Erol Gelenbe, Fabio Massacci, Paddy Nixon, Fabrice Saffre, Nikita Schmidt, and Franco Zambonelli. A survey of autonomic communications. *ACM Transactions on Autonomous and Adaptive Systems*, 1(2):223–259, December 2006.
- [DG08] Jeffrey Dean and Sanjay Ghemawat. MapReduce: simplified data processing on large clusters. *Communications of the ACM*, 51(1):107–113, 2008.
- [DS04] Marco Dorigo and Thomas Stützle. *Ant Colony Optimization*. MIT Press, Cambridge, MA, July 2004.
- [EAWS12] Thaddeus Eze, Richard Anthony, Chris Walshaw, and Alan Soper. Autonomic computing in the first decade: trends and direction. In *ICAS 2012, The Eighth International Conference on Autonomic and Autonomous Systems*, pages 80–85, 2012.
- [EB10] Moritz Eysholdt and Heiko Behrens. Xtext: implement your language faster than the quick and dirty way. In *OOPSLA*, pages 307–309. ACM, 2010.
- [EMRU07] Roland Ewald, Carsten Maus, Arndt Rolfs, and Adelinde M. Uhrmacher. Discrete event modeling and simulation in systems biology. *Journal of Simulation*, 1(2):81–96, 2007.
- [FHA99] Eric Freeman, Susanne Hupfer, and Ken Arnold. *JavaSpaces Principles, Patterns, and Practice: Principles, Patterns and Practices*. The Jini Technology Series. Addison-Wesley Longman, June 1999.

- [FMDMSM<sup>+</sup>12] Jose Luis Fernandez-Marquez, Giovanna Di Marzo Serugendo, Sara Montagna, Mirko Viroli, and Josep Lluis Arcos. Description and composition of bio-inspired design patterns: a complete overview. *Natural Computing*, May 2012. Online First.
- [FSS10] Regina Frei, Giovanna Di Marzo Serugendo, and Traian-Florin Serbanuta. Ambient intelligence in self-organising assembly systems using the chemical reaction model. *Journal of Ambient Intelligence and Humanized Computing*, 1(3):163–184, 2010.
- [GB00] Michael A. Gibson and Jehoshua Bruck. Efficient exact stochastic simulation of chemical systems with many species and many channels. *J. Phys. Chem. A*, 104:1876–1889, 2000.
- [Gel85] David Gelernter. Generative communication in Linda. *ACM Transactions on Programming Languages and Systems*, 7(1):80–112, January 1985.
- [GGG05] Ramakrishna Gummadi, Omprakash Gnawali, and Ramesh Govindan. Macro-programming wireless sensor networks using kairos. In *Distributed Computing in Sensor Systems (DCOSS)*, pages 126–140, 2005.
- [Gil77] Daniel T. Gillespie. Exact stochastic simulation of coupled crfreactions. *The Journal of Physical Chemistry*, 81(25):2340–2361, December 1977.
- [Gil06] Scott F. Gilbert. *Developmental Biology, Eighth Edition*. Sinauer Associates Inc., Eighth edition, March 2006.
- [GJK<sup>+</sup>04] Vitaly V. Gursky, Johannes Jaeger, Konstantin N. Kozlov, John Reinitz, and Alexander M. Samsonov. Pattern formation and nuclear divisions are uncoupled in drosophila segmentation: comparison of spatially discrete and continuous models. *Physica D: Nonlinear Phenomena*, 197(3-4):286–302, October 2004.
- [GJM12] Hongliang Guo, Yaochu Jin, and Yan Meng. A morphogenetic framework for self-organized multirobot pattern formation and boundary coverage. *ACM Transactions on Autonomous and Adaptive Systems*, 7(1):15, 2012.
- [GPO13] Maria Ganzha, Marcin Paprzycki, and Andrea Omicini. Software agents: Twenty years and counting. *Computing Now*, 6(11), November 2013.
- [GPOS13] Pedro Pablo González Pérez, Andrea Omicini, and Marco Sbaraglia. A biochemically-inspired coordination-based model for simulating intracellular signalling pathways. *Journal of Simulation*, 7(3):216–226, August 2013. Special Issue: Agent-based Modeling and Simulation.
- [Gra59] Pierre-Paul Grassé. La reconstruction du nid et les coordinations interindividuelles chez Bellicositermes natalensis et Cubitermes sp. la théorie de la stigmergie: Essai d’interprétation du comportement des termites constructeurs. *Insectes Sociaux*, 6(1):41–80, March 1959.

- [GSP11] Paul Gardner-Stephen and Swapna Palaniswamy. Serval mesh software-wifi multi model management. In *Proceedings of the 1st International Conference on Wireless Technologies for Humanitarian Relief*, ACWR '11, pages 71–77, New York, NY, USA, 2011. ACM.
- [GSW06] Dina Q. Goldin, Scott A. Smolka, and Peter Wegner, editors. *Interactive Computation: The New Paradigm*. Springer, September 2006.
- [GVO07] Luca Gardelli, Mirko Viroli, and Andrea Omicini. Design patterns for self-organising systems. In Hans-Dieter Burkhard, Rineke Verbrugge, and László Zsolt Varga, editors, *Multi-Agent Systems and Applications V*, volume 4696 of *LNAI*, pages 123–132. Springer, September 2007. 5th International Central and Eastern European Conference on Multi-Agent Systems (CEEMAS'07), Leipzig, Germany, 25–27 September 2007, Proceedings.
- [HAES09] M. Shamim Hossain, Atif Alamri, and Abdulmotaleb El Saddik. A biologically inspired framework for multimedia service management in a ubiquitous environment. *Concurrency and Computation: Practice and Experience*, 21(11):1450–1466, 2009.
- [Har87] David Harel. Statecharts: A visual formalism for complex systems. *Sci. Comput. Program.*, 8(3):231–274, June 1987.
- [HDBDM12] Dries Harnie, Theo D'Hondt, Elisa Gonzalez Boix, and Wolfgang De Meuter. Programming urban-area applications. In *Proceedings of the 27th Annual ACM Symposium on Applied Computing*, pages 1516–1521. ACM, 2012.
- [HDF<sup>+</sup>11] Philip N. Howard, Aiden Duffy, Deen Freelon, Muzammil Hussain, Will Mari, and Marwa Mazaid. Opening closed regimes: What was the role of social media during the arab spring? 2011.
- [HQL<sup>+</sup>11] Fei Hu, Meikang Qiu, Jiayin Li, Travis Grant, Drew Taylor, Seth McCaleb, Lee Butler, and Richard Hamner. A review on cloud computing: Design challenges in architecture and security. *CIT. Journal of Computing and Information Technology*, 19(1):25–55, 2011.
- [HSG<sup>+</sup>06] Stefan Hoops, Sven Sahle, Ralph Gauges, Christine Lee, Jörgen Pahle, Natalia Simus, Mudita Singhal, Liang Xu, Pedro Mendes, and Ursula Kummer. Copasi – a complex pathway simulator. *Bioinformatics*, 22(24):3067–3074, 2006.
- [HVUM07] Alexander Helleboogh, Giuseppe Vizzari, Adelinde Uhrmacher, and Fabien Michel. Modeling dynamic environments in multi-agent simulation. *Autonomous Agents and Multi-Agent Systems*, 14(1):87–116, February 2007.
- [HW08] Mordechai (Muki) Haklay and Patrick Weber. Openstreetmap: User-generated street maps. *IEEE Pervasive Computing*, 7(4):12–18, October 2008.

- [IPW01] Atsushi Igarashi, Benjamin C. Pierce, and Philip Wadler. Featherweight java: A minimal core calculus for Java and GJ. *ACM Transactions on Programming Languages and Systems*, 23(3), 2001.
- [Jaz05] Mehdi Jazayeri. Species evolve, individuals age. In Motoshi Saeki, Gerardo Canfora, and Shuichiro Yamamoto, editors, *8th International Workshop on Principles of Software Evolution (IWPSE 2005)*, pages 3–9, Lisbon, Portugal, 5–6 September 2005.
- [Kie02] Andrzej M. Kierzek. Stocks: Stochastic kinetic simulations of biochemical systems with gillespie algorithm. *Bioinformatics*, 18(3):470–481, 2002.
- [KR08] Lila Kari and Grzegorz Rozenberg. The many facets of natural computing. *Communications of the ACM*, 51:72–83, 2008.
- [LCRP<sup>+</sup>05] Sean Luke, Claudio Cioffi-Revilla, Liviu Panait, Keith Sullivan, and Gabriel Catalin Balan. Mason: A multiagent simulation environment. *Simulation*, 81(7):517–527, 2005.
- [LDB09] William J.R. Longabaugh, Eric H. Davidson, and Hamid Bolouri. Visualization, documentation, analysis, and communication of large-scale gene regulatory networks. *Biochimica et Biophysica Acta (BBA) - Gene Regulatory Mechanisms*, 1789(4):363–374, 2009.
- [LIDP10] Paola Lecca, Adaoha E. C. Ihekweaba, Lorenzo Dematté, and Corrado Priami. Stochastic simulation of the spatio-temporal dynamics of reaction-diffusion systems: the case for the bicoid gradient. *J. Integrative Bioinformatics*, 7(1), 2010.
- [LK87] Frank C. H. Lin and Robert M. Keller. The gradient model load balancing method. *IEEE Trans. Softw. Eng.*, 13(1):32–38, 1987.
- [LMG<sup>+</sup>09] Uichin Lee, Eugenio Magistretti, Mario Gerla, Paolo Bellavista, Pietro Lió, and Kang-Won Lee. Bio-inspired multi-agent data harvesting in a proactive urban monitoring environment. *Ad Hoc Networks*, 7:725–741, 2009.
- [LPF12] P. Lukowicz, S. Pentland, and A. Ferscha. From context awareness to socially aware computing. *IEEE Pervasive Computing*, 11(1):32 –41, 2012.
- [LT06] Jiming Liu and Kwok Ching Tsui. Toward nature-inspired computing. *Communications of the ACM*, 49(10):59–64, October 2006.
- [Mac08] Rebecca MacKinnon. Flatter world and thicker walls? blogs, censorship and civic discourse in china. *Public Choice*, 134(1-2):31–46, 2008.
- [MCOV11] Ambra Molesini, Matteo Casadei, Andrea Omicini, and Mirko Viroli. Simulation in agent-oriented software engineering: The SODA case study. *Science of Computer Programming*, August 2011. Special Issue on Agent-oriented Design methods and Programming Techniques for Distributed Computing in Dynamic and Complex Environments.

- [MFHH05] S. Madden, M. Franklin, J. Hellerstein, and W. Hong. Tinydb: An acquisitional query processing system for sensor networks. In *ACM TODS*, 2005.
- [Mil99] Robin Milner. *Communicating and Mobile Systems: The Pi-Calculus*. Cambridge University Press, 1999.
- [MM04] Frank Manola and Eric Miller. RDF primer. W3C recommendation, World Wide Web Consortium, February 2004.
- [MMTL13] Massimo Monti, Thomas Meyer, Christian F. Tschudin, and Marco Luise. Stability and sensitivity analysis of traffic-shaping algorithms inspired by chemical engineering. *IEEE Journal on Selected Areas in Communications*, 31(6):1105–1114, 2013.
- [MMTZ06] M. Mamei, R. Menezes, R. Tolksdorf, and F. Zambonelli. Case studies for self-organization in computer science. *Journal of Systems Architecture*, 52:443–460, 2006.
- [MN98] Makoto Matsumoto and Takuji Nishimura. Mersenne twister: A 623-dimensionally equidistributed uniform pseudo-random number generator. *ACM Trans. Model. Comput. Simul.*, 8(1):3–30, 1998.
- [MN10] Charles M. Macal and Michael J. North. Tutorial on agent-based modelling and simulation. *Journal of Simulation*, 4(3):151–162, 2010.
- [MO13] Stefano Mariani and Andrea Omicini. Molecules of knowledge: Self-organisation in knowledge-intensive environments. *Intelligent Distributed Computing VI*, pages 17–22, 2013.
- [MOV09] Ambra Molesini, Andrea Omicini, and Mirko Viroli. Environment in Agent-Oriented Software Engineering methodologies. *Multiagent and Grid Systems*, 5(1):37–57, 2009. Special Issue “Engineering Environments in Multi-Agent Systems”.
- [MPR06] Amy L. Murphy, Gian Pietro Picco, and Gruia-Catalin Roman. Lime: A model and middleware supporting mobility of hosts and agents. *ACM Trans. on Software Engineering and Methodology*, 15(3):279–328, 2006.
- [MPV12] Sara Montagna, Danilo Pianini, and Mirko Viroli. Gradient-based self-organisation patterns of anticipative adaptation. In *Proceedings of 6th IEEE International Conference on Self-Adaptive and Self-Organizing Systems (SASO 2012)*, pages 169–174, September 2012.
- [MT03] Ronaldo Menezes and Robert Tolksdorf. Adaptiveness in Linda-based coordination models. In Giovanna Di Marzo Serugendo, Anthony Karageorgos, Omer F. Rana, and Franco Zambonelli, editors, *Engineering Self-Organising Systems. Nature-Inspired Approaches to Software Engineering*, volume 2977, pages 212–232. Springer, 2003.

- [MTUG09] Aneil Mallavarapu, Matthew Thomson, Benjamin Ullian, and Jeremy Gunawardena. Programming with models: modularity and abstraction provide powerful capabilities for systems biology. *Journal of the Royal Society, Interface / the Royal Society*, 6(32):257–270, March 2009.
- [Mur89] T. Murata. Petri nets: Properties, analysis and applications. *Proceedings of the IEEE*, 77(4):541–580, April 1989.
- [MV10] Sara Montagna and Mirko Viroli. A framework for modelling and simulating networks of cells. *Electronic Notes in Theoretical Computer Science*, 268:115–129, December 2010. Proceedings of the 1st International Workshop on Interactions between Computer Science and Biology (CS2Bio’10).
- [MVF<sup>M</sup><sup>+</sup>13] Sara Montagna, Mirko Viroli, Jose Luis Fernandez-Marquez, Giovanna Di Marzo Serugendo, and Franco Zambonelli. Injecting self-organisation into pervasive service ecosystems. *Mobile Networks and Applications*, 18(3):398–412, 2013.
- [MYT07] Thomas Meyer, Lidia Yamamoto, and Christian F. Tschudin. An artificial chemistry for networking. In *First Workshop on Bio-Inspired Design of Networks, BIOWIRE 2007, Cambridge, UK*, pages 45–57, 2007.
- [MZ05] Marco Mamei and Franco Zambonelli. Programming stigmergic coordination with the TOTA middleware. In *4th International Joint Conference on Autonomous Agents and Multiagent Systems (AAMAS 2005)*, pages 415–422, New York, NY, USA, 2005. ACM.
- [MZ06a] Marco Mamei and Franco Zambonelli. *Field-Based Coordination for Pervasive Multiagent Systems. Models, Technologies, and Applications*. Springer Series in Agent Technology. Springer, March 2006.
- [MZ06b] Marco Mamei and Franco Zambonelli. Self-maintained distributed tuples for field-based coordination in dynamic networks. *Concurrency and Computation: Practice and Experience*, 18(4):427–443, 2006.
- [MZ07] Marco Mamei and Franco Zambonelli. Pervasive pheromone-based interaction with RFID tags. *TAAS*, 2(2), 2007.
- [MZ09] Marco Mamei and Franco Zambonelli. Programming pervasive and mobile computing applications: The tota approach. *ACM Trans. Softw. Eng. Methodol.*, 18(4):1–56, 2009.
- [MZL03] Marco Mamei, Franco Zambonelli, and Letizia Leonardi. Co-fields: Towards a unifying approach to the engineering of swarm intelligent systems. In Paolo Petta, Robert Tolksdorf, and Franco Zambonelli, editors, *Engineering Societies in the Agents World III*, volume 2577 of *LNCS*, pages 68–81. Springer, April 2003. 3rd International Workshop (ESAW 2002), Madrid, Spain, 16–17 September 2002. Revised Papers.
- [Nag01] Radhika Nagpal. *Programmable Self-Assembly: Constructing Global Shape using Biologically-inspired Local Interactions and Origami Mathematics*. PhD thesis, MIT, Cambridge, MA, USA, 2001.

- [NHCV07] Michael J. North, Tom R. Howe, Nicholson T. Collier, and Jerry R. Vos. A declarative model assembly infrastructure for verification and validation. In Shingo Takahashi, David Sallach, and Juliette Rouchier, editors, *Advancing Social Simulation: The First World Congress*, pages 129–140. Springer Japan, 2007.
- [NLPT14] Rocco De Nicola, Michele Loreti, Rosario Pugliese, and Francesco Tiezzi. A formal approach to autonomic systems programming: The scel language. *ACM Trans. Auton. Adapt. Syst.*, 9(2):7:1–7:29, July 2014.
- [NW04] Ryan Newton and Matt Welsh. Region streams: Functional macroprogramming for sensor networks. In *First International Workshop on Data Management for Sensor Networks (DMSN)*, pages 78–87, August 2004.
- [NZ01] Michael S. Noble and Stoyanka Zlateva. Scientific computation with javaspaces. In *Proceedings of the 9th International Conference on High-Performance Computing and Networking*, pages 657–666. Springer-Verlag, 2001.
- [OP07] Jan Ondrus and Yves Pigneur. An assessment of nfc for future mobile payment systems. In *Proceedings of the International Conference on the Management of Mobile Business*, ICMB ’07, pages 43–, Washington, DC, USA, 2007. IEEE Computer Society.
- [ORAI11] Bogdan Oancea, Ion Gh. Rosca, Tudorel Andrei, and Andreea Iluzia Iacob. Evaluating java performance for linear algebra numerical computations. *Procedia CS*, 3:474–478, 2011.
- [ORV06] Andrea Omicini, Alessandro Ricci, and Mirko Viroli. The multidisciplinary patterns of interaction from sciences to Computer Science. In *Interactive Computation: The New Paradigm*, pages 395–414. Springer, September 2006.
- [ORV08] Andrea Omicini, Alessandro Ricci, and Mirko Viroli. Artifacts in the A&A meta-model for multi-agent systems. *Autonomous Agents and Multi-Agent Systems*, 17(3):432–456, December 2008. Special Issue on Foundations, Advanced Topics and Industrial Perspectives of Multi-Agent Systems.
- [OZ99] Andrea Omicini and Franco Zambonelli. Coordination for Internet application development. *Autonomous Agents and Multi-Agent Systems*, 2(3):251–269, September 1999.
- [Par06] H. Van Dyke Parunak. A survey of environments and mechanisms for human-human stigmergy. In Danny Weyns, H. Van Dyke Parunak, and Fabien Michel, editors, *Environments for Multi-Agent Systems II*, volume 3830 of *LNCS*, pages 163–186. Springer, 2006.
- [PBS02] H. Van Dyke Parunak, Sven Brueckner, and John Sauter. Digital pheromone mechanisms for coordination of unmanned vehicles. In Cristiano Castelfranchi and W. Lewis

- Johnson, editors, *1st International Joint Conference on Autonomous Agents and Multiagent systems*, volume 1, pages 449–450, New York, NY, USA, 15–19July 2002. ACM.
- [PH14] T.-T.T. Pham and J.C. Ho. What are the core drivers in consumer adoption of nfc-based mobile payments?: A proposed research framework. In *Management of Engineering Technology (PICMET), 2014 Portland International Conference on*, pages 3041–3049, July 2014.
  - [PJRGO6] Theodore J Perkins, Johannes Jaeger, John Reinitz, and Leon Glass. Reverse engineering the gap gene network of *Drosophila Melanogaster*. *PLoS Computational Biology*, 2(5):e51, 05 2006.
  - [PMV13] Danilo Pianini, Sara Montagna, and Mirko Viroli. Chemical-oriented simulation of computational systems with Alchemist. *Journal of Simulation*, 2013.
  - [Pol89] Lissa F. Pollacia. A survey of discrete event simulation and state-of-the-art discrete event languages. *SIGSIM Simul. Dig.*, 20(3):8–25, September 1989.
  - [PPSR09] Andrei Pisarev, Ekaterina Poustelnikova, Maria Samsonova, and John Reinitz. Flyex, the quantitative atlas on segmentation gene expression at cellular resolution. *Nucleic Acids Research*, 37(Database-Issue):560–566, 2009.
  - [Pri95] Corrado Priami. Stochastic pi-calculus. *Comput. J.*, 38(7):578–589, 1995.
  - [PRSS01] Corrado Priami, Aviv Regev, Ehud Shapiro, and William Silverman. Application of a stochastic name-passing calculus to representation and simulation of molecular processes. *Information Processing Letters*, 80:25–31, October 2001.
  - [PVM<sup>+</sup>10] Danilo Pianini, Sascia Virruso, Ronaldo Menezes, Andrea Omicini, and Mirko Viroli. Self organization in coordination systems using a wordnet-based ontology. In *SASO*, pages 114–123, 2010.
  - [RCD01] Davide Rossi, Giacomo Cabri, and Enrico Denti. Tuple-based technologies for coordination. In Andrea Omicini, Franco Zambonelli, Matthias Klusch, and Robert Tolksdorf, editors, *Coordination of Internet Agents: Models, Technologies, and Applications*, chapter 4, pages 83–109. Springer-Verlag, March 2001.
  - [RCKZ13] Vaskar Raychoudhury, Jiannong Cao, Mohan Kumar, and Daqiang Zhang. Middleware for pervasive computing: A survey. *Pervasive and Mobile Computing*, 9(2):177 – 200, 2013.
  - [RHR<sup>+</sup>08] Peter Van Roy, Seif Haridi, Alexander Reinefeld, Jean-Bernard Stefany, Roland Yap, and Thierry Coupaye. Self-management for large-scale distributed systems: an overview of the selfman project. In *Formal Methods for Components and Objects, LNCS No. 5382*, pages 153–178. Springer Verlag, 2008.

- [RPJ96] Rolando Rivera-Pomar and Herbert Jackle. From gradients to stripes in drosophila embryogenesis: filling in the gaps. *Trends in Genetics*, 12(11):478 – 483, 1996.
- [RPV11] Alessandro Ricci, Michele Piunti, and Mirko Viroli. Environment programming in multi-agent systems – an artifact-based perspective. *Autonomous Agents and Multi-Agent Systems*, 23(2):158–192, September 2011.
- [RS95] John Reinitz and David H. Sharp. Mechanism of eve stripe formation. *Mechanisms of Development*, 49(1-2):133–158, January 1995.
- [SGJ07] Michael Schumacher, Laurent Grangier, and Radu Jurca. Governing environments for agent-based traffic simulations. In *Proceedings of the 5th international Central and Eastern European conference on Multi-Agent Systems and Applications V*, CEEMAS ’07, pages 163–172, Berlin, Heidelberg, 2007. Springer-Verlag.
- [Sha04] Nigel Shadbolt. From the Editor in Chief: Nature-Inspired Computing. *IEEE Intelligent Systems*, 19(1):2–3, January–February 2004.
- [SJ07] Drew Stovall and Christine Julien. Resource discovery with evolving tuples. In *International Workshop on Engineering of Software Services for Pervasive Environments: In Conjunction with the 6th ESEC/FSE Joint Meeting*, ESSPE ’07, pages 1–10, New York, NY, USA, 2007. ACM.
- [SKK<sup>+</sup>08] Svetlana Surkova, David Kosman, Konstantin Kozlov, Manu, Ekaterina Myasnikova, Anastasia A. Samsonova, Alexander Spirov, Carlos E. Vanario-Alonso, Maria Samsonova, and John Reinitz. Characterization of the Drosophila segment determination morphome. *Developmental Biology*, 313(2):844 – 862, 2008.
- [Skl07] Elizabeth Sklar. Netlogo, a multi-agent simulation environment. *Artificial Life*, 13(3):303–311, 2007.
- [SMA05] Koushik Sen, Viswanathan M., and Gul Agha. Vesta: A statistical model-checker and analyzer for probabilistic systems. In *Quantitative Evaluation of Systems, 2005.*, pages 251–252, 2005.
- [SRM<sup>+</sup>13] Daniel Schuster, Alberto Rosi, Marco Mamei, Thomas Springer, Markus Endler, and Franco Zambonelli. Pervasive social context: Taxonomy and survey. *ACM Transactions on Intelligent Systems and Technology*, 4(3), 2013.
- [STP08] Alexander Slepoy, Aidan P. Thompson, and Steven J. Plimpton. A constant-time kinetic monte carlo algorithm for simulation of large biochemical reaction networks. *The Journal of Chemical Physics*, 128(20):205101, 2008.
- [SV13] Stefano Sebastio and Andrea Vandin. MultiVeStA: Statistical model checking for discrete event simulators. In *VALUETOOLS*, 2013.

- [SYD<sup>+</sup>13] Graeme Stevenson, Juan Ye, Simon Dobson, Danilo Pianini, Sara Montagna, and Mirko Viroli. Combining self-organisation, context-awareness and semantic reasoning: the case of resource discovery in opportunistic networks. In Sung Y. Shin and José Carlos Maldonado, editors, *SAC*, pages 1369–1376, Coimbra, Portugal, 18–22 March 2013. ACM.
- [TB99] Guy Theraulaz and Eric Bonabeau. A brief history of stigmergy. *Artificial Life*, 5(2):97–116, Spring 1999.
- [TM03] Robert Tolksdorf and Ronaldo Menezes. Using swarm intelligence in linda systems. In Andrea Omicini, Paolo Petta, and Jeremy Pitt, editors, *ESAW*, volume 3071 of *Lecture Notes in Computer Science*, pages 49–65. Springer, 2003.
- [TOCH14] Garry Wei-Han Tan, Keng-Boon Ooi, Siong-Choy Chong, and Teck-Soon Hew. Nfc mobile credit card: The next frontier of mobile payment? *Telemat. Inf.*, 31(2):292–307, May 2014.
- [UP05] Adelinde M. Uhrmacher and Corrado Priami. Discrete event systems specification in systems biology - a discussion of stochastic pi calculus and devs. In *Proceedings of the 37th conference on Winter simulation*, WSC ’05, pages 317–326. Winter Simulation Conference, 2005.
- [VB08] Cristian Versari and Nadia Busi. Efficient stochastic simulation of biological systems with multiple variable volumes. *Electronic Notes in Theoretical Computer Science*, 194(3):165–180, 2008.
- [VC09] Mirko Viroli and Matteo Casadei. Biochemical tuple spaces for self-organising coordination. In John Field and Vasco Thudichum Vasconcelos, editors, *COORDINATION*, volume 5521 of *Lecture Notes in Computer Science*, pages 143–162. Springer, 2009.
- [VCMZ11] Mirko Viroli, Matteo Casadei, Sara Montagna, and Franco Zambonelli. Spatial coordination of pervasive services through chemical-inspired tuple spaces. *ACM Transactions on Autonomous and Adaptive Systems*, 6(2):14:1 – 14:24, June 2011.
- [VCNO10] Mirko Viroli, Matteo Casadei, Elena Nardini, and Andrea Omicini. Towards a chemical-inspired infrastructure for self-\* pervasive applications. In Danny Weyns, Sam Malek, Rogério de Lemos, and Jesper Andersson, editors, *Self-Organizing Architectures*, volume 6090 of *LNCS*, chapter 8, pages 152–176. Springer, July 2010. 1st International Workshop on Self-Organizing Architectures (SOAR 2009), Cambridge, UK, 14-17 September 2009, Revised Selected and Invited Papers.
- [VCO09] Mirko Viroli, Matteo Casadei, and Andrea Omicini. A framework for modelling and implementing self-organising coordination. In Sung Y. Shin and Sascha Ossowski, editors, *SAC*, pages 1353–1360. ACM, 2009.

- [VD14] Mirko Viroli and Ferruccio Damiani. A calculus of self-stabilising computational fields. In eva Kühn and Rosario Pugliese, editors, *16th Conference on Coordination Languages and Models (Coordination 2014)*, volume 8459 of *LNCS*, pages 163–178. Springer-Verlag, June 2014. Proceedings of the 16th Conference on Coordination Models and Languages (Coordination 2014), Berlin (Germany), 3-5 June.
- [VDB13] Mirko Viroli, Ferruccio Damiani, and Jacob Beal. A calculus of computational fields. In Carlos Canal and Massimo Villari, editors, *Advances in Service-Oriented and Cloud Computing*, volume 393 of *Communications in Computer and Information Science*, pages 114–128. Springer Berlin Heidelberg, 2013.
- [VHR<sup>+</sup>07] Mirko Viroli, Tom Holvoet, Alessandro Ricci, Kurt Schelfhout, and Franco Zambonelli. Infrastructures for the environment of multiagent systems. *Autonomous Agents and Multi-Agent Systems*, 14(1):49–60, July 2007.
- [VPB12] Mirko Viroli, Danilo Pianini, and Jacob Beal. Linda in space-time: an adaptive coordination model for mobile ad-hoc environments. In Marjan Sirjani, editor, *Coordination Languages and Models*, volume 7274 of *LNCS*, pages 212–229. Springer-Verlag, June 2012. Proceedings of the 14th Conference of Coordination Models and Languages (Coordination 2012), Stockholm (Sweden), 14-15 June.
- [VPMS12] Mirko Viroli, Danilo Pianini, Sara Montagna, and Graeme Stevenson. Pervasive ecosystems: a coordination model based on semantic chemistry. In Sascha Ossowski, Paola Lecca, Chih-Cheng Hung, and Jiman Hong, editors, *27th Annual ACM Symposium on Applied Computing (SAC 2012)*, Riva del Garda, TN, Italy, 26–30 March 2012. ACM.
- [VZ10] Mirko Viroli and Franco Zambonelli. A biochemical approach to adaptive service ecosystems. *Information Sciences*, 180(10):1876–1892, May 2010.
- [WBH06] Danny Weyns, Nelly Boucké, and Tom Holvoet. Gradient field-based task assignment in an agv transportation system. In *Proceedings of the fifth international joint conference on Autonomous agents and multiagent systems*, AAMAS ’06, pages 842–849, New York, NY, USA, 2006. ACM.
- [WBYI08] Justin Werfel, Yaneer Bar-Yam, and Donald Ingber. Bioinspired environmental coordination in spatial computing systems. In *1st International SASO Workshop on Spatial Computing*, pages 338–343, Washington, DC, USA, 2008. IEEE Computer Society.
- [Weg97] Peter Wegner. Why interaction is more powerful than algorithms. *Commun. ACM*, 40(5):80–91, May 1997.
- [WLS<sup>+</sup>02] Brian White, Jay Lepreau, Leigh Stoller, Robert Ricci, Shashi Guruprasad, Mac Newbold, Mike Hibler, Chad Barb, and Abhijeet Joglekar. An integrated experimental environment for distributed systems and networks. pages 255–270, Boston, MA, December 2002.

- [WMLF98] Peter Wyckoff, Stephen W. McLaughry, Tobin J. Lehman, and Daniel A. Ford. T Spaces. *IBM Systems Journal*, 37(3):454–474, 1998.
- [Yam07] Daniel Yamins. *A Theory of Local-to-Global Algorithms for One-Dimensional Spatial Multi-Agent Systems*. PhD thesis, Harvard, Cambridge, MA, USA, December 2007.
- [Zam12] Franco Zambonelli. Toward sociotechnical urban superorganisms. *IEEE Computer*, 47(8), 2012.
- [ZCF<sup>+</sup>11] Franco Zambonelli, Gabriella Castelli, Laura Ferrari, Marco Mamei, Alberto Rosi, Giovanna Di Marzo, Matteo Risoldi, Akla-Essø Tchao, Simon Dobson, Graeme Stevenson, Yuan Ye, Elena Nardini, Andrea Omicini, Sara Montagna, Mirko Viroli, Alois Ferscha, Sascha Maschek, and Bernhard Wally. Self-aware pervasive service ecosystems. *Procedia Computer Science*, 7:197–199, December 2011. Proceedings of the 2nd European Future Technologies Conference and Exhibition 2011 (FET 11).
- [Zei76] Bernard P. Zeigler. *Theory of Modeling and Simulation*. John Wiley, 1976.
- [Zei84] B. Zeigler. *Multifaceted modelling and Discrete Event Simulation*. Academic Press, 1984.
- [ZO04] Franco Zambonelli and Andrea Omicini. Challenges and research directions in agent-oriented software engineering. *Autonomous Agents and Multi-Agent Systems*, 9(3):253–283, 2004.
- [ZOA<sup>+</sup>14] Franco Zambonelli, Andrea Omicini, Bernhard Anzengruber, Gabriella Castelli, Francesco L. DeAngelis, Giovanna Di Marzo Serugendo, Simon Dobson, Jose Luis Fernandez-Marquez, Alois Ferscha, Marco Mamei, Stefano Mariani, Ambra Molesini, Sara Montagna, Jussi Nieminen, Danilo Pianini, Matteo Risoldi, Alberto Rosi, Graeme Stevenson, Mirko Viroli, and Juan Ye. Developing pervasive multi-agent systems with nature-inspired coordination. *Pervasive and Mobile Computing*, 2014. Special Issue on “10 years of Pervasive Computing” in honor of Chatschik Bisdikian.
- [ZV11] Franco Zambonelli and Mirko Viroli. A survey on nature-inspired metaphors for pervasive service ecosystems. *International Journal of Pervasive Computing and Communications*, 7(3):186–204, 2011.